Get rich quick: exact solutions reveal how unbalanced initializations promote rapid feature learning

Daniel Kunin[*1] **Allan Raventós**[*1] **Clémentine Dominé**[2] **Feng Chen**[1]
David Klindt[3] **Andrew Saxe**[2] **Surya Ganguli**[1]

[1]Stanford University [2]University College London [3]Cold Spring Harbor Laboratory

Abstract

While the impressive performance of modern neural networks is often attributed to their capacity to efficiently extract task-relevant features from data, the mechanisms underlying this *rich feature learning regime* remain elusive, with much of our theoretical understanding stemming from the opposing *lazy regime*. In this work, we derive exact solutions to a minimal model that transitions between lazy and rich learning, precisely elucidating how unbalanced *layer-specific* initialization variances and learning rates determine the degree of feature learning. Our analysis reveals that they conspire to influence the learning regime through a set of conserved quantities that constrain and modify the geometry of learning trajectories in parameter and function space. We extend our analysis to more complex linear models with multiple neurons, outputs, and layers and to shallow nonlinear networks with piecewise linear activation functions. In linear networks, rapid feature learning only occurs with balanced initializations, where all layers learn at similar speeds. While in nonlinear networks, unbalanced initializations that promote faster learning in earlier layers can accelerate rich learning. Through a series of experiments, we provide evidence that this unbalanced rich regime drives feature learning in deep finite-width networks, promotes interpretability of early layers in CNNs, reduces the sample complexity of learning hierarchical data, and decreases the time to grokking in modular arithmetic. Our theory motivates further exploration of unbalanced initializations to enhance efficient feature learning.

1 Introduction

Deep learning has transformed machine learning, demonstrating remarkable capabilities in a myriad of tasks ranging from image recognition to natural language processing. It's widely believed that the impressive performance of these models lies in their capacity to efficiently extract task-relevant features from data. However, understanding this feature acquisition requires unraveling a complex interplay between datasets, network architectures, and optimization algorithms. Within this framework, two distinct regimes, determined at initialization, have emerged: the lazy and the rich.

Lazy regime. Various investigations have revealed a notable phenomenon in overparameterized neural networks, where throughout training the networks remain close to their linearization [1, 2, 3, 4, 5]. Seminal work by Jacot et al. [6], demonstrated that in the infinite-width limit, the Neural Tangent Kernel (NTK), which describes the evolution of the neural network through training, converges to a deterministic limit. Consequently, the network learns a solution akin to kernel regression with the NTK matrix. Termed the *lazy* or *kernel* regime, this domain has been characterized by a deterministic NTK [6, 7], convex dynamics with minimal movement in parameter space [8], static hidden representations, exponential learning curves, and implicit biases aligned with a reproducing kernel Hilbert space (RKHS) norm [9]. However, Chizat et al. [8] challenged this understanding,

---

* Equal contribution. Correspondence to `kunin@stanford.edu` and `aravento@stanford.edu`.

(a) Scale and geometry influence the learning regime   (b) A complex phase portrait of feature learning
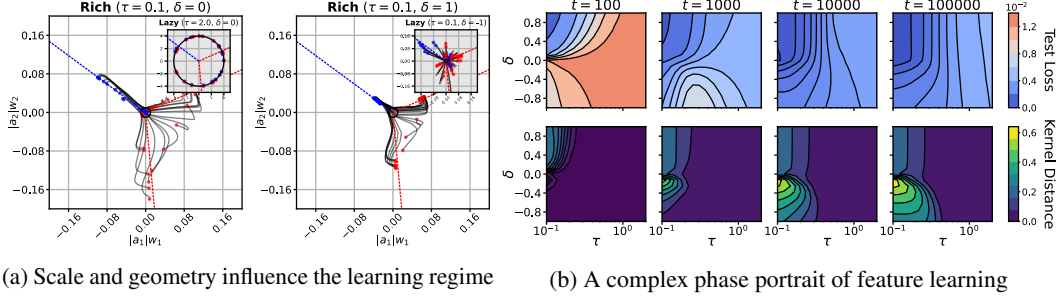
Figure 1: **Unbalanced initializations lead to rapid rich learning and generalization.** We follow the experimental setup used in Fig. 1 of Chizat et al. [8] – a wide two-layer student ReLU network $f(x; \theta) = \sum_{i=1}^{h} a_i \max\{0, w_i^\mathsf{T} x\}$ trained on a dataset generated from a narrow two-layer teacher ReLU network. The student parameters are initialized as $w_i \sim \text{Unif}(\mathbb{S}^{d-1}(\frac{\tau}{\alpha}))$ and $a_i = \pm\alpha\tau$ such that $\tau > 0$ controls the *scale* of the function at initialization, while $\alpha > 0$ controls the *geometry*, the relative magnitude of the first and second layer parameters. The conserved quantity is $\delta = \tau^2(\alpha^2 - \alpha^{-2})$. (a) Shows the training trajectories of $|a_i|w_i$ (color denotes $\text{sgn}(a_i)$) when $d = 2$ for four different settings of τ, δ. The left plot confirms that small scale leads to rich and large scale to lazy. The right plot shows that even at small scale, the initialization geometry can move the network between rich and lazy as well. Here an upstream initialization $\delta > 0$ shows striking alignment to the teacher (dotted lines), while a downstream initialization $\delta < 0$ shows no alignment. (b) Shows the test loss and kernel distance from initialization computed through training over a sweep of τ and δ when $d = 100$. Lazy learning happens when τ is large, rich learning happens when τ is small, and rapid rich learning happens when *both* τ is small and δ is large – an upstream initialization. This initialization also leads to the smallest test loss. See Fig. 10 in Appendix D.1 for supporting figures.

asserting that the lazy regime isn't a product of the infinite-width architecture, but is contingent on the *scale* of the network at initialization. They demonstrated that given any finite-width model $f(x; \theta)$ whose output is zero at initialization, a scaled version of the model $\tau f(x; \theta)$ will enter the lazy regime as the scale τ diverges. However, they also noted that these scaled models often perform worse in test error. While the lazy regime offers insights into the network's convergence to a global minimum, it does not fully capture the generalization capabilities of neural networks trained with standard initializations. It is thus widely believed that a different regime, driven by small or vanishing initializations, underlies the many successes of neural networks.

Rich regime. In contrast to the lazy regime, the *rich* or *feature-learning* or *active* regime is distinguished by a learned NTK that evolves through training, non-convex dynamics traversing between saddle points [10, 11, 12], sigmoidal learning curves, and simplicity biases such as low-rankness [13] or sparsity [14]. Yet, the exact characterization of rich learning and the features it learns frequently depends on the specific problem at hand, with its definition commonly simplified as what it is not: lazy. Recent analyses have shown that beyond scale, other aspects of the initialization can substantially impact the extent of feature learning, such as the effective rank [15], layer-specific initialization variances [16, 17, 18], and large learning rates [19, 20, 21, 22]. Azulay et al. [9] demonstrated that in two-layer linear networks, the relative difference in weight magnitudes between the first and second layer, termed the *initialization geometry* in our work, can impact feature learning, with balanced initializations yielding rich learning dynamics, while unbalanced ones tend to induce lazy dynamics. However, as shown in Fig. 1, for nonlinear networks unbalanced initializations can induce both rich and lazy dynamics, creating a complex phase portrait of learning regimes influenced by both scale and geometry. Building on these observations, our study aims to precisely understand how layer-specific initialization variances and learning rates determine the transition between lazy and rich learning in finite-width networks. Moreover, we endeavor to gain insights into the inductive biases of both regimes, and the transition between them, during training and at interpolation, with the ultimate goal of elucidating how the rich regime acquires features that facilitate generalization.

Our contributions. Our work begins with an exploration of the two-layer single-neuron linear network proposed by Azulay et al. [9] as a minimal model displaying both lazy and rich learning. By employing a combination of hyperbolic and spherical coordinate transformations, we derive exact solutions for the gradient flow dynamics with layer-specific learning rates under all initializations.

Alongside recent work by Xu and Ziyin [23][1], our analysis stands out as one of the few analytically tractable models for the transition between lazy and rich learning in a finite-width network, marking a notable contribution to the field. Our analysis reveals that the layer-specific initialization variances and learning rates, which we collectively refer to as the *initialization geometry*, conspire to influence the learning regime through a simple set of conserved quantities that constrain the geometry of learning trajectories. Additionally, it reveals that a crucial aspect of the initialization geometry overlooked in prior analysis is the directionality. While a *balanced initialization* results in all layers learning at similar rates, an *unbalanced initialization* can cause faster learning in either earlier layers, referred to as an *upstream initialization*, or later layers, referred to as a *downstream initialization*. Due to the depth-dependent expressivity of layers in a network, upstream and downstream initializations often exhibit fundamentally distinct learning trajectories. We extend our analysis of the initialization geometry developed in the single-neuron model to more complex linear models with multiple neurons, outputs, and layers and to two-layer nonlinear networks with piecewise linear activation functions. We find that in linear networks, rapid rich learning can only occur with balanced initializations, while in nonlinear networks, upstream initializations can actually accelerate rich learning. Finally, through a series of experiments, we provide evidence that upstream initializations drive feature learning in deep finite-width networks, promote interpretability of early layers in CNNs, reduce the sample complexity of learning hierarchical data, and decrease the time to grokking in modular arithmetic.

Notation. In this work, we consider a feedforward network $f(x; \theta) : \mathbb{R}^d \to \mathbb{R}$ parameterized by $\theta \in \mathbb{R}^p$. The network is trained by gradient flow $\dot{\theta}_i = -\eta_{\theta_i} \nabla_{\theta_i} \mathcal{L}(\theta)$, with an initialization θ_0 and layer-specific learning rate $\eta_{\theta_i} > 0 \in \mathbb{R}^p$, to minimize the mean squared error $\mathcal{L}(\theta) = \frac{1}{2} \sum_{i=1}^n (f(x_i; \theta) - y_i)^2$ computed over a dataset $\{(x_1, y_1), \ldots, (x_n, y_n)\}$ of size n. We denote the training data matrix as $X \in \mathbb{R}^{n \times d}$ with rows $x_i \in \mathbb{R}^d$ and the label vector as $y \in \mathbb{R}^n$. The network's output $f(x; \theta)$ evolves according to the differential equation, $\partial_t f(x; \theta) = -\sum_{i=1}^n \Theta(x, x_i; \theta)(f(x_i; \theta) - y_i)$, where $\Theta(x, x'; \theta) : \mathbb{R}^d \times \mathbb{R}^d \to \mathbb{R}$ is the *Neural Tangent Kernel (NTK)*, defined as $\Theta(x, x'; \theta) = \sum_{i=1}^p \eta_{\theta_i} \partial_{\theta_i} f(x; \theta) \partial_{\theta_i} f(x'; \theta)$. The NTK quantifies how one gradient step with data point x' affects the evolution of the networks's output evaluated at another data point x. When η_{θ_i} is shared by all parameters, the NTK is the kernel associated with the feature map $\nabla_\theta f(x; \theta) \in \mathbb{R}^p$. We also define the *NTK matrix* $K \in \mathbb{R}^{n \times n}$, which is computed across the training data such that $K_{ij} = \Theta(x_i, x_j; \theta)$. The NTK matrix evolves from its initialization K_0 to convergence K_∞ through training. Lazy and rich learning exist on a spectrum, with the extent of this evolution serving as the distinguishing factor. Various studies have proposed different metrics to track the evolution of the NTK matrix [24, 25, 26, 27]. In this work, we use *kernel distance* [25], defined as $S(t_1, t_2) = 1 - \langle K_{t_1}, K_{t_2} \rangle / (\|K_{t_1}\|_F \|K_{t_2}\|_F)$, which is a scale invariant measure of similarity between the NTK at two points in time. In the lazy regime $S(0, t) \approx 0$, while in the rich regime $0 \ll S(0, t) \leq 1$.

2 Related Work

Linear networks. Significant progress in studying the rich regime has been achieved in the context of linear networks. In this setting, $f(x; \theta) = \beta(\theta)^\intercal x$ is linear in its input x, but can exhibit highly nonlinear dynamics in parameter θ and function $\beta(\theta)$ space. Foundational work by Saxe et al. [10] provided exact solutions to gradient flow dynamics in linear networks with task-aligned initializations. They achieved this by solving a system of Bernoulli differential equations that prioritize learning the most salient features first, which can be beneficial for generalization [28]. This analysis has been extended to wide [29, 30] and deep [31, 32, 33] linear networks with more flexible initialization schemes [34, 35, 36]. It has also been applied to study the evolution of the NTK [37] and the influence of the scale on the transition between lazy and rich learning [12, 23]. In this work, we present novel exact solutions for a minimal model utilizing a mix of Bernoulli and Riccati equations to showcase a complex phase portrait of lazy and rich learning with separate alignment and fitting phases.

Implicit bias. An effective analysis approach to understanding the rich regime studies how the initialization influences the inductive bias at interpolation. The aim is to identify a function $Q(\theta)$ such that the network converges to a first-order KKT point minimizing $Q(\theta)$ among all possible interpolating solutions. Foundational work by Soudry et al. [38] pioneered this approach for a linear classifier trained with gradient descent, revealing a max margin bias. These findings have been extended to deep linear networks [39, 40, 41], homogeneous networks [42, 43, 44], and quasi-homogeneous networks [45]. A similar line of research expresses the learning dynamics of networks

---

[1]Xu and Ziyin [23] presented exact NTK dynamics for a linear model trained with one-dimensional data.

trained with mean squared error as a *mirror flow* for some potential $\Phi(\beta)$, such that the inductive bias can be expressed as a *Bregman divergence* [46]. This approach has been applied to diagonal linear networks, revealing an inductive bias that interpolates between ℓ^1 and ℓ^2 norms in the rich and lazy regimes respectively [14]. However, finding the potential $\Phi(\beta)$ is problem-specific and requires solving a second-order differential equation, which may not be solvable even in simple settings [47, 48]. Azulay et al. [9] extended this analysis to a time-warped mirror flow, enabling the study of a broader class of architectures. In this work we derive exact expressions for the inductive bias of our minimal model and extend the results in Azulay et al. [9] to wide and deep linear networks.

Two-layer networks. Two-layer, or single-hidden layer, piecewise linear networks have emerged as a key setting for advancing our understanding of the rich regime. Maennel et al. [49] observed that in training two-layer ReLU networks from small initializations, the first-layer weights concentrate along fixed directions determined by the training data, irrespective of network width. This phenomenon, termed *quantization*, has been proposed as a *simplicity bias* inherent to the rich regime, driving the network towards low-rank solutions when feasible. Subsequent studies have aimed to precisely elucidate this effect by introducing structural constraints on the training data [50, 51, 52, 53, 54, 55]. Across these analyses, a consistent observation is that the learning dynamics involve distinct phases: an initial alignment phase characterized by quantization, followed by fitting phases where the task is learned. All of these studies assumed a balanced (or nearly balanced) initialization between the first and second layer. In this study, we explore how unbalanced initializations influence the phases of learning, demonstrating that it can eliminate or augment the quantization effect.

Infinite-width networks. Many recent advancements in understanding the rich regime have come from studying how the initialization variance and layer-wise learning rates should scale in the infinite-width limit to ensure constant movement in the activations, gradients, and outputs. In this limit, analyzing dynamics becomes simpler in several respects: random variables concentrate, nonlinearities act linearly, and quantities will either vanish to zero, remain constant, or diverge to infinity [17]. A set of works used tools from statistical mechanics to provide analytic solutions for the rich population dynamics of two-layer nonlinear neural networks initialized according to the *mean field* parameterization [56, 57, 58, 59]. These ideas were extended to deeper networks through a *tensor program* framework, leading to the derivation of *maximal update parametrization* (μP) [16, 18]. The μP parameterization has also been derived through a self-consistent dynamical mean field theory [60] and a spectral scaling analysis [61]. In this study, we focus on finite-width neural networks.

3   A Minimal Model of Lazy and Rich Learning with Exact Solutions

Here we explore an illustrative setting simple enough to admit exact gradient flow dynamics, yet complex enough to showcase lazy and rich learning regimes. We study a two-layer linear network with a single hidden neuron defined by the map $f(x; \theta) = aw^\intercal x$ where $a \in \mathbb{R}$, $w \in \mathbb{R}^d$ are the parameters. We examine how the parameter initializations a_0, w_0 and the layer-wise learning rates η_a, η_w influence the training trajectory in parameter space, function space (defined by the product $\beta = aw$), and the evolution of the the NTK matrix,

$$K = X \left(\eta_w a^2 \mathbf{I}_d + \eta_a w w^\intercal \right) X^\intercal. \quad (1)$$

Except for a measure zero set of initializations which converge to saddle points[2], all gradient flow trajectories will converge to a global minimum, de-

(a) $\delta = -2$ (b) $\delta = 0$ (c) $\delta = 2$

Figure 2: **Balance determines geometry of trajectory.** The quantity $\delta = \eta_w a^2 - \eta_a \|w\|^2$ is conserved through gradient flow, which constrains the trajectory to: (a) a one-sheeted hyperboloid for downstream initializations, (b) a double cone for balanced initializations, and (c) a two-sheeted hyperboloid for upstream initializations. Gradient flow dynamics for three different initializations a_0, w_0 with the same product $\beta_0 = a_0 w_0$ are shown. The minima manifold is shown in red and the manifold of equivalent β_0 initializations in gray. The surface is colored according to training loss, with blue representing higher loss and red representing lower loss.

termined by the normal equations $X^\intercal X aw = X^\intercal y$. However, even when $X^\intercal X$ is invertible such that the global minimum β_* is unique, the rescaling symmetry between a and w results in a manifold of minima in parameter space. The minima manifold is a one-dimensional hyperbola where $w \propto \beta_*$

---

[2]The set of saddle points $\{(a, w)\}$ is the $d - 1$ dimensional subspace satisfying $a = 0$ and $w^\intercal X^\intercal y = 0$.

and has two distinct branches for positive and negative a. The symmetry also imposes a constraint on the network's trajectory, maintaining the difference $\delta = \eta_w a^2 - \eta_a \|w\|^2 \in \mathbb{R}$ throughout training (see Appendix A.1 for details). This confines the parameter dynamics to the surface of a hyperboloid where the magnitude and sign of the conserved quantity determines the geometry, as shown in Fig. 2. An upstream initialization occurs when $\delta > 0$, a balanced initialization when $\delta = 0$, and a downstream initialization when $\delta < 0$.

Deriving exact solutions in parameter space. We assume whitened input $X^\intercal X = \mathbf{I}_d$ such that the ordinary least squares solution is $\beta_* = X^\intercal y$, and the gradient flow dynamics simplify to $\dot{a} = \eta_a \left(w^\intercal \beta_* - a\|w\|^2\right), \dot{w} = \eta_w \left(a\beta_* - a^2 w\right)$. We note that $w(t) \in \text{span}(\{w_0, \beta_*\})$, and through training, w aligns in direction to $\pm \beta_*$ depending on the basin of attraction[3] the parameters are initialized in. Therefore, we can monitor the dynamics by tracking the hyperbolic geometry between a and $\|w(t)\|$ and the spherical angle between $w(t)$ and β_*. We study the variables $\mu = a\|w\|$, an invariant under the rescale symmetry, and $\phi = \frac{w^\intercal \beta_*}{\|w\|\|\beta_*\|}$, the cosine of the spherical angle. From these two scalar quantities $\mu(t), \phi(t)$ and the initialization a_0, w_0, we can determine the trajectory $a(t)$ and $w(t)$ in parameter space. The dynamics for μ, ϕ are given by the coupled nonlinear ODEs,

$$\dot{\mu} = \sqrt{\delta^2 + 4\eta_a\eta_w\mu^2}\left(\phi\|\beta_*\| - \mu\right), \qquad \dot{\phi} = \frac{\eta_a\eta_w 2\mu\|\beta_*\|}{\sqrt{\delta^2 + 4\eta_a\eta_w\mu^2} - \delta}\left(1 - \phi^2\right). \qquad (2)$$

Amazingly, this system can be solved exactly, as discussed in Appendix A.2, and shown in Fig. 3. Without delving into the specifics, we can develop an intuitive understanding of the solutions by examining the influence of the initialization geometry δ.

Upstream. When $\delta \gg 0$, the updates for both μ and ϕ diverge, but ϕ updates much more rapidly. We can decouple the dynamics of μ and ϕ by separation of their time scales and assume ϕ has reached its steady-state of ± 1 before μ has updated. Then, the dynamics of μ is linear and proceeds exponentially to $\pm\|\beta_*\|$. This regime exhibits minimal kernel movement (see Fig. 3 (c)) because the kernel is dominated by the $\eta_w a^2 \mathbf{I}_d$ term, whereas it is mainly w that updates.

Balanced. When $\delta = 0$, μ follows a Bernoulli differential equation driven by a time-dependent signal $\phi\|\beta_*\|$, and ϕ follows a Riccati equation evolving from an initial value to ± 1 depending on the basin of attraction. For vanishing initialization $\|\beta_0\| \to 0$, the temporal dynamics of μ and ϕ decouple such that there are two phases of

(a) $\mu(t)$ (b) $\phi(t)$ (c) $S(0, t)$

Figure 3: **Exact solutions for the single hidden neuron model.** Our theoretical predictions (black dashed lines) agree with gradient flow simulations (solid lines, color-coded based on δ values), shown here for three key metrics: μ (left), ϕ (middle), and $S(0, t)$ (right). Each metric starts at the same value for all δ, but varying δ has a pronounced effect on the metric's dynamics. For upstream initializations ($\delta \gg 0$), μ changes only slightly, ϕ exponentially aligns, and S remains near zero, indicative of the lazy regime. For balanced initializations ($\delta = 0$), both μ and ϕ change significantly and S quickly moves away from zero, indicative of the rich regime. For downstream initializations ($\delta \ll 0$), μ quickly drops to zero, then μ and ϕ slowly climb back to one. Similarly, S remains small before a sudden transition towards one, indicative of a delayed rich regime. See Appendix A.2 for further details.

learning: an initial alignment phase where $\phi \to \pm 1$, followed by a fitting phase where $\mu \to \pm\|\beta_*\|$. In the first phase, w aligns to β_* resulting in a rank-one update to the NTK, identical to the silent alignment effect described in Atanasov et al. [37]. In the second phase, the dynamics of μ simplify to the Bernoulli equation studied in Saxe et al. [10] and the kernel evolves solely in overall scale.

Downstream. When $\delta \ll 0$, the updates for μ diverge, while the updates for ϕ vanishes. In this regime the dynamics proceed by an initial fast phase where μ converges exponentially to its steady state of $\phi\|\beta_*\|$. Plugging this steady state into the dynamics of ϕ gives a Bernoulli differential equation $\dot{\phi} = \eta_a\eta_w\|\beta_*\|^2|\delta|^{-1}\phi(1 - \phi^2)$. Due to the coefficient $|\delta|^{-1}$, the second alignment phase proceeds very slowly as ϕ approaches ± 1, assuming $\phi, \mu \neq 0$, which is a saddle point. In this regime, the

---

[3]The basin is given by the sign of a_0 for $\delta \geq 0$ or $w_0^\intercal \beta_* + \frac{a_0}{2}(\delta + \sqrt{\delta^2 + 4\|\beta_*\|^2})$ for $\delta < 0$. See A.2.5.

dynamics proceed by an initial lazy fitting phase, followed by a rich alignment phase, where the delay is determined by the magnitude of δ.

Identifying regimes of learning in function space. Here we take an alternative route towards understanding the influence of the initialization geometry by directly examining the dynamics in function space, an analysis strategy we will generalize to broader setups in Sections 4 and 5. The network's function is determined by the product $\beta = aw$ and governed by the ODE,

$$\dot{\beta} = -\underbrace{\left(\eta_w a^2 \mathbf{I}_d + \eta_a ww^\intercal\right)}_{M} X^\intercal \rho, \quad (3)$$

where $\rho = X\beta - y$ is the residual. These dynamics can be interpreted as preconditioned gradient flow on the loss in function space where the preconditioning matrix M depends on time through its dependence on a^2 and $ww^\intercal$. Whenever $\|\beta\| \neq 0$, we can express M in terms of β and δ as

$$M = \frac{\kappa + \delta}{2} I_d + \frac{\kappa - \delta}{2} \frac{\beta\beta^\intercal}{\|\beta\|^2}, \quad (4)$$

where $\kappa = \sqrt{\delta^2 + 4\eta_a\eta_w\|\beta\|^2}$ (see Appendix A.3 for a derivation). This establishes a *self-consistent* equation for the dynamics of β regulated by δ. Additionally,

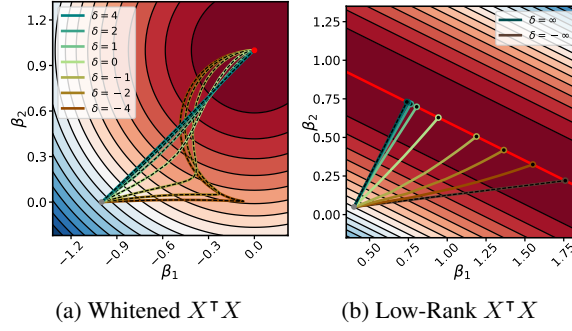(a) Whitened $X^\intercal X$ (b) Low-Rank $X^\intercal X$

Figure 4: **Balance modulates β dynamics and implicit bias.** Here we show the dynamics of $\beta = aw$ with different values of δ, but the same initial β_0. When $X^\intercal X$ is whitened (left), we can solve for the dynamics exactly using our expressions for μ, ϕ (black dashed lines). Upstream initializations follow the trajectory of gradient flow on β, downstream initializations first move in the direction of β_0 before sweeping around towards β_*, and balanced initializations take an intermediate trajectory between these two. When $X^\intercal X$ is low-rank (right), then we can only predict the trajectories in the limit of $\delta = \pm\infty$. If the interpolating manifold is one-dimensional, then we can solve for the solution in terms of δ exactly (black dots). See Appendix A.4 for details.

notice that M characterizes the NTK matrix Eq. (1). Thus, understanding the evolution of M along the trajectory β_0 to β_* offers a method to discern between lazy and rich learning. ***Upstream.*** When $\delta \gg 0$, $M \approx \delta\mathbf{I}_d$, and the dynamics of β converge to the trajectory of linear regression trained by gradient flow. Along this trajectory the NTK matrix remains constant, confirming the dynamics are lazy. ***Balanced.*** When $\delta = 0$, $M = \sqrt{\eta_a\eta_w}\|\beta\|(\mathbf{I}_d + \frac{\beta\beta^\intercal}{\|\beta\|^2})$. Here the dynamics balance between following the lazy trajectory and attempting to fit the task by only changing in norm. As a result the NTK changes in both magnitude and direction through training, confirming the dynamics are rich. ***Downstream.*** When $\delta \ll 0$, $M \approx |\delta|\frac{\beta\beta^\intercal}{\|\beta\|^2}$, and β follows a projected gradient descent trajectory, attempting to reach β_* in the direction of β_0. Along this trajectory the NTK matrix doesn't evolve. However, if β_0 is not aligned to β_*, then at some point the dynamics of β will slowly align. In this second alignment phase the NTK matrix will change, confirming the dynamics are initially lazy followed by a delayed rich phase.

Determining the implicit bias via mirror flow. So far we have considered whitened or full rank $X^\intercal X$, ensuring the existence of a unique least squares solution β_*. In this setting, δ influences the trajectory the model takes from β_0 to β_*, as shown in Fig. 4 (a). Now we consider low-rank $X^\intercal X$, such that there exist infinitely many interpolating solutions in function space. By studying the structure of M, we can characterize how δ determines the interpolating solution the dynamics converge to. Extending a time-warped mirror flow analysis strategy pioneered by Azulay et al. [9] to allow $\delta < 0$ (see Appendix A.4 for details), we prove the following theorem, which shows a tradeoff between reaching the minimum norm solution and preserving the direction of the initialization β_0.

Theorem 3.1 (Extending Theorem 2 in Azulay et al. [9])**.** *For a single hidden neuron linear network, for any $\delta \in \mathbb{R}$, and initialization β_0 such that $\beta(t) \neq 0$ for all $t \geq 0$, if the gradient flow solution $\beta(\infty)$ satisfies $X\beta(\infty) = y$, then,*

$$\beta(\infty) = \underset{\beta \in \mathbb{R}^d}{\arg\min}\ \Psi_\delta(\beta) - \psi_\delta \frac{\beta_0}{\|\beta_0\|}^\intercal \beta \quad s.t. \quad X\beta = y \quad (5)$$

where $\Psi_\delta(\beta) = \frac{1}{3}\left(\sqrt{\delta^2 + 4\|\beta\|^2} - 2\delta\right)\sqrt{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}$ *and* $\psi_\delta = \sqrt{\sqrt{\delta^2 + 4\|\beta_0\|^2} - \delta}$.

6

We observe that for vanishing initializations there is functionally no difference between the inductive bias of the upstream ($\delta \gg 0$) and balanced ($\delta = 0$) settings. However, in the downstream setting ($\delta \ll 0$), it is the second term preserving the direction of the initialization that dominates the inductive bias. This tradeoff in inductive bias as a function of δ is presented in Fig. 4 (b), where if the null space of $X^\mathsf{T}X$ is one-dimensional, we can solve for $\beta(\infty)$ in closed form (see Appendix A.4).

4 Wide and Deep Linear Networks

We now show how the analysis techniques used to study the influence of initialization geometry in the single-neuron setting can be applied to linear networks with multiple neurons, outputs, and layers.

Wide linear networks. We consider the dynamics of a two-layer linear network with h hidden neurons and c outputs, $f(x; \theta) = A^\mathsf{T}Wx$, where $W \in \mathbb{R}^{h \times d}$ and $A \in \mathbb{R}^{h \times c}$. We assume that $\min\{d, c\} \leq h \leq \max\{d, c\}$, such that this parameterization can represent all linear maps from $\mathbb{R}^d \to \mathbb{R}^c$. As in the single-neuron setting, the rescaling symmetry in this model implies that $\Delta = \eta_w A_0 A_0^\mathsf{T} - \eta_a W_0 W_0^\mathsf{T}$ remains conserved throughout gradient flow [62]. The NTK matrix is $K = (\mathbf{I}_c \otimes X)(\eta_w A^\mathsf{T}A \oplus \eta_a W^\mathsf{T}W)(\mathbf{I}_c \otimes X^\mathsf{T})$, where $\otimes$ and $\oplus$ denote the Kronecker product and sum[4] respectively. Drawing insights from our analysis of the single-neuron scenario ($h = c = 1$), we consider the dynamics of $\beta = W^\mathsf{T}A \in \mathbb{R}^{d \times c}$,

$$\text{vec}\left(\dot{\beta}\right) = - \underbrace{(\eta_w A^\mathsf{T}A \oplus \eta_a W^\mathsf{T}W)}_{M} \text{vec}(X^\mathsf{T}X\beta - X^\mathsf{T}Y), \tag{6}$$

where $\text{vec}(\cdot)$ denotes the vectorization operator. As in the single-neuron setting, we find that the dynamics of β are preconditioned by a matrix M that depends on quadratics of A and W. M can be expressed in terms of the matrices $\beta_i = w_i a_i^\mathsf{T} \in \mathbb{R}^{d \times c}$, which represent the contribution to the input-output map of a single hidden neuron with parameters w_i, a_i and conserved quantity $\delta_i = \Delta_{ii}$.

Theorem 4.1. *Assume* $\|\beta_i\|_F \neq 0$ *for all* $i \in [h]$ *and let* $\kappa_i = \sqrt{\delta_i^2 + 4\eta_a\eta_w\|\beta_i\|_F^2}$, *then the matrix* M *can be expressed as the sum* $M = \sum_{i=1}^h M_i$ *over hidden neurons where* M_i *is defined as,*

$$M_i = \left(\frac{\kappa_i + \delta_i}{2}\right)\frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2} \oplus \left(\frac{\kappa_i - \delta_i}{2}\right)\frac{\beta_i\beta_i^\mathsf{T}}{\|\beta_i\|_F^2}. \tag{7}$$

By studying the dependence of M on the conserved quantities δ_i and the relative sizes of dimensions d, h and c, we can determine the influence of the initialization geometry on the learning regime[5].

Funnel and inverted-funnel networks. We consider *funnel networks*, which narrow from input to output ($d > h \geq c$), and *inverted-funnel networks*, which expand from input to output ($d \leq h < c$). Except for a measure zero set of initializations, funnel networks always enter the lazy regime as all $\delta_i \to \infty$, whereas inverted funnel networks do so as $\delta_i \to -\infty$. As elaborated in Appendix B, this occurs because in these limits a solution β_* must exist within the space spanned by M at initialization. In the opposite limits of δ, these networks enter a lazy followed by rich regime, assuming $\beta_0 \not\propto \beta_*$. Conversely, as the $\delta_i \to 0$, all networks transition into the rich regime.

Square networks. In the setting of *square networks* ($d = h = c$), we can precisely identify the influence δ_i has on the β trajectory. First, we establish that as all δ_i tend towards $\pm\infty$, the network symmetrically transitions into the lazy regime. Second, by leveraging the task-aligned initialization proposed in Saxe et al. [10], we can directly express M in terms of β, and show that its dynamics are given by mirror flow with a *hyperbolic entropy* potential evaluated on the singular values of β (see Appendix B and Theorem B.6 for details). The potential smoothly interpolates between an ℓ^1 and ℓ^2 penalty on the singular values for the rich and lazy regimes respectively. This potential was first identified as the inductive bias for diagonal linear networks by Woodworth et al. [14].

Deep linear networks. As detailed in Appendix B.2, we generalize the inductive bias derived for rich two-layer linear networks by Azulay et al. [9] to encompass deep linear networks. For a depth-$(l+1)$ linear network, $f(x; \theta) = A^\mathsf{T}\prod_{i=1}^l W_i x$, where $\beta = \prod_{i=1}^l W_i^\mathsf{T}A$, we find that the inductive bias of the rich regime is $Q(\beta) = (\frac{l+1}{l+2})\|\beta\|^{\frac{l+2}{l+1}} - \|\beta_0\|^{-\frac{l}{l+1}}\beta_0^\mathsf{T}\beta$ (see Theorem B.9). This inductive bias strikes a balance between attaining the minimum norm solution and preserving the initialization direction, which with increased depth emphasizes the latter.

---

[4]The Kronecker sum is defined for square matrices $A \in \mathbb{R}^{c \times c}$ and $B \in \mathbb{R}^{d \times d}$ as $A \oplus B = A \otimes \mathbf{I}_d + \mathbf{I}_c \otimes B$.

[5]As shown in Appendix B, we can recover Eq. (4) presented in the single-neuron setting directly from Eq. (7).

5 Piecewise Linear Networks

We now take a first step towards extending our analysis from linear networks to piecewise linear networks with activation functions of the form $\sigma(z) = \max\{z, \gamma z\}$. The input-output map of a piecewise linear network with l hidden layers and h hidden neurons per layer is comprised of potentially $O(h^{dl})$ convex *activation regions* [63]. Each region is defined by a unique *activation pattern* of the hidden neurons. The input-output map is linear within each region and continuous at the boundary between regions. Collectively, the activation regions form a 2-colorable[6] convex partition of input space, as shown in Fig. 5. We investigate how the initialization geometry influences the evolution of this partition and the linear maps within each region.

Two-layer network. We consider the dynamics of a two-layer piecewise linear network without biases, $f(x; \theta) = a^\mathsf{T}\sigma(Wx)$, where $W \in \mathbb{R}^{h \times d}$ and $a \in \mathbb{R}^h$. Following the approach in Section 4, we consider the contribution to the input-output map from a single hidden neuron $i \in [h]$ with parameters $w_i \in \mathbb{R}^d$, $a_i \in \mathbb{R}$ and conserved quantity $\delta_i = \eta_w a_i^2 - \eta_a \|w_i\|^2$ [62]. However, unlike the linear setting, the neuron's contribution to $f(x_j; \theta)$ is regulated by whether the input x_j is in the neuron's *active halfspace*. Let $C \in \mathbb{R}^{h \times n}$ be the matrix with elements $c_{ij} = \sigma'(w_i^\mathsf{T} x_j)$, which determines the activation of the i^{th} neuron for the j^{th} data point, and let $\rho_j = f(x_j; \theta) - y_j$. Then we can express the dynamics of $\beta_i = a_i w_i$ as,

$$\dot{\beta}_i = -\underbrace{\left(\eta_w a_i^2 I_d + \eta_a w_i w_i^\mathsf{T}\right)}_{M_i}\underbrace{\sum_{j=1}^n c_{ij} x_j \rho_j}_{\xi_i}. \quad (8)$$

The matrix $M_i \in \mathbb{R}^{d \times d}$ is a preconditioning matrix on the dynamics, and when $\beta_i \neq 0$, it can be expressed in terms of β_i and δ_i. Unlike the linear setting, $\xi_i \in \mathbb{R}^d$ driving the dynamics is not shared for all neurons because of its dependence on c_{ij}. Additionally, the NTK matrix in this setting


Figure 5: **Evolution of a ReLU network.** Here we depict the surface of a two-layer ReLU network with twenty hidden neurons, starting from a nearly zero input-output map $f(x; \theta_0) \approx 0$, trained on a task akin to *exclusive or*. The surface consists of convex conic regions, each with a distinct activation pattern, colored by the parity of active neurons. With a lazy downstream initialization (bottom), the activation partition is fixed, and the network fits the data by reweighing the hidden neurons. With a rich balanced or upstream initialization (top), the network undergoes an initial alignment phase where the partition map changes rapidly while the input-output map remains close to zero, followed by a data fitting phase.

depends on M_i and C, with elements $K_{jk} = \sum_{i=1}^h c_{ij} x_j^\mathsf{T} M_i x_k c_{ik}$. To examine the evolution of K, we consider a *signed spherical coordinate* transformation separating the dynamics of β_i into its directional $\hat{\beta}_i = \text{sgn}(a_i)\frac{\beta_i}{\|\beta_i\|}$ and radial $\mu_i = \text{sgn}(a_i)\|\beta_i\|$ components, such that $\beta_i = \mu_i \hat{\beta}_i$. $\hat{\beta}_i$ determines the direction and orientation of the halfspace where the i^{th} neuron is active, while μ_i determines the slope of the contribution in this halfspace. These coordinates evolve according to,

$$\dot{\mu}_i = -\sqrt{\delta_i^2 + 4\eta_a\eta_w\mu_i^2}\,\hat{\beta}_i^\mathsf{T}\xi_i, \qquad \dot{\hat{\beta}}_i = -\frac{\sqrt{\delta_i^2 + 4\eta_a\eta_w\mu_i^2} + \delta_i}{2\mu_i}\left(I_d - \hat{\beta}_i\hat{\beta}_i^\mathsf{T}\right)\xi_i. \qquad (9)$$

Downstream. When $\delta_i \ll 0$, $M_i \approx |\delta_i|\hat{\beta}_i\hat{\beta}_i^\mathsf{T}$. The dynamics are approximately $\partial_t\hat{\beta}_i = 0$ and $\partial_t\mu_i = -|\delta_i|\hat{\beta}_i^\mathsf{T}\xi_i$. Irrespective of ξ_i, $\hat{\beta}_i(t) = \hat{\beta}_i(0)$, which implies the overall partition map doesn't change (Fig. 5, bottom), nor the activation patterns C, nor M_i. Only μ_i changes to fit the data, while the NTK remains constant. If the number of hidden neurons is insufficient to fit the data, there will be a second, slow rich alignment phase where the kernel will change, with $|\delta_i|$ controlling the delay.

Balanced. When $\delta_i = 0$, $M_i = \sqrt{\eta_a\eta_w}|\mu_i|(\mathbf{I}_d + \hat{\beta}_i\hat{\beta}_i^\mathsf{T})$, and the dynamics simplify to, $\partial_t\hat{\beta}_i = -\sqrt{\eta_a\eta_w}\text{sgn}(\mu_i)(\mathbf{I}_d - \hat{\beta}_i\hat{\beta}_i^\mathsf{T})\xi_i$ and $\partial_t\mu_i = -2\sqrt{\eta_a\eta_w}|\mu_i|\hat{\beta}_i^\mathsf{T}\xi_i$. Here both the direction and magnitude of β_i evolve, resulting in changes to the activation regions, patterns C, and NTK K. For vanishing initializations where $\|\beta_i(0)\| \to 0$ for all $i \in [h]$, we can decouple the dynamics into two distinct phases of training, analogous to the rich regime discussed in Section 3 and shown in Fig. 5, top. *Phase I: Partition alignment.* At vanishing scale, the output $f(x; \theta_0) \approx 0$ for all input x,

---

[6]To our knowledge, this property has not been recognized before. See Appendix C.1 for a formal statement.

(a) Kernel Distance   (b) Convolutional Filters   (c) Random Hierarchy   (d) Transformer Grokking
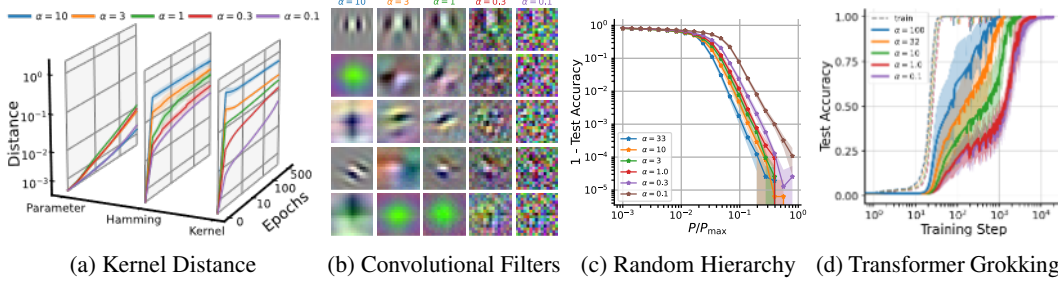
Figure 6: **Impact of upstream initializations in practice.** Here we provide evidence that an upstream initialization (a) drives feature learning through changing activation patterns, (b) promotes interpretability of early layers in CNNs, (c) reduces the sample complexity of learning hierarchical data, and (d) decreases the time to grokking in modular arithmetic. In these experiments, we regulate the learning speed of the first layer relative to the rest of the network by scaling the initialization without altering the input-output map. Here, $\alpha = 1$ denotes a standard parameterization. $\alpha \gg 1$ and $\alpha \ll 1$ correspond to upstream and downstream initializations, respectively. See Appendix D.3.

such that the vector driving the dynamics $\xi_i \approx -\sum_{j=1}^{n} c_{ij} x_j y_j$ is independent of the other hidden neurons. At the same time, the radial dynamics slow down relative to the directional dynamics, and the function's output will remain small as each neuron aligns decoupled from the rest. Prior works have introduced structural constraints on the training data, such as orthogonally separable [50, 53, 54], pair-wise orthonormal [52], linearly separable and symmetric [51] or small angle [55], to analytically determine the fixed points of this alignment phase. *Phase II: Data fitting.* After enough time, the magnitudes of β_i have grown such that we can no longer assume $f(x; \theta) \approx 0$ and thus the residual will depend on all β_i. In this phase, the radial dynamics dominate the learning driving the network to fit the data. However, it is possible for the directions to continue to change, and thus prior works have further decomposed this phase into multiple stages.

Upstream. When $\delta_i \gg 0$, $M_i \approx \delta_i \mathbf{I}_d$, and the dynamics are approximately $\partial_t \hat{\beta}_i = -\delta_i \mu_i^{-1}(\mathbf{I}_d - \hat{\beta}_i \hat{\beta}_i^{\mathsf{T}}) \xi_i$ and $\partial_t \mu_i = -\delta_i \hat{\beta}_i^{\mathsf{T}} \xi_i$. Again, both the direction and magnitude of β_i change. However, unlike the balanced setting, in this setting M_i is independent of β_i and stays constant through training. Yet, as β_i change in direction, so can C, and thus the NTK. This setting is unique, because it is rich due to a changing activation pattern, but the dynamics do not move far in parameter space. Furthermore, unlike in the balanced scenario where scale adjusts the speed of radial dynamics, here it regulates the speed of directional dynamics, with vanishing initializations prompting an extremely fast alignment phase, as observed in Fig. 1.

Unbalanced initializations in diverse domains. In our analysis, we find that upstream initializations can lead to rapid rich learning in nonlinear networks, explaining results shown in Fig. 1. Further experiments in Fig. 6 suggest that upstream initializations has an impact across various domains of deep learning: (a) Standard initializations see significant NTK evolution early in training [25]. We show the movement is linked to changes in activation patterns (Hamming distance) rather than large parameter shifts. Adjusting the initialization variance of the first and last layers can amplify or diminish this movement. (b) Filters in CNNs trained on image classification tasks often align with edge detectors [64]. We show that adjusting the learning speed of the first layer can enhance or degrade this alignment. (c) Deep learning models are believed to avoid the curse of dimensionality and learn with limited data by exploiting hierarchical structures in real-world tasks. Using the Random Hierarchy Model, introduced by Petrini et al. [65] as a framework for synthetic hierarchical tasks, we show that modifying the initialization geometry can decrease or increase the sample complexity of learning. (d) Networks trained on simple modular arithmetic tasks will suddenly generalize long after memorizing their training data [66]. This behavior, termed grokking, is thought to result from a transition from lazy to rich learning [67, 68] and believed to be important towards understanding emergent phenomena [69]. We show that decreasing the variance of the embedding in a single-layer transformer ($< 6\%$ of all parameters) significantly reduces the time to grokking. Overall, our experiments suggest that upstream initializations may play a crucial role in neural network behaviors.

9

6 Conclusion

In this work, we derived exact solutions to a minimal model that can transition between lazy and rich learning to precisely elucidate how unbalanced layer-specific initialization variances and learning rates determine the degree of feature learning. We further extended our analysis to wide and deep linear networks and shallow piecewise linear networks. We find through theory and empirics that unbalanced initializations, which promote faster learning at earlier layers, can actually accelerate rich learning. **Limitations.** The primary limitation lies in the difficulty to extend our theory to deeper nonlinear networks. In contrast to linear networks, where additional symmetries simplify dynamics, nonlinear networks require consideration of the activation pattern's impact on subsequent layers. One potential solution involves leveraging the path framework used in Saxe et al. [70]. Another limitation is our omission of discretization and stochastic effects of SGD, which disrupt the conservation laws central to our study and introduce additional simplicity biases [71, 72, 73]. **Future work.** Our theory encourages further investigation into unbalanced initializations to optimize efficient feature learning. In deep networks there are many ways to be unbalanced beyond upstream and downstream. Understanding how the *learning speed profile* across layers impacts feature learning, inductive biases, and generalization is an important direction for future work.

Acknowledgments and Disclosure of Funding

Author Contributions

This project originated from conversations between Daniel and Allan at the Kavli Institute for Theoretical Physics. Daniel, Allan, and Feng are primarily responsible for the single neuron analysis in Section 3. Clem and Daniel are primarily responsible for the wide and deep linear analysis in Section 4. Daniel is primarily responsible for the nonlinear analysis in Section 5. Allan, Feng, and David are primarily responsible for the empirics in Fig. 1 and Fig. 6. Daniel is primarily responsible for writing the main sections. All authors contributed to the writing of the appendix and the polishing of the manuscript.

References

[1] Simon S Du, Xiyu Zhai, Barnabas Poczos, and Aarti Singh. Gradient descent provably optimizes over-parameterized neural networks. *arXiv preprint arXiv:1810.02054*, 2018.

[2] Simon Du, Jason Lee, Haochuan Li, Liwei Wang, and Xiyu Zhai. Gradient descent finds global minima of deep neural networks. In *International conference on machine learning*, pages 1675–1685. PMLR, 2019.

[3] Zeyuan Allen-Zhu, Yuanzhi Li, and Yingyu Liang. Learning and generalization in overparameterized neural networks, going beyond two layers. *Advances in neural information processing systems*, 32, 2019.

[4] Zeyuan Allen-Zhu, Yuanzhi Li, and Zhao Song. A convergence theory for deep learning via over-parameterization. In *International conference on machine learning*, pages 242–252. PMLR, 2019.

[5] Difan Zou, Yuan Cao, Dongruo Zhou, and Quanquan Gu. Gradient descent optimizes over-parameterized deep relu networks. *Machine learning*, 109:467–492, 2020.

[6] Arthur Jacot, Franck Gabriel, and Clément Hongler. Neural tangent kernel: Convergence and generalization in neural networks. *Advances in neural information processing systems*, 31, 2018.

[7] Greg Yang. Tensor programs ii: Neural tangent kernel for any architecture. *arXiv preprint arXiv:2006.14548*, 2020.

[8] Lenaic Chizat, Edouard Oyallon, and Francis Bach. On lazy training in differentiable programming. *Advances in neural information processing systems*, 32, 2019.

[9] Shahar Azulay, Edward Moroshko, Mor Shpigel Nacson, Blake E Woodworth, Nathan Srebro, Amir Globerson, and Daniel Soudry. On the implicit bias of initialization shape: Beyond infinitesimal mirror descent. In *International Conference on Machine Learning*, pages 468–477. PMLR, 2021.

[10] Andrew M Saxe, James L McClelland, and Surya Ganguli. Exact solutions to the nonlinear dynamics of learning in deep linear neural networks. *arXiv preprint arXiv:1312.6120*, 2013.

[11] Andrew M Saxe, James L McClelland, and Surya Ganguli. A mathematical theory of semantic development in deep neural networks. *Proceedings of the National Academy of Sciences*, 116 (23):11537–11546, 2019.

[12] Arthur Jacot, François Ged, Berfin Şimşek, Clément Hongler, and Franck Gabriel. Saddle-to-saddle dynamics in deep linear networks: Small initialization training, symmetry, and sparsity. *arXiv preprint arXiv:2106.15933*, 2021.

[13] Zhiyuan Li, Yuping Luo, and Kaifeng Lyu. Towards resolving the implicit bias of gradient descent for matrix factorization: Greedy low-rank learning. *arXiv preprint arXiv:2012.09839*, 2020.

[14] Blake Woodworth, Suriya Gunasekar, Jason D Lee, Edward Moroshko, Pedro Savarese, Itay Golan, Daniel Soudry, and Nathan Srebro. Kernel and rich regimes in overparametrized models. In *Conference on Learning Theory*, pages 3635–3673. PMLR, 2020.

[15] Yuhan Helena Liu, Aristide Baratin, Jonathan Cornford, Stefan Mihalas, Eric Shea-Brown, and Guillaume Lajoie. How connectivity structure shapes rich and lazy learning in neural circuits. *ArXiv*, 2023.

[16] Greg Yang and Edward J Hu. Feature learning in infinite-width neural networks. *arXiv preprint arXiv:2011.14522*, 2020.

[17] Tao Luo, Zhi-Qin John Xu, Zheng Ma, and Yaoyu Zhang. Phase diagram for two-layer relu neural networks at infinite-width limit. *Journal of Machine Learning Research*, 22(71):1–47, 2021.

[18] Greg Yang, Edward J Hu, Igor Babuschkin, Szymon Sidor, Xiaodong Liu, David Farhi, Nick Ryder, Jakub Pachocki, Weizhu Chen, and Jianfeng Gao. Tensor programs v: Tuning large neural networks via zero-shot hyperparameter transfer. *arXiv preprint arXiv:2203.03466*, 2022.

[19] Aitor Lewkowycz, Yasaman Bahri, Ethan Dyer, Jascha Sohl-Dickstein, and Guy Gur-Ari. The large learning rate phase of deep learning: the catapult mechanism. *arXiv preprint arXiv:2003.02218*, 2020.

[20] Jimmy Ba, Murat A Erdogdu, Taiji Suzuki, Zhichao Wang, Denny Wu, and Greg Yang. High-dimensional asymptotics of feature learning: How one gradient step improves the representation. *Advances in Neural Information Processing Systems*, 35:37932–37946, 2022.

[21] Libin Zhu, Chaoyue Liu, Adityanarayanan Radhakrishnan, and Mikhail Belkin. Catapults in sgd: spikes in the training loss and their impact on generalization through feature learning. *arXiv preprint arXiv:2306.04815*, 2023.

[22] Hugo Cui, Luca Pesce, Yatin Dandi, Florent Krzakala, Yue M Lu, Lenka Zdeborová, and Bruno Loureiro. Asymptotics of feature learning in two-layer networks after one gradient-step. *arXiv preprint arXiv:2402.04980*, 2024.

[23] Yizhou Xu and Liu Ziyin. When does feature learning happen? perspective from an analytically solvable model. *arXiv preprint arXiv:2401.07085*, 2024.

[24] Corinna Cortes, Mehryar Mohri, and Afshin Rostamizadeh. Algorithms for learning kernels based on centered alignment. *The Journal of Machine Learning Research*, 13(1):795–828, 2012.

[25] Stanislav Fort, Gintare Karolina Dziugaite, Mansheej Paul, Sepideh Kharaghani, Daniel M Roy, and Surya Ganguli. Deep learning versus kernel learning: an empirical study of loss landscape geometry and the time evolution of the neural tangent kernel. *Advances in Neural Information Processing Systems*, 33:5850–5861, 2020.

[26] Mario Geiger, Stefano Spigler, Arthur Jacot, and Matthieu Wyart. Disentangling feature and lazy training in deep neural networks. *Journal of Statistical Mechanics: Theory and Experiment*, 2020(11):113301, 2020.

[27] Aristide Baratin, Thomas George, César Laurent, R Devon Hjelm, Guillaume Lajoie, Pascal Vincent, and Simon Lacoste-Julien. Implicit regularization via neural feature alignment. In *International Conference on Artificial Intelligence and Statistics*, pages 2269–2277. PMLR, 2021.

[28] Andrew K Lampinen and Surya Ganguli. An analytic theory of generalization dynamics and transfer learning in deep linear networks. *arXiv preprint arXiv:1809.10374*, 2018.

[29] Kenji Fukumizu. Effect of batch learning in multilayer neural networks. *Gen*, 1(04):1E–03, 1998.

[30] Lukas Braun, Clémentine Carla Juliette Dominé, James E Fitzgerald, and Andrew M Saxe. Exact learning dynamics of deep linear networks with prior knowledge. In *Advances in Neural Information Processing Systems*, 2022.

[31] Sanjeev Arora, Nadav Cohen, and Elad Hazan. On the optimization of deep networks: Implicit acceleration by overparameterization. In *International conference on machine learning*, pages 244–253. PMLR, 2018.

[32] Sanjeev Arora, Nadav Cohen, Wei Hu, and Yuping Luo. Implicit regularization in deep matrix factorization. *Advances in Neural Information Processing Systems*, 32, 2019.

[33] Liu Ziyin, Botao Li, and Xiangming Meng. Exact solutions of a deep linear network. *Advances in Neural Information Processing Systems*, 35:24446–24458, 2022.

[34] Gauthier Gidel, Francis Bach, and Simon Lacoste-Julien. Implicit regularization of discrete gradient dynamics in linear neural networks. *Advances in Neural Information Processing Systems*, 32, 2019.

[35] Salma Tarmoun, Guilherme Franca, Benjamin D Haeffele, and Rene Vidal. Understanding the dynamics of gradient flow in overparameterized linear models. In *International Conference on Machine Learning*, pages 10153–10161. PMLR, 2021.

[36] Daniel Gissin, Shai Shalev-Shwartz, and Amit Daniely. The implicit bias of depth: How incremental learning drives generalization. *arXiv preprint arXiv:1909.12051*, 2019.

[37] Alexander Atanasov, Blake Bordelon, and Cengiz Pehlevan. Neural networks as kernel learners: The silent alignment effect. In *International Conference on Learning Representations*, 2021.

[38] Daniel Soudry, Elad Hoffer, Mor Shpigel Nacson, Suriya Gunasekar, and Nathan Srebro. The implicit bias of gradient descent on separable data. *The Journal of Machine Learning Research*, 19(1):2822–2878, 2018.

[39] Ziwei Ji and Matus Telgarsky. Gradient descent aligns the layers of deep linear networks. *arXiv preprint arXiv:1810.02032*, 2018.

[40] Suriya Gunasekar, Jason D Lee, Daniel Soudry, and Nati Srebro. Implicit bias of gradient descent on linear convolutional networks. *Advances in neural information processing systems*, 31, 2018.

[41] Edward Moroshko, Blake E Woodworth, Suriya Gunasekar, Jason D Lee, Nati Srebro, and Daniel Soudry. Implicit bias in deep linear classification: Initialization scale vs training accuracy. *Advances in neural information processing systems*, 33:22182–22193, 2020.

[42] Kaifeng Lyu and Jian Li. Gradient descent maximizes the margin of homogeneous neural networks. *arXiv preprint arXiv:1906.05890*, 2019.

[43] Mor Shpigel Nacson, Suriya Gunasekar, Jason Lee, Nathan Srebro, and Daniel Soudry. Lexico-graphic and depth-sensitive margins in homogeneous and non-homogeneous deep models. In *International Conference on Machine Learning*, pages 4683–4692. PMLR, 2019.

[44] Lenaic Chizat and Francis Bach. Implicit bias of gradient descent for wide two-layer neural networks trained with the logistic loss. In *Conference on learning theory*, pages 1305–1338. PMLR, 2020.

[45] Daniel Kunin, Atsushi Yamamura, Chao Ma, and Surya Ganguli. The asymmetric maximum margin bias of quasi-homogeneous neural networks. *arXiv preprint arXiv:2210.03820*, 2022.

[46] Suriya Gunasekar, Jason Lee, Daniel Soudry, and Nathan Srebro. Characterizing implicit bias in terms of optimization geometry. In *International Conference on Machine Learning*, pages 1832–1841. PMLR, 2018.

[47] Suriya Gunasekar, Blake Woodworth, and Nathan Srebro. Mirrorless mirror descent: A natural derivation of mirror descent. In *International Conference on Artificial Intelligence and Statistics*, pages 2305–2313. PMLR, 2021.

[48] Zhiyuan Li, Tianhao Wang, Jason D Lee, and Sanjeev Arora. Implicit bias of gradient descent on reparametrized models: On equivalence to mirror descent. *Advances in Neural Information Processing Systems*, 35:34626–34640, 2022.

[49] Hartmut Maennel, Olivier Bousquet, and Sylvain Gelly. Gradient descent quantizes relu network features. *arXiv preprint arXiv:1803.08367*, 2018.

[50] Mary Phuong and Christoph H Lampert. The inductive bias of relu networks on orthogonally separable data. In *International Conference on Learning Representations*, 2020.

[51] Kaifeng Lyu, Zhiyuan Li, Runzhe Wang, and Sanjeev Arora. Gradient descent on two-layer nets: Margin maximization and simplicity bias. *Advances in Neural Information Processing Systems*, 34, 2021.

[52] Etienne Boursier, Loucas Pillaud-Vivien, and Nicolas Flammarion. Gradient flow dynamics of shallow relu networks for square loss and orthogonal inputs. *Advances in Neural Information Processing Systems*, 35:20105–20118, 2022.

[53] Mingze Wang and Chao Ma. Early stage convergence and global convergence of training mildly parameterized neural networks. *Advances in Neural Information Processing Systems*, 35: 743–756, 2022.

[54] Hancheng Min, René Vidal, and Enrique Mallada. Early neuron alignment in two-layer relu networks with small initialization. *arXiv preprint arXiv:2307.12851*, 2023.

[55] Mingze Wang and Chao Ma. Understanding multi-phase optimization dynamics and rich nonlinear behaviors of relu networks. *Advances in Neural Information Processing Systems*, 36, 2024.

[56] Song Mei, Andrea Montanari, and Phan-Minh Nguyen. A mean field view of the landscape of two-layer neural networks. *Proceedings of the National Academy of Sciences*, 115(33): E7665–E7671, 2018.

[57] Lenaic Chizat and Francis Bach. On the global convergence of gradient descent for over-parameterized models using optimal transport. *Advances in neural information processing systems*, 31, 2018.

[58] Justin Sirignano and Konstantinos Spiliopoulos. Mean field analysis of neural networks: A law of large numbers. *SIAM Journal on Applied Mathematics*, 80(2):725–752, 2020.

[59] Grant Rotskoff and Eric Vanden-Eijnden. Trainability and accuracy of artificial neural networks: An interacting particle system approach. *Communications on Pure and Applied Mathematics*, 75(9):1889–1935, 2022.

[60] Blake Bordelon and Cengiz Pehlevan. Self-consistent dynamical field theory of kernel evolution in wide neural networks. *Advances in Neural Information Processing Systems*, 35:32240–32256, 2022.

[61] Greg Yang, James B Simon, and Jeremy Bernstein. A spectral condition for feature learning. *arXiv preprint arXiv:2310.17813*, 2023.

[62] Simon S Du, Wei Hu, and Jason D Lee. Algorithmic regularization in learning deep homogeneous models: Layers are automatically balanced. *Advances in Neural Information Processing Systems*, 31, 2018.

[63] Maithra Raghu, Ben Poole, Jon Kleinberg, Surya Ganguli, and Jascha Sohl-Dickstein. On the expressive power of deep neural networks. In *international conference on machine learning*, pages 2847–2854. PMLR, 2017.

[64] Alex Krizhevsky, Ilya Sutskever, and Geoffrey E Hinton. Imagenet classification with deep convolutional neural networks. *Communications of the ACM*, 60(6):84–90, 2017.

[65] Leonardo Petrini, Francesco Cagnetta, Umberto M Tomasini, Alessandro Favero, and Matthieu Wyart. How deep neural networks learn compositional data: The random hierarchy model. *arXiv preprint arXiv:2307.02129*, 2023.

[66] Alethea Power, Yuri Burda, Harri Edwards, Igor Babuschkin, and Vedant Misra. Grokking: Generalization beyond overfitting on small algorithmic datasets. *arXiv preprint arXiv:2201.02177*, 2022.

[67] Tanishq Kumar, Blake Bordelon, Samuel J Gershman, and Cengiz Pehlevan. Grokking as the transition from lazy to rich training dynamics. *arXiv preprint arXiv:2310.06110*, 2023.

[68] Kaifeng Lyu, Jikai Jin, Zhiyuan Li, Simon Shaolei Du, Jason D Lee, and Wei Hu. Dichotomy of early and late phase implicit biases can provably induce grokking. In *The Twelfth International Conference on Learning Representations*, 2023.

[69] Neel Nanda, Lawrence Chan, Tom Lieberum, Jess Smith, and Jacob Steinhardt. Progress measures for grokking via mechanistic interpretability. *arXiv preprint arXiv:2301.05217*, 2023.

[70] Andrew Saxe, Shagun Sodhani, and Sam Jay Lewallen. The neural race reduction: Dynamics of abstraction in gated networks. In *International Conference on Machine Learning*, pages 19287–19309. PMLR, 2022.

[71] Daniel Kunin, Javier Sagastuy-Brena, Surya Ganguli, Daniel LK Yamins, and Hidenori Tanaka. Neural mechanics: Symmetry and broken conservation laws in deep learning dynamics. *arXiv preprint arXiv:2012.04728*, 2020.

[72] Hidenori Tanaka and Daniel Kunin. Noether's learning dynamics: Role of symmetry breaking in neural networks. *Advances in Neural Information Processing Systems*, 34:25646–25660, 2021.

[73] Feng Chen, Daniel Kunin, Atsushi Yamamura, and Surya Ganguli. Stochastic collapse: How gradient noise attracts sgd dynamics towards simpler subnetworks. *Advances in Neural Information Processing Systems*, 36, 2024.

[74] Gal Vardi and Ohad Shamir. Implicit regularization in relu networks with the square loss. In *Conference on Learning Theory*, pages 4224–4258. PMLR, 2021.

[75] Razvan Pascanu, Guido Montufar, and Yoshua Bengio. On the number of response regions of deep feed forward networks with piece-wise linear activations. *arXiv preprint arXiv:1312.6098*, 2013.

[76] Guido F Montufar, Razvan Pascanu, Kyunghyun Cho, and Yoshua Bengio. On the number of linear regions of deep neural networks. *Advances in neural information processing systems*, 27, 2014.

[77] Matus Telgarsky. Representation benefits of deep feedforward networks. *arXiv preprint arXiv:1509.08101*, 2015.

[78] Raman Arora, Amitabh Basu, Poorya Mianjy, and Anirbit Mukherjee. Understanding deep neural networks with rectified linear units. *arXiv preprint arXiv:1611.01491*, 2016.

[79] Thiago Serra, Christian Tjandraatmadja, and Srikumar Ramalingam. Bounding and counting linear regions of deep neural networks. In *International Conference on Machine Learning*, pages 4558–4566. PMLR, 2018.

[80] Boris Hanin and David Rolnick. Complexity of linear regions in deep networks. In *International Conference on Machine Learning*, pages 2596–2604. PMLR, 2019.

15

[81] Boris Hanin and David Rolnick. Deep relu networks have surprisingly few activation patterns. *Advances in neural information processing systems*, 32, 2019.

[82] Yann LeCun, Léon Bottou, Yoshua Bengio, and Patrick Haffner. Gradient-based learning applied to document recognition. *Proceedings of the IEEE*, 86(11):2278–2324, 1998.

[83] Kaiming He, Xiangyu Zhang, Shaoqing Ren, and Jian Sun. Delving deep into rectifiers: Surpassing human-level performance on imagenet classification. In *Proceedings of the IEEE international conference on computer vision*, pages 1026–1034, 2015.

A   Single-Neuron Linear Network

In this section, we provide a detailed analysis of the two-layer linear network with a single hidden neuron discussed in Section 3. The network is defined by the function $f(x; \theta) = aw^\mathsf{T}x$, where $a \in \mathbb{R}$ and $w \in \mathbb{R}^d$ are the parameters. We aim to understand the impact of the initializations a_0, w_0 and the layer-wise learning rates η_a, η_w on the training trajectory in parameter space, function space (defined by the product $\beta = aw$), and the evolution of the Neural Tangent Kernel (NTK) matrix K:

$$K = X \left(\eta_w a^2 \mathbf{I}_d + \eta_a ww^\mathsf{T} \right) X^\mathsf{T}. \tag{10}$$

The gradient flow dynamics are governed by the following coupled ODEs:

$$\dot{a} = -\eta_a w^\mathsf{T} \left(X^\mathsf{T} X aw - X^\mathsf{T}y \right), \qquad\qquad a(0) = a_0, \tag{11}$$
$$\dot{w} = -\eta_w a \left(X^\mathsf{T} X aw - X^\mathsf{T}y \right), \qquad\qquad w(0) = w_0. \tag{12}$$

The global minima of this problem are determined by the normal equations $X^\mathsf{T} X aw = X^\mathsf{T}y$. Even when $X^\mathsf{T} X$ is invertible, yielding a unique global minimum in function space $\beta_* = (X^\mathsf{T}X)^{-1}X^\mathsf{T}y$, the symmetry between a and w, permitting scaling transformations, $a \to a\alpha$ and $w \to w/\alpha$ for any $\alpha \neq 0$ without changing the product aw, results in a manifold of minima in parameter space. This minima manifold is a one-dimensional hyperbola where $aw = \beta_*$, with two distinct branches for positive and negative a. The set of saddle points $\{(a, w)\}$ forms a $(d - 1)$-dimensional subspace satisfying $a = 0$ and $w^\mathsf{T}X^\mathsf{T}y = 0$. Except for a measure zero set of initializations that converge to the saddle points, all gradient flow trajectories will converge to a global minimum. In Appendix A.2.5, we detail the basin of attraction for each branch of the minima manifold and the d-dimensional surface of initializations that converge to saddle points, separating the two basins.

A.1   Conserved quantity

The scaling symmetry between a and w results in a conserved quantity $\delta \in \mathbb{R}$ throughout training, as noted in many prior works [10, 62, 71], where

$$\delta = \eta_w a^2 - \eta_a \|w\|^2. \tag{13}$$

This can be easily verified by explicitly writing out the dynamics of δ. Define $\rho = \left(X^\mathsf{T} X aw - X^\mathsf{T}y \right)$ for succinct notation, such that

$$\begin{aligned}
\dot{\delta} &= 2\eta_w a\dot{a} - 2\eta_a w^\mathsf{T}\dot{w} \\
&= 2\eta_w a \left(-\eta_a w^\mathsf{T}\rho \right) - 2\eta_a w^\mathsf{T} \left(-\eta_w a\rho \right) \\
&= 0.
\end{aligned}$$

The conserved quantity confines the parameter dynamics to the surface of a hyperboloid where the magnitude and sign of the conserved quantity determines the geometry, as shown in Fig. 2. A hyperboloid of the form $\sum_{i=1}^{k} x_i^2 - \sum_{i=k+1}^{n} x_i^2 = \alpha$, with $\alpha \geq 0$, exhibits varied topology and geometry based on k and α. It has two sheets when $k = 1$ and one sheet otherwise. Its geometry is primarily dictated by α: as α tends to infinity, curvature decreases, while at $\alpha = 0$, a singularity occurs at the origin.

A.2   Exact solutions

To derive exact dynamics we assume the input data is whitened such that $X^\mathsf{T}X = \mathbf{I}_d$ and $\beta_* = X^\mathsf{T}y$ such that $\beta_* \neq 0$. The dynamics of a and w can then be simplified as

$$\dot{a} = \eta_a \left(w^\mathsf{T}\beta_* - a\|w\|^2 \right), \qquad\qquad a(0) = a_0 \tag{14}$$
$$\dot{w} = \eta_w \left(a\beta_* - a^2 w \right), \qquad\qquad w(0) = w_0. \tag{15}$$

A.2.1   Deriving the dynamics for μ and ϕ

As discussion in Section 3 we study the variables $\mu = a\|w\|$, an invariant under the rescale symmetry, and $\phi = \frac{w^\mathsf{T}\beta_*}{\|w\|\|\beta_*\|}$, the cosine of the angle between w and β_*. This change of variables can also be understood as a signed spherical decomposition of β: μ is the signed magnitude of β and ϕ is the

cosine angle between β and β_*. Through chain rule, we obtain the dynamics for μ and ϕ, which can be expressed as

$$\dot{\mu} = \sqrt{\delta^2 + 4\eta_a\eta_w\mu^2}\left(\phi\|\beta_*\| - \mu\right), \qquad\qquad \mu(0) = a_0\|w_0\|, \qquad\qquad (16)$$

$$\dot{\phi} = \frac{\eta_a\eta_w 2\mu\|\beta_*\|}{\sqrt{\delta^2 + 4\eta_a\eta_w\mu^2} - \delta}\left(1 - \phi^2\right), \qquad\qquad \phi(0) = \frac{w_0^\mathsf{T}\beta_*}{\|w_0\|\|\beta_*\|}. \qquad\qquad (17)$$

We leave the derivation to the reader, but emphasize that a key simplification used is to express the sum $\eta_w a^2 + \eta_a\|w\|^2$ in terms of δ,

$$\eta_w a^2 + \eta_a\|w\|^2 = \sqrt{\delta^2 + 4\eta_a\eta_w\mu^2}. \qquad\qquad (18)$$

Additionally, notice that η_a and η_w only appear in the dynamics for μ and ϕ as the product $\eta_a\eta_w$ or in the expression for δ. If we were to define $\mu' = \sqrt{\eta_a\eta_w}\mu$ and $\beta'_* = \sqrt{\eta_a\eta_w}\beta_*$, then it is not hard to show that the product $\eta_a\eta_w$ is absorbed into the dynamics. Thus, without loss of generality we can assume the product $\eta_a\eta_w = 1$, resulting in the following coupled system of nonlinear ODEs,

$$\dot{\mu} = \sqrt{\delta^2 + 4\mu^2}\left(\phi\|\beta_*\| - \mu\right), \qquad\qquad \mu(0) = a_0\|w_0\| \qquad\qquad (19)$$

$$\dot{\phi} = \frac{2\mu\|\beta_*\|}{\sqrt{\delta^2 + 4\mu^2} - \delta}\left(1 - \phi^2\right), \qquad\qquad \phi(0) = \frac{w_0^\mathsf{T}\beta_*}{\|w_0\|\|\beta_*\|} \qquad\qquad (20)$$

We will now show how to solve this system of equations for μ and ϕ. We will solve this system when $\delta = 0$, $\delta > 0$, and $\delta < 0$ separately. We will then in Appendix A.2.6 show a general treatment on how to obtain the individual coordinates of a and w from the solutions for μ and ϕ.

A.2.2 Balanced $\delta = 0$

When $\delta = 0$, the dynamics for μ, ϕ are,

$$\dot{\mu} = \mathrm{sgn}(\mu)2\mu(\phi\|\beta_*\| - \mu), \qquad\qquad \mu(0) = a_0\|w_0\|, \qquad\qquad (21)$$

$$\dot{\phi} = \mathrm{sgn}(\mu)\|\beta_*\|(1 - \phi^2), \qquad\qquad \phi(0) = \frac{w_0^\mathsf{T}\beta_*}{\|w_0\|\|\beta_*\|}. \qquad\qquad (22)$$

First, we show that the sign of μ cannot change through training and $\mathrm{sgn}(\mu) = \mathrm{sgn}(a)$. Because $\delta = 0$, the dynamics of a and w are constrained to a double cone with a singularity at the origin $(a = 0, w = 0)$. This point is a saddle point of the dynamics, so the trajectory cannot pass through this point to move from one cone to the other. In other words, the cone where the dynamics are initialized on is the cone they remain on. Without loss of generality, we assume $a_0 > 0$, and solve the dynamics. The dynamics of μ is a Bernoulli differential equation driven by a time-dependent signal $\phi\|\beta_*\|$. The dynamics of ϕ is decoupled from μ and is in the form of a Riccati equation evolving from an initial value ϕ_0 to 1, as we have assumed an initialization with positive a_0. This ODE is separable with the solution,

$$\phi(t) = \tanh\left(c_\phi + \|\beta_*\|t\right), \qquad\qquad (23)$$

where $c_\phi = \tanh^{-1}(\phi_0)$. Plugging this solution into the dynamics for μ gives a Bernoulli differential equation,

$$\dot{\mu} = 2\|\beta_*\|\tanh\left(c_\phi + \|\beta_*\|t\right)\mu - 2\mu^2, \qquad\qquad (24)$$

with the solution,

$$\mu(t) = \frac{2\cosh^2\left(c_\phi + \|\beta_*\|t\right)}{2\left(c_\phi + \|\beta_*\|t\right) + \sinh\left(2(c_\phi + \|\beta_*\|t)\right) + c_\mu}, \qquad\qquad (25)$$

where $c_\mu = 2\mu_0^{-1}\cosh^2(c_\phi) - (2c_\phi + \sinh(2c_\phi))$. Note, if $\phi_0 = -1$, then $\dot{\phi} = 0$, and the dynamics of μ will be driven to 0, which is a saddle point.

A.2.3 Upstream $\delta > 0$

When $\delta > 0$, the dynamics are constrained to a hyperboloid composed of two identical sheets determined by the sign of a_0 (as shown in Fig. 2 (c)). Without loss of generality we assume $a_0 > 0$, which ensures $a(t) > 0$ for all $t \geq 0$. However, unlike in the balanced setting, the dynamics of μ

and ϕ do not decouple, making it difficult to solve. Instead, we consider $\nu = \frac{w^\intercal \beta_*}{a}$, which evolves according to the Riccati equation,

$$\dot{\nu} = \|\beta_*\|^2 - \delta\nu - \nu^2, \qquad\qquad \nu(0) = \frac{w_0^\intercal \beta_*}{a_0}. \qquad (26)$$

The solution is given by,

$$\nu(t) = \frac{2R\nu_0 \cosh\left(Rt\right) + \left(2\|\beta_*\|^2 - \delta\nu_0\right)\sinh\left(Rt\right)}{2R\cosh\left(Rt\right) + (2\nu_0 + \delta)\sinh\left(Rt\right)}, \qquad (27)$$

where $R = \frac{1}{2}\sqrt{\delta^2 + 4\|\beta_*\|^2}$. The trajectory of $a(t)$ is given by the Bernoulli equation,

$$\dot{a} = a(\nu(t) + \delta - a^2), \qquad\qquad a(0) = a_0, \qquad (28)$$

which can be solved analytically using $\nu(t)$. We omit the solution due to its complexity, but provide a notebook used to generate our figures encoding the solution. From the solutions for ν, a, we can easily obtain dynamics for μ, ϕ.

A.2.4   Downstream $\delta < 0$

When $\delta < 0$, the dynamics are constrained to a hyperboloid composed of a single sheet (as shown in Fig. 2 (a)). However, unlike in the upstream setting, a may change sign. A zero-crossing in a leads to a finite time blowup in ν. Consequently, applying the approach used to solve for the dynamics in the upstream setting becomes more intricate. First we show the following lemma:

Lemma A.1. *If $a_0 \neq 0$ or $w_0^\intercal \beta_* \neq 0$, then $a(t)w(t)^\intercal \beta_* = 0$ has at most one solution for $t \geq 0$.*

Proof. Let $\omega(t) = w(t)^\intercal \beta_*$. The two-dimensional dynamics of $a(t)$ and $\omega(t)$ are given by,

$$\dot{a} = \omega - a(a^2 - \delta), \qquad (29)$$
$$\dot{\omega} = a\|\beta_*\|^2 - a^2\omega. \qquad (30)$$

Consider the orthant $O^+ = \{(a, \omega)|a > 0, \omega > 0\}$. The boundary ∂O^+ is formed by two orthogonal subspaces. On $\{(a, \omega)|a = 0, \omega \geq 0\}$, $\dot{a} \geq 0$. On $\{(a, \omega)|a \geq 0, \omega = 0\}$, $\dot{\omega} \geq 0$. Therefore, O^+ is a positively invariant set. Similarly, $O^- = \{(a, \omega)|a < 0, \omega < 0\}$ is a positively invariant set. On the boundary $\partial O^+ \cup \partial O^- = \{(a, \omega)|a\omega = 0\}$, the flow is contained only at the origin $a = 0, \omega = 0$, which represents all saddle points of the dynamics of (a, w). By assumption, (a, w) is not initialized at a saddle point, and thus the origin is not reachable for $t \geq 0$. As a result, the trajectory $(a(t), \omega(t))$ will at most intersect the boundary $\partial O_+ \cup \partial O_-$ once. $\qquad\square$

From Lemma A.1, we conclude that either a crosses zero, $w^\intercal \beta_*$ crosses zero, or neither crosses zero. When a doesn't cross zero, then ν is well-defined for $t \geq 0$, and our argument from Appendix A.2.3 still holds, leading to solutions for μ, ϕ. When a does cross zero, instead of ν, we consider $\upsilon = \frac{a}{w^\intercal \beta_*}$, the inverse of ν. In this case, we know from Lemma A.1 that $w^\intercal \beta_*$ does not cross zero and thus υ is well-defined for $t \geq 0$ and evolves according to the Riccatti equation,

$$\dot{\upsilon} = 1 + \delta\upsilon - \|\beta_*\|^2\upsilon^2, \qquad\qquad \upsilon(0) = \frac{a_0}{w_0^\intercal \beta_*}. \qquad (31)$$

These dynamics have a solution similar to Eq. (27), which we leave to the reader. With $\upsilon(t)$, we can then solve for the dynamics of $w^\intercal \beta_*$. Let $\omega = w^\intercal \beta_*$, then ω evolves according to the Bernoulli equation,

$$\dot{\omega} = \upsilon\|\beta\|^2\omega - \upsilon^2\omega^3, \qquad\qquad \omega(0) = w(0)^\intercal \beta_*, \qquad (32)$$

which can be solved analytically using $\upsilon(t)$. Again, we omit the solution due to its complexity, but provide a notebook used to generate our figures encoding the solution. From the solutions for υ, ω, we can easily obtain dynamics for μ, ϕ.

A.2.5 Basins of attraction

From Lemma A.1 we know that a can cross zero no more than once during its trajectory. Consequently, we can identify the basin of attraction by determining the conditions under which a changes sign. This analysis is crucial because initial conditions leading to a sign change in a correspond to scenarios where initial positive and negative values of a_0 are drawn towards the negative and positive branches of the minima manifold, respectively. From Eq. (27) we can immediately see that a will change sign when the denominator vanishes. This can happen if $\sqrt{\delta^2 + 4\|\beta_*\|^2} < -2\nu_0 - \delta$. For $\delta < 0$, this is satisfied if $\nu_0 < \frac{1}{2}\left(-\delta - \sqrt{\delta^2 + 4\|\beta_*\|^2}\right)$, which gives the hyperplane $w_0^\mathsf{T}\beta_* + \frac{a_0}{2}\left(\delta + \sqrt{\delta^2 + 4\|\beta_*\|^2}\right) = 0$ that separates between initializations for which a changes sign and initializations for which it does not (Fig. 7). Consequently, letting S^+ be the set of initializations attracted to the minimum manifold with $a > 0$, we have that:

$$S^+ = \left\{ (w_0, a_0) \;\middle|\; \begin{array}{ll} a_0 > 0 & \text{if } \delta \geq 0 \\ w_0^\mathsf{T}\beta_* > -\frac{a_0}{2}\left(\delta + \sqrt{\delta^2 + 4\|\beta_*\|^2}\right) & \text{if } \delta < 0 \end{array} \right\} \tag{33}$$

where the bottom inequality means that β_0 is sufficiently aligned to β_* in the case of $a_0 \geq 0$ or sufficiently misaligned in the case of $a_0 \leq 0$. We can similarly define the analogous S^-. An initialization on the separating hyperplane will converge to a saddle point where $w^\mathsf{T}\beta_* = a = 0$.


Figure 7: **Two basins of attraction.** For this model, parameter space is partitioned into two basins of attraction, one for the positive and negative branch of the minima manifold. The surface separating the basins of attraction is determined by the equation $w_0^\mathsf{T}\beta_* + \frac{a_0}{2}\left(\delta + \sqrt{\delta^2 + 4\|\beta_*\|^2}\right) = 0$. For a given δ, this equation describes a hyperplane through the origin. However, a given δ can only be achieved on the surface of some hyperboloid. Thus, the separating surface is the union of the intersections of a hyperplane and a hyperboloid, both parameterized by δ. This intersection is empty if $\delta > 0$. Initializations exactly on the separating surface will travel along the surface to a saddle point where $w^\mathsf{T}\beta_* = a = 0$.

A.2.6 Recovering parameters (a, w) from (μ, ϕ)

We now discuss how to recover the dynamics of the parameters (a, w) from our solutions for (μ, ϕ). We can recover a and $\|w\|$ from μ. Using Eq. (18) discussed previously, we can show

$$a = \operatorname{sgn}(\mu)\sqrt{\frac{\sqrt{\delta^2 + 4\mu^2} + \delta}{2}}, \qquad \|w\| = \sqrt{\frac{\sqrt{\delta^2 + 4\mu^2} - \delta}{2}}. \tag{34}$$

We now discuss how to obtain the vector w from ϕ. The key observation, as discussed in Section 3, is that w only moves in the span of w_0 and β_*. This means we can express $w(t)$ as

$$w(t) = c_1(t) \left(\frac{\beta_*}{\|\beta_*\|} \right) + c_2(t) \left(\frac{\left(\mathbf{I}_d - \frac{\beta_* \beta_*^\intercal}{\|\beta_*\|^2} \right) w_0}{\sqrt{\|w_0\|^2 - \left(\frac{\beta_*^\intercal w_0}{\|\beta_*\|} \right)^2}} \right) \tag{35}$$

where $c_1(t)$ is the coefficient in the direction of β_* and $c_2(t)$ is the coefficient in the direction orthogonal to β_* on the two-dimensional plane defined by w_0. From the definition of ϕ we can easily obtain the coefficients $c_1 = \|w\|\phi$ and $c_2 = \sqrt{\|w\|^2 - c_1^2}$. We always choose the positive square root for c_2, as $c_2(t) \geq 0$ for all t.
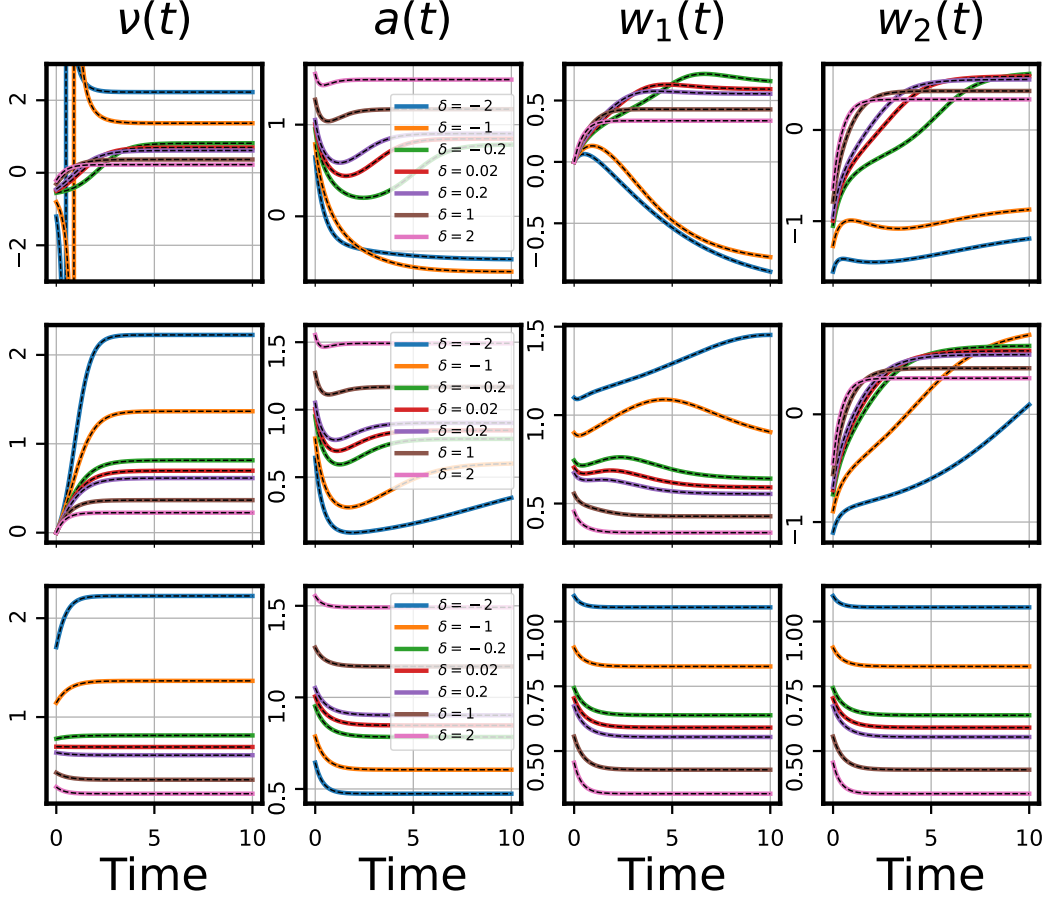

Figure 8: **Exact temporal dynamics of relevant variables in single-hidden neuron model.** Our theory recovers the time evolution under gradient flow of the quantities considered in this section, specifically ν, φ, and ζ, as well as the resulting dynamics of the model parameters $\{a, w_1, w_2\}$. The true $\beta*$ is a unit vector pointing in $\pi/4$ direction; $\beta(0)$ is a unit vector pointing towards $3\pi/2$, $-\pi/4$, and $\pi/4$ directions, respectively, for each of the three rows. δ then defines how $a(0)$ and $\|w(0)\|$ are chosen for a particular $\beta(0)$ where by convention we choose $a(0) > 0$.

A.3 Function space dynamics of β

The network's function is determined by the product $\beta = aw$ and governed by the ODE,

$$\dot{\beta} = a\dot{w} + \dot{a}w = - \underbrace{\left(\eta_w a^2 I_d + \eta_a ww^\intercal \right)}_{M} \underbrace{\left(X^\intercal X \beta - X^\intercal y \right)}_{X^\intercal \rho}. \tag{36}$$

Notice, that the vector $X^\intercal \rho$ driving the dynamics of β is the gradient of the loss with respect to β, $X^\intercal \rho = \nabla_\beta \mathcal{L}$. Thus, these dynamics can be interpreted as preconditioned gradient flow on the loss in

β space where the preconditioning matrix M depends on time through its dependence on a^2 and ww^T. The matrix M also characterizes the NTK matrix, $K = XMX^\mathsf{T}$. As discussed in Section 3, our goal is to understand the evolution of M along a trajectory $\{\beta(t) \in \mathbb{R}^d \;:\; t \geq 0\}$ solving Eq. (36).

First, notice that by expanding $\|\beta\|^2 = a^2\|w\|^2$ in terms of the conservation law, we can show

$$a^2 = \frac{\sqrt{\delta^2 + 4\eta_a\eta_w\|\beta\|^2} + \delta}{2\eta_w}, \tag{37}$$

which is the unique positive solution of the quadratic expression $\eta_w a^4 - \delta a^2 - \eta_a\|\beta\|^2 = 0$. When $a^2 > 0$ we can use this solution and the outer product $\beta\beta^\mathsf{T} = a^2 ww^\mathsf{T}$ to solve for ww^T in terms of β,

$$ww^\mathsf{T} = \frac{\sqrt{\delta^2 + 4\eta_a\eta_w\|\beta\|^2} - \delta}{2\eta_a}\frac{\beta\beta^\mathsf{T}}{\|\beta\|^2}. \tag{38}$$

Plugging these expressions into M gives

$$M = \frac{\sqrt{\delta^2 + 4\eta_a\eta_w\|\beta\|^2} + \delta}{2}I_d + \frac{\sqrt{\delta^2 + 4\eta_a\eta_w\|\beta\|^2} - \delta}{2}\frac{\beta\beta^\mathsf{T}}{\|\beta\|^2}. \tag{39}$$

Thus, given any initialization a_0, w_0 such that $a(t)^2 > 0$ for all $t \geq 0$, we can express the dynamics of β entirely in terms of β. This is true for all initializations with $\delta \geq 0$, except if initialized on the saddle point at the origin. It is also true for all initializations with $\delta < 0$ where the sign of a does not switch signs. In the next section we will show how to interpret these trajectories as time-warped mirror flows for a potential that depends on δ. As a means of keeping the analysis entirely in β space, we will make the slightly more restrictive assumption to only study trajectories given any initialization β_0 such that $\|\beta(t)\| > 0$ for all $t \geq 0$.

Notice, that η_a and η_w only appear in the dynamics for β as the product $\eta_a\eta_w$ or in the expression for δ. By defining $\beta' = \sqrt{\eta_a\eta_w}\beta$ and $y' = \sqrt{\eta_a\eta_w}y$ and studying the dynamics of β', we can absorb $\eta_a\eta_w$ into the β terms in M and the additional factor $\sqrt{\eta_a\eta_w}$ into the β and y terms in ρ. This transformation of β and y merely rescales β space without changing the loss landscape or location of critical points. As a result, from here on we will, without loss of generality, study the dynamics of β assuming $\eta_a\eta_w = 1$.

A.4 Deriving the inductive bias

Until now, we have primarily considered that $X^\mathsf{T}X$ is either whitened or full rank, ensuring the existence of a unique least squares solution β_*. In this setting, δ influences the trajectory the model takes from initialization to convergence, but all models eventually converge to the same point, as shown in Fig. 4. Now we consider the over-parameterized setting where we have more features d than observations n such that $X^\mathsf{T}X$ is low-rank and there exists infinitely many interpolating solutions in function space. By studying the structure of M we can characterize or even predict how δ determines which interpolating solution the dynamics converge to among all possible interpolating solutions. To do this we will extend a time-warped mirror flow analysis strategy pioneered by Azulay et al. [9].

A.4.1 Overview of time-warped mirror flow analysis

Here we recap the standard analysis for determining the implicit bias of a linear network through mirror flow. As first introduced in Gunasekar et al. [46], if the learning dynamics of the predictor β can be expressed as a *mirror flow* for some strictly convex potential $\Phi_\alpha(\beta)$,

$$\dot\beta = -\left(\nabla^2\Phi_\alpha(\beta)\right)^{-1}X^\mathsf{T}\rho, \tag{40}$$

where $\rho = (X\beta - y)$ is the residual, then the limiting solution of the dynamics is determined by the constrained optimization problem,

$$\beta(\infty) = \arg\min_{\beta\in\mathbb{R}^d} D_{\Phi_\alpha}(\beta, \beta(0)) \quad \text{s.t.} \quad X\beta = y, \tag{41}$$

where $D_{\Phi_\alpha}(p, q) = \Phi_\alpha(p) - \Phi_\alpha(q) - \langle\nabla\Phi_\alpha(q), p - q\rangle$ is the Bregman divergence defined with Φ_α. To understand the relationship between mirror flow Eq. (40) and the optimization problem Eq. (41), we consider an equivalent constrained optimization problem

$$\beta(\infty) = \arg\min_{\beta\in\mathbb{R}^d} Q(\beta) \quad \text{s.t.} \quad X\beta = y, \tag{42}$$

22

where $Q(\beta) = \Phi_\alpha(\beta) - \nabla\Phi_\alpha(\beta(0))^\mathsf{T}\beta$, which is often referred to as the *implicit bias*. $Q(\beta)$ is strictly convex, and thus it is sufficient to show that $\beta(\infty)$ is a first order KKT point of the constrained optimization (42). This is true iff there exists $\nu \in \mathbb{R}^n$ such that $\nabla Q(\beta(\infty)) = X^\mathsf{T}\nu$. The goal is to derive ν from the mirror flow Eq. (40). Notice, we can rewrite Eq. (40) as, $(\nabla\dot{\Phi}_\alpha(\beta)) = -X^\mathsf{T}\rho$, which integrated over time gives

$$\nabla\Phi_\alpha(\beta(\infty)) - \nabla\Phi_\alpha(\beta(0)) = -X^\mathsf{T}\int_0^\infty \rho(t)dt. \tag{43}$$

The LHS is $\nabla Q(\beta(\infty))$. Thus, by defining $\nu = \int_0^\infty \rho(t)dt$, which assumes the residual decays fast enough such that this is well defined, then we have shown the desired KKT condition. Crucial to this analysis is that there exists a solution to the second-order differential equation

$$\nabla^2\Phi_\alpha(\beta) = \left(\nabla_\theta\beta\nabla_\theta\beta^\mathsf{T}\right)^{-1}, \tag{44}$$

which even for extremely simple Jacobian maps may not be true [47]. Azulay et al. [9] showed that if there exists a smooth scalar function $g(\beta) : \mathbb{R}^d \to \mathbb{R}$ such that the ODE,

$$\nabla^2\Phi_\alpha(\beta) = g(\beta)\left(\nabla_\theta\beta\nabla_\theta\beta^\mathsf{T}\right)^{-1}, \tag{45}$$

has a solution, then the previous interpretation holds for $\Phi_\alpha(\beta)$ with $\nu = \int_0^\infty g(\beta(t'))\rho(t')dt'$. As before, it is crucial that this integral exists and is finite. Azulay et al. [9] further explained that this scalar function $g(\beta)$ can be considered as warping time $\tau(t) = \int_0^t g(\beta(t'))dt'$ on the trajectory taken in predictor space $\beta(\tau(t))$. So long as this warped time doesn't "stall out", that is we require that $\tau(\infty) = \infty$, then this will not change the interpolating solution.

A.4.2 Applying time-warped mirror flow analysis

Here show how to apply the time-warped mirror flow analysis to the dynamics of β derived in Appendix A.3 where $\nabla_\theta\beta\nabla_\theta\beta^\mathsf{T} = M$. We will only consider initializations β_0 such that $\|\beta(t)\| > 0$ for all $t \geq 0$, such that M can be expressed as

$$M = \frac{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}{2}I_d + \frac{\sqrt{\delta^2 + 4\|\beta\|^2} - \delta}{2}\frac{\beta\beta^\mathsf{T}}{\|\beta\|^2}. \tag{46}$$

Computing M^{-1}. Whenever $\|\beta\| > 0$, then M is a positive definite matrix with a unique inverse that can be derived using the Sherman–Morrison formula, $(A + uv^\mathsf{T})^{-1} = A^{-1} - \frac{A^{-1}uv^\mathsf{T}A^{-1}}{1+u^\mathsf{T}A^{-1}v}$. Here we can define A, u, and v as

$$A = \left(\frac{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}{2}\right)I_d, \quad u = \left(\frac{\sqrt{\delta^2 + 4\|\beta\|^2} - \delta}{2\|\beta\|^2}\right)\beta, \quad v = \beta \tag{47}$$

First notice the following simplification, $u^\mathsf{T}A^{-1}v = \frac{\sqrt{\delta^2+4\|\beta\|^2}-\delta}{\sqrt{\delta^2+4\|\beta\|^2}+\delta}$. After some algebra, M^{-1} is

$$M^{-1} = \left(\frac{2}{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}\right)I_d - \left(\frac{\frac{\sqrt{\delta^2+4\|\beta\|^2}-\delta}{\sqrt{\delta^2+4\|\beta\|^2}+\delta}}{\|\beta\|^2\sqrt{\delta^2+4\|\beta\|^2}}\right)\beta\beta^\mathsf{T} \tag{48}$$

To make notation simpler we will define the following two scalar functions,

$$f_\delta(x) = \frac{2}{\sqrt{\delta^2 + 4x} + \delta}, \qquad h_\delta(x) = \frac{\sqrt{\delta^2 + 4x} - \delta}{x\sqrt{\delta^2 + 4x}\left(\sqrt{\delta^2 + 4x} + \delta\right)}, \tag{49}$$

such that we can express $M^{-1} = f_\delta\left(\|\beta\|^2\right)I_d - h_\delta\left(\|\beta\|^2\right)\beta\beta^\mathsf{T}$.

Proving M^{-1} is not a Hessian map. If M^{-1} is the Hessian of some potential, then we can show that the dynamics of β are a mirror flow. However, from our expression for M^{-1} we can actually prove that it is *not* a Hessian map. As discussed in Gunasekar et al. [47], a symmetric matrix $H(\beta)$ is the Hessian of some potential $\Phi(\beta)$ if and only if it satisfies the condition,

$$\forall\beta \in \mathbb{R}^m, \quad \forall i, j, k \in [m] \quad \frac{\partial H_{ij}(\beta)}{\partial\beta_k} = \frac{\partial H_{ik}(\beta)}{\partial\beta_j}. \tag{50}$$

23

We will use this property to show M^{-1} is not a Hessian map. First, notice this condition is trivially true when $i = j = k$. Second, notice that for all $i \neq j \neq k$,

$$\frac{\partial M_{ij}^{-1}}{\partial \beta_k} = \frac{\partial M_{ik}^{-1}}{\partial \beta_j} = -2\nabla h_\delta\left(\|\beta\|^2\right)\beta_i\beta_j\beta_k \tag{51}$$

Thus, M^{-1} is a Hessian map if and only if for all $i \neq j$, $\frac{\partial M_{ii}^{-1}}{\partial \beta_j} = \frac{\partial M_{ij}^{-1}}{\partial \beta_i}$. Using our expression for M^{-1}, the LHS is

$$\frac{\partial M_{ii}^{-1}}{\partial \beta_j} = 2\nabla f_\delta\left(\|\beta\|^2\right)\beta_j - 2\nabla h_\delta\left(\|\beta\|^2\right)\beta_j\beta_i^2 \tag{52}$$

while the RHS is

$$\frac{\partial M_{ij}^{-1}}{\partial \beta_i} = -h_\delta\left(\|\beta\|^2\right)\beta_j - 2\nabla h_\delta\left(\|\beta\|^2\right)\beta_j\beta_i^2 \tag{53}$$

Thus, M^{-1} is a Hessian map if and only if $2\nabla f_\delta(x) + h_\delta(x) = 0$. Plugging in our definitions of $f_\delta(x)$ and $h_\delta(x)$ we find

$$2\nabla f_\delta(x) + h_\delta(x) = \frac{-4}{\sqrt{\delta^2 + 4x}(\sqrt{\delta^2 + 4x} + \delta)^2}, \tag{54}$$

which does not equal zero and thus M^{-1} is not a Hessian map.

Finding a scalar function $g_\delta(x)$ such that $g_\delta(\|\beta\|^2)M^{-1}$ is a Hessian map. While we have shown that M^{-1} is not a Hessian map, it is very close to a Hessian map. Here we will show that there exists a scalar function $g_\delta(x)$ such that $g_\delta\left(\|\beta\|^2\right)M^{-1}$ is a Hessian map. For any $g_\delta(x)$ can define $g_\delta\left(\|\beta\|^2\right)M^{-1}$ in terms of two new functions $\tilde{f}_\delta(x)$ and $\tilde{h}_\delta(x)$ evaluated at $x = \|\beta\|^2$,

$$g_\delta\left(\|\beta\|^2\right)M^{-1} = \underbrace{g_\delta\left(\|\beta\|^2\right)f_\delta\left(\|\beta\|^2\right)}_{\tilde{f}_\delta(\|\beta\|^2)}I_d - \underbrace{g_\delta\left(\|\beta\|^2\right)h_\delta\left(\|\beta\|^2\right)}_{\tilde{h}_\delta(\|\beta\|^2)}\beta\beta^\mathsf{T}. \tag{55}$$

Thus, as derived in the previous section, we get the analogous condition on $\tilde{f}_\delta(x)$ and $\tilde{h}_\delta(x)$ for $g_\delta\left(\|\beta\|^2\right)M^{-1}$ to be a Hessian map,

$$2\underbrace{\left(\nabla g_\delta(x)f_\delta(x) + g(x)\nabla f_\delta(x)\right)}_{\nabla \tilde{f}_\delta(x)} + \underbrace{g_\delta(x)h_\delta(x)}_{\tilde{h}_\delta(x)} = 0 \tag{56}$$

Rearranging terms we find that $g_\delta(x)$ must solve the ODE

$$\nabla g_\delta(x) = -\left(2f_\delta(x)\right)^{-1}\left(2\nabla f_\delta(x) + h_\delta(x)\right)g_\delta(x). \tag{57}$$

Using our previous expressions (Eq. (49) and Eq. (54)) we find

$$-\left(2f_\delta(x)\right)^{-1}\left(2\nabla f_\delta(x) + h_\delta(x)\right) = \frac{1}{\sqrt{\delta^2 + 4x}(\sqrt{\delta^2 + 4x} + \delta)}, \tag{58}$$

which implies $g_\delta(x)$ solves the differential equation, $\nabla g_\delta(x) = \frac{g_\delta(x)}{\sqrt{\delta^2 + 4x}(\sqrt{\delta^2 + 4x} + \delta)}$. The solution is $g_\delta(x) = c\sqrt{\sqrt{\delta^2 + 4x} + \delta}$, where $c \in \mathbb{R}$ is a constant. Let $c = 1$. Plugging in our expressions for $g_\delta\left(\|\beta\|^2\right), f_\delta\left(\|\beta\|^2\right), h_\delta\left(\|\beta\|^2\right)$, we get that

$$g_\delta\left(\|\beta\|^2\right)M^{-1} = \left(\frac{2}{\sqrt{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}}\right)I_d - \left(\frac{\sqrt{\frac{\sqrt{\delta^2 + 4\|\beta\|^2} - \delta}{\sqrt{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}}}}{\|\beta\|^2\sqrt{\delta^2 + 4\|\beta\|^2}}\right)\beta\beta^\mathsf{T} \tag{59}$$

is a Hessian map for some unknown potential $\Phi_\delta(\beta)$.

Solving for the potential $\Phi_\delta(\beta)$. Take the ansatz that there exists some function scalar $q(x)$ such that $\Phi_\delta(\beta) = q_\delta(\|\beta\|) + c_\delta$ where c_δ is a constant such that $\Phi_\delta(\beta) > 0$ for all $\beta \neq 0$ and $\Phi_\delta(0) = 0$. The Hessian of this ansatz takes the form,

$$\nabla^2\Phi_\delta(\beta) = \left(\frac{\nabla q(\|\beta\|)}{\|\beta\|}\right)I_d - \left(\frac{\nabla q(\|\beta\|)}{\|\beta\|^3} - \frac{\nabla^2 q(\|\beta\|)}{\|\beta\|^2}\right)\beta\beta^\mathsf{T}. \tag{60}$$

Equating terms from our expression for $g_\delta\left(\|\beta\|^2\right)M^{-1}$ (equation 59) we get the expression for $\nabla q(\|\beta\|)$

$$\nabla q(\|\beta\|) = \frac{2\|\beta\|}{\sqrt{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}}, \tag{61}$$

which plugged into the second term gives the expression for $\nabla^2 q(\|\beta\|)$,

$$\nabla^2 q(\|\beta\|) = \frac{2}{\sqrt{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta}} - \left(\frac{\frac{\sqrt{\delta^2+4\|\beta\|^2}-\delta}{\sqrt{\sqrt{\delta^2+4\|\beta\|^2}+\delta}}}{\sqrt{\delta^2+4\|\beta\|^2}}\right) = \frac{\sqrt{\sqrt{\delta^2+4\|\beta\|^2}+\delta}}{\sqrt{\delta^2+4\|\beta\|^2}}. \tag{62}$$

We now look for a function $q(x)$ such that both these conditions (Eq. (61) and Eq. (62)) are true. Consider the following function and its derivatives,

$$q(x) = \frac{1}{3}\left(\sqrt{\delta^2 + 4x^2} - 2\delta\right)\sqrt{\sqrt{\delta^2 + 4x^2} + \delta} \tag{63}$$

$$\nabla q(x) = \frac{2x}{\sqrt{\sqrt{\delta^2 + 4x^2} + \delta}} \tag{64}$$

$$\nabla^2 q(x) = \frac{\sqrt{\sqrt{\delta^2 + 4x^2} + \delta}}{\sqrt{\delta^2 + 4x^2}} \tag{65}$$

Letting $x = \|\beta\|$ notice $\nabla q(\|\beta\|)$ and $\nabla^2 q(\|\beta\|)$ satisfies the previous conditions. Furthermore, $\nabla^2 q(x) > 0$ for all δ as long as $x \neq 0$ and thus $q(x)$ is a convex function which achieves its minimum at $x = 0$. Thus, the constant $c_\delta = -q(0)$ is

$$c_\delta = \begin{cases} 0 & \text{if } \delta \leq 0 \\ \frac{\sqrt{2}|\delta|^{\frac{3}{2}}}{3} & \text{if } \delta > 0 \end{cases} = \max\left\{0, \text{sgn}(\delta)\frac{\sqrt{2}|\delta|^{\frac{3}{2}}}{3}\right\}, \tag{66}$$

and the potential $\Phi_\delta(\beta)$ is

$$\Phi_\delta(\beta) = \frac{1}{3}\left(\sqrt{\delta^2 + 4\|\beta\|^2} - 2\delta\right)\sqrt{\sqrt{\delta^2 + 4\|\beta\|^2} + \delta} + \max\left\{0, \text{sgn}(\delta)\frac{\sqrt{2}|\delta|^{\frac{3}{2}}}{3}\right\}. \tag{67}$$

Finally, putting it all together, we can express the inductive bias as in Theorem 3.1.

A.4.3 Connection to theorem 2 in Azulay et al. [9]

We discuss how Theorem 3.1 connects to Theorem 2 in Azulay et al. [9], which we rewrite:

Theorem A.2 (Theorem 2 from Azulay et al. [9]). *For a depth 2 fully connected network with a single hidden neuron ($h = 1$), any $\delta \geq 0$, and initialization β_0 such that $\beta_0 \neq 0$, if the gradient flow solution $\beta(\infty)$ satisfies $X\beta(\infty) = y$, then,*

$$\beta(\infty) = \underset{\beta \in \mathbb{R}^d}{\arg\min}\, q_\delta(\|\beta\|) + z^\mathsf{T}\beta \quad \text{s.t.} \quad X\beta = y \tag{68}$$

where $q_\delta(x) = \frac{\left(x^2 - \frac{\delta}{2}\left(\frac{\delta}{2} + \sqrt{x^2 + \frac{\delta^2}{4}}\right)\right)\sqrt{\sqrt{x^2 + \frac{\delta^2}{4}} - \frac{\delta}{2}}}{x}$ *and* $z = -\frac{3}{2}\sqrt{\sqrt{\|\beta_0\|^2 + \frac{\delta^2}{4}} - \frac{\delta}{2}}\frac{\beta_0}{\|\beta_0\|}$.

The most striking difference is in the expressions for the inductive bias. Azulay et al. [9] take an alternative route towards deriving the inductive bias by inverting M in terms of the original parameters a and w and then simplifying M^{-1} in terms of β, which results in quite a different expression for their inductive bias. However, they are actually functionally equivalent. It requires a bit of algebra, but one can show that

$$\Phi_\delta(\beta) = \frac{2\sqrt{2}}{3}q_\delta(\|\beta\|) + c_\delta. \tag{69}$$

Another important distinction between our two theorems lies in the assumptions we make. Azulay et al. [9] consider only initializations such that $\delta \geq 0$ and $\beta_0 \neq 0$. We make a less restrictive

25

assumption by considering initializations β_0 such that $\|\beta(t)\| > 0$ for all $t \geq 0$, which allows for both positive and negative δ. Except for a measure zero set of initializations, all initializations considered by Azulay et al. [9] also satisfy our assumptions. In both cases, our assumptions ensure that M is invertible for the entire trajectory from initialization to interpolating solution. However, it is worth considering whether the theorems would hold even when there exists a point on the trajectory where M is low-rank. As discussed in Appendix A.3, this can only happen for an initialization with $\delta < 0$ and where the sign of a changes. Only at the point where $a(t) = 0$ does M become low-rank. A similar challenge arose in this setting when deriving the exact solutions presented in Appendix A.2.4. We were able to circumvent the issue in part by introducing Lemma A.1 proving that this sign change could only happen at most once given any initialization. This lemma was based on the setting with whitened input, but a similar statement likely holds for the general setting. If this were the case, we could define M at this unique point on the trajectory in terms of the limit of M as it approached this point. This could potentially allow us to extend the time-warped mirror flow analysis to all initializations such that $\|\beta_0\| > 0$.

A.4.4 Exact solution when interpolating manifold is one-dimensional

When the null space of $X^\mathsf{T} X$ is one-dimensional, the constrained optimization problems in Theorem 3.1 and Theorem A.2 have an exact analytic solution. In this case we can parameterize all interpolating solutions β with a single scalar $\alpha \in \mathbb{R}$ such that $\beta = \beta_* + \alpha v$ where $X^\mathsf{T} X v = 0$ and $\|v\| = 1$. Using this description of β, we can then differentiate the inductive bias with respect to α, set to zero, and solve for α. We will use the following expressions,

$$\nabla_x q(x) = \frac{3}{2}\mathrm{sign}(x)\sqrt{\sqrt{x^2 + \frac{\delta^2}{4}} - \frac{\delta}{2}}, \quad \nabla_\alpha \|\beta\| = \frac{\alpha}{\|\beta\|}, \quad \nabla_\alpha z^\mathsf{T}\beta = z^\mathsf{T} v. \tag{70}$$

We will also use the expression, $\|\beta\|^2 = \|\beta_*\|^2 + \alpha^2$. Pulling these expressions together we get the following equation for α,

$$\sqrt{\sqrt{\|\beta_*\|^2 + \alpha^2 + \frac{\delta^2}{4}} - \frac{\delta}{2}} \frac{\alpha}{\sqrt{\|\beta_*\|^2 + \alpha^2}} = -\frac{2z^\mathsf{T} v}{3}. \tag{71}$$

If we let $k = -\frac{2z^\mathsf{T} v}{3}$, the solution for α is

$$\alpha = k\sqrt{\frac{k^2 + \delta}{2} + \sqrt{\left(\frac{k^2 + \delta}{2}\right)^2 + \|\beta_*\|^2}}. \tag{72}$$

This solution always works for the initializations we considered in Theorem 3.1. Interestingly, it appears that $\beta = \beta_* - \alpha v$ also works for initializations not previously considered. This includes trajectories that pass through the origin, resulting in a change in the sign of a.

B Wide and Deep Linear Networks

In the previous section we demonstrated how the balancedness δ influences the regime of learning in a single-neuron linear network by studying the dynamics in parameter space, function space, and the implicit bias. Throughout our analysis, we identified three learning regimes – lazy, rich, and delayed rich – that correspond to different values of δ. The driving cause of this distinction is the change in the geometry and topology of the conserved surface. Here we discuss how our analysis techniques can be extended to linear networks with multiple neurons, layers, and outputs. As we move towards more complex networks, the number of conserved quantities will grow, one for each hidden-neuron. As a result, the analysis in this section will get more complex, but overall the main points identified in the single-neuron setting still hold.

B.1 Two layer function space dynamics.

We consider the dynamics of a two-layer linear network with h hidden neurons and c outputs, $f(x; \theta) = A^\mathsf{T} W x$, where $W \in \mathbb{R}^{h \times d}$ and $A \in \mathbb{R}^{h \times c}$. We assume that $\min\{d, c\} \le h \le \max\{d, c\}$, such that this parameterization can represent all linear maps from $\mathbb{R}^d \to \mathbb{R}^c$. As in the single-neuron setting, the rescaling symmetry in this model between the first and second layer implies the $h \times h$ matrix $\Delta = A_0 A_0^\mathsf{T} - W_0 W_0^\mathsf{T}$ determined at initialization remains conserved throughout gradient flow [62]. The NTK matrix can be expressed as $K = (\mathbf{I}_c \otimes X) (\eta_w A^\mathsf{T} A \oplus \eta_a W^\mathsf{T} W) (\mathbf{I}_c \otimes X^\mathsf{T})$, where $\otimes$ and $\oplus$ denote the Kronecker product and sum[7] respectively. We consider the dynamics of $\beta = W^\mathsf{T} A \in \mathbb{R}^{d \times c}$ in function space. The network function β is governed by the ODE,

$$\dot{\beta} = \dot{W}^\mathsf{T} A + W^\mathsf{T} \dot{A}. \tag{73}$$

Respectively W and A follow the temporal dynamics given by the ODE

$$\dot{W}^\mathsf{T} = -\eta_w X^\mathsf{T} (X\beta - Y) A^\mathsf{T}, \tag{74}$$

and

$$\dot{A} = -\eta_a W X^\mathsf{T} (X\beta - Y). \tag{75}$$

Replacing equations 75 and 74 in equation 73 we get

$$\dot{\beta} = -\left(\eta_w X^\mathsf{T} (X\beta - Y) A^\mathsf{T} A + \eta_a W^\mathsf{T} W X^\mathsf{T} (X\beta - Y) \right). \tag{76}$$

Vectorising using the identity $vec(\mathbf{ABC}) = (\mathbf{C}^\mathsf{T} \otimes \mathbf{A}) vec(\mathbf{B})$ equation 76 becomes

$$vec(\dot{\beta}) = -vec\left(\eta_w I_d X^\mathsf{T} \underbrace{(X\beta - Y)}_{P} A^\mathsf{T} A + \eta_a W^\mathsf{T} W X^\mathsf{T} \underbrace{(X\beta - Y)}_{P} I_c \right), \tag{77}$$

$$= -(\eta_w A^\mathsf{T} A \otimes I_c + \eta_a I_d \otimes W^\mathsf{T} W) vec(X^\mathsf{T} P), \tag{78}$$

$$= -\underbrace{(\eta_w A^\mathsf{T} A \oplus \eta_a W^\mathsf{T} W)}_{M} vec(X^\mathsf{T} P). \tag{79}$$

The vectorised form of the network function is given by

$$vec\left(\dot{\beta} \right) = -\underbrace{(\eta_w A^\mathsf{T} A \oplus \eta_a W^\mathsf{T} W)}_{M} vec(X^\mathsf{T} X\beta - X^\mathsf{T} Y). \tag{80}$$

Interpreting $M(\beta)$ in different limit and architectures

As in the single-neuron setting, we find that the dynamics of β can be expressed as gradient flow preconditioned by a matrix M that depends on quadratics of A and W.

Consider a single hidden neuron $i \in [h]$ of the multi-output model defined by the parameters $w_i \in \mathbb{R}^d$ and $a_i \in \mathbb{R}^c$. Let $\beta_i = w_i a_i^\mathsf{T}$ be the $\mathbb{R}^{d \times c}$ matrix representing the contribution of this hidden neuron

---

[7]The Kronecker sum is defined for square matrices $A \in \mathbb{R}^{c \times c}$ and $B \in \mathbb{R}^{d \times d}$ as $A \oplus B = A \otimes \mathbf{I}_d + \mathbf{I}_c \otimes B$.

to the input-output map of the network. As in the previous section, we will consider the two gram matrices $\beta_i^\mathsf{T}\beta_i \in \mathbb{R}^{c \times c}$ and $\beta_i\beta_i^\mathsf{T} \in \mathbb{R}^{d \times d}$,

$$\beta_i^\mathsf{T}\beta_i = \|w_i\|^2 a_i a_i^\mathsf{T}, \qquad \beta_i\beta_i^\mathsf{T} = \|a_i\|^2 w_i w_i^\mathsf{T}. \tag{81}$$

Notice that we can express $\|\beta_i\|_F^2$ as

$$\|\beta_i\|_F^2 = \mathrm{Tr}(\beta_i^\mathsf{T}\beta_i) = \mathrm{Tr}(\beta_i\beta_i^\mathsf{T}) = \|a_i\|^2\|w_i\|^2 \tag{82}$$

At each hidden neuron we have the conserved quantity[8] $\eta_w\|a_i\|^2 - \eta_a\|w_i\|^2 = \delta_i$ where $\delta_i \in \mathbb{R}$. Using this quantity we can invert the expression for $\|\beta_i\|_F^2$ to get

$$\|a_i\|^2 = \frac{\sqrt{\delta_i^2 + \eta_a\eta_w 4\|\beta_i\|_F^2} + \delta_i}{2}, \tag{83}$$

$$\|w_i\|^2 = \frac{\sqrt{\delta_i^2 + \eta_a\eta_w 4\|\beta_i\|_F^2} - \delta_i}{2}. \tag{84}$$

When $\|\beta_i\|_F^2 > 0$, we can use these expressions to solve for the outer products $a_i a_i^\mathsf{T}$ and $w_i w_i^\mathsf{T}$ entirely in terms of β_i,

$$a_i a_i^\mathsf{T} = \frac{\sqrt{\delta_i^2 + \eta_a\eta_w 4\|\beta_i\|_F^2} + \delta_i}{2} \frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2}, \tag{85}$$

$$w_i w_i^\mathsf{T} = \frac{\sqrt{\delta_i^2 + \eta_a\eta_w 4\|\beta_i\|_F^2} - \delta_i}{2} \frac{\beta_i\beta_i^\mathsf{T}}{\|\beta_i\|_F^2}. \tag{86}$$

Without making any assumptions on the initialization (such as the isotropic initialization) we can express the NTK in terms of the β_i and consider the effect the vector of conserved quantities $\delta \in \mathbb{R}^h$ has on the dynamics.

Lemma B.1. *Assuming $\beta_i \neq 0$ for all $i \in [h]$ and let $\kappa_i = \sqrt{\delta_i^2 + 4\eta_a\eta_w\|\beta_i\|_F^2}$, then the matrix M can be expressed as the sum $M = \sum_{i=1}^h M_i$ over hidden neurons where M_i is defined as,*

$$M_i = \left(\frac{\kappa_i + \delta_i}{2}\right)\frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2} \oplus \left(\frac{\kappa_i - \delta_i}{2}\right)\frac{\beta_i\beta_i^\mathsf{T}}{\|\beta_i\|_F^2}. \tag{87}$$

We show how to express M in terms of the matrices $\beta_i = w_i a_i^\mathsf{T} \in \mathbb{R}^{d \times c}$, which represent the contribution to the input-output map of a single hidden neuron $i \in [h]$ of the network with parameters $w_i \in \mathbb{R}^d$, $a_i \in \mathbb{R}^c$, and conserved quantity $\delta_i = \Delta_{ii}$.

B.1.1 Funnel Networks

We consider *funnel networks*, which narrow from input to output ($d > h \geq c$), and *inverted-funnel networks*, which expand from input to output ($d \leq h < c$).

As $\delta_i \to \infty$, $M \to \sum_{i=1}^h |\delta_i| \frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2} \otimes \mathbf{I}_d$.

Lemma B.2. *Consider the rank-one matrices $\frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2}$ in the space $\mathbb{R}^{c \times c}$. The rank of M is bounded by*

$$\mathrm{rank}\,(M) \leq \begin{cases} d \times h & \text{if } h < c, \text{ Low-rank} \\ d \times c & \text{if } h \geq c. \end{cases} \tag{88}$$

Proof. In this limit, the rank of M is given by

$$d \times \mathrm{rank}(\sum_{i=1}^h \frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2}) \leq d \times \sum_{i=1}^h \mathrm{rank}(\frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2}). \tag{89}$$

It follows that

$$\sum_{i=1}^h \mathrm{rank}\left(\frac{\beta_i^\mathsf{T}\beta_i}{\|\beta_i\|_F^2}\right) \leq \begin{cases} h & \text{if } h < c. \\ c & \text{if } h \geq c. \end{cases} \tag{90}$$

---

[8]As long as $c > 1$, then the surface of this $d + c$ hyperboloid is always connected, however its topology will depend on the relationship between d and c.

as the rank of the sum $\sum_{i=1}^{h} \frac{\beta_i^{\mathsf{T}} \beta_i}{\|\beta_i\|_F^2}$ is also at most $\operatorname{rank} c$. Therefore,

$$\operatorname{rank}(M) \leq \begin{cases} d \times h & \text{if } h < c. \\ d \times c & \text{if } h \geq c. \end{cases} \tag{91}$$

$\square$

According to lemma B.2

- If $h < c$, $\operatorname{rank}(M)$ is bounded by $d \times h$ and remains below $d \times c$, categorizing it as a low-rank matrix. As a result, the solution $\beta_*^{\mathsf{T}} \beta_*$ might be in the null space of M. The network may enter either the lazy or a lazy followed by rich regime, depending on the relationship between β_0 and β_*. If $\beta_0 \propto \beta_*$ the network will enter the lazy regime.

- Assuming the $\frac{\beta_i^{\mathsf{T}} \beta_i}{\|\beta_i\|_F^2}$ terms are linearly independent, the matrix $\operatorname{rank}(M)$ achieves full rank and spans the solution space. The network can learn the task by only changing their norm while keeping their direction and the NTK matrix fixed. Thus, *funnel networks* defined by $(d > h \geq c)$ will transition into the lazy regime in this limit.

A similar assertion applies as $\delta_i \to -\infty$. In this limit, $M \to \mathbf{I}_c \otimes \sum_{i=1}^{h} |\delta_i| \frac{\beta_i \beta_i^{\mathsf{T}}}{\|\beta_i\|_F^2}$ with the $\operatorname{rank}(M)$ being constrained by the relationship between the number of hidden layers h and the input layer dimensions d.

- If $h < d$, the matrix M is low rank and bounded by $c \times h$. These networks may enter either the lazy or a lazy followed by rich regime, depending on the relationship between β_0 and β_*.

- When $h \geq d$, M is full rank if $\frac{\beta_i^{\mathsf{T}} \beta_i}{\|\beta_i\|_F^2}$ are linearly independent. Except for a measure zero set of initializations, *inverted funnel* networks always enter the lazy regime in this limit.

When $\delta_i = 0$, $M = \sqrt{\eta_a \eta_w} |\beta_i| \sum_{i=1}^{h} \frac{\beta_i^{\mathsf{T}} \beta_i}{\|\beta_i\|_F^2} \oplus \sum_{i=1}^{h} \frac{\beta_i \beta_i^{\mathsf{T}}}{\|\beta_i\|_F^2}$ all networks transition into the rich regime. Employing a similar rationale as before, assuming that all terms $\frac{\beta_i^{\mathsf{T}} \beta_i}{|\beta_i|_F^2}$ and $\frac{\beta_i \beta_i^{\mathsf{T}}}{|\beta_i|_F^2}$ are respectively linearly separable, then one term of M will be low-rank, while the other will be full-rank, contingent on the relationship between c, d, and h. Consequently, the dynamics of the network balance between low-rank and full-rank elements, leading to changes in both the magnitude and direction of the Neural Tangent Kernel (NTK) during training, a hallmark of the rich regime.

B.1.2 Single-Neuron

For *funnel network* with a single hidden neuron ($h = c = 1$), we recover equation 4 from equation 7. We extend this analysis to *inverted-funnel network* with a single hidden neuron ($h = d = 1$). Assuming $h = c = 1$, the rank one matrix $\frac{\beta_i^{\mathsf{T}} \beta_i}{\|\beta_i\|_F^2} = 1$. Therefore, equation 7 becomes

$$M = \left(\frac{\kappa + \delta}{2}\right) \oplus \left(\frac{\kappa - \delta}{2}\right) \frac{\beta \beta^{\mathsf{T}}}{\|\beta\|_F^2}. \tag{92}$$

$$= \frac{\kappa + \delta}{2} I_d + \frac{\kappa - \delta}{2} \frac{\beta \beta^{\mathsf{T}}}{\|\beta\|^2}. \tag{93}$$

where $\kappa = \sqrt{\delta^2 + \eta_a \eta_w 4 \|\beta\|^2}$. We recover the *funnel network* equation 4 for a single-neuron. In the main text, we analyze how the expression for $M(\beta)$ simplifies when δ approaches $-\infty$, 0, and ∞. This analysis helps us develop a deeper understanding of $M_\delta(\beta)$.

We now turn to single neuron *inverted funnel network* where $h = d = 1$, the rank one matrices $\frac{\beta_i \beta_i^{\mathsf{T}}}{\|\beta_i\|_F^2} = 1$. Therefore, equation 7 becomes

$$M = \left(\frac{\kappa + \delta}{2}\right) \frac{\beta^{\mathsf{T}} \beta}{\|\beta\|_F^2} \oplus \left(\frac{\kappa - \delta}{2}\right) \tag{94}$$

$$= \frac{\kappa + \delta}{2} \frac{\beta^{\mathsf{T}} \beta}{\|\beta\|_F^2} + \frac{\kappa - \delta}{2} I_c. \tag{95}$$

From our expression for $M(\beta)$ we will consider how it simplifies when $\delta \to -\infty, 0, \infty$.

$$M \to \begin{cases} \delta I_c & \delta \to -\infty. \\ \sqrt{\eta_a \eta_w} \|\beta\| \left(\frac{\beta^\intercal \beta}{\|\beta\|^2} + I_c \right) & \delta = 0. \\ |\delta| \frac{\beta^\intercal \beta}{\|\beta\|^2} & \delta \to \infty. \end{cases} \tag{96}$$

As anticipated, the single-neuron *inverted funnel network* also transitions into the ***Rich*** regime when $\delta = 0$. Under this condition, $M = \sqrt{\eta_a \eta_w} \|\beta\| (\frac{\beta^\intercal \beta}{\|\beta\|^2} + I_c)$, where the initial term denotes a projection matrix. Here the dynamics balance between following the lazy trajectory and attempting to fit the task by only changing in norm. As a result the NTK changes in both magnitude and direction through training, confirming the dynamics are rich.

Additionally, the single-neuron *inverted funnel network* follows analogous regime transitions to the single-neuron *funnel network*, albeit in opposite directions. As $\delta \to -\infty$, $M \to \delta I_c$, the single-neuron *inverted funnel network* network enters the ***Lazy*** regime. In this regime, the dynamics of β converge to the trajectory of linear regression trained by gradient flow and along this trajectory the NTK matrix remains constant.

Conversely, as $\delta \to +\infty$, $M \to |\delta| \frac{\beta^\intercal \beta}{\|\beta\|^2}$, transitions into the ***Lazy-to-Rich*** regime for the network. Here the dynamics of β are constrained to learn the task by only changing their norm while keeping the direction and the NTK matrix fixed – an initial lazy phase. However, if β must change direction to fit the task, and assuming finite δ or that $t \to \infty$ faster than $\delta \to -\infty$, then at some point there will be a slow alignment of β to β_*. In this second phase the NTK matrix will change, confirming the dynamics are lazy-to-rich.

In summary, networks with scalar outputs enter the lazy regime as $\delta \to \infty$, while networks with scalar inputs enter the lazy regime as $\delta \to -\infty$. Conversely, both types of networks enter the "active" regime in opposite directions. Finally, both cases of scalar output and scalar input enter the rich regime when $\delta \to 0$

B.1.3 Multi-Output

A fruitful setting for analysis is found in *square networks*, where the dimensions of the input, hidden, and output layers coincide ($d = h = c$). By studying the dependence of M on the conserved quantity $\mathrm{diag}(\Delta)$ and the shape of the network, defined by the the dimensions d, h and c, we can identify the lazy, rich, and lazy-to-rich regimes. We establish that as δ tends towards $\pm\infty$, the network symmetrically transitions into the lazy regime, while approaching zero, it converges into a rich regime. Furthermore, in this setting, we can precisely identify the influence δ_i has on the inductive bias.

We consider the isotropic initialization defined as $\Delta = \delta I_h$ in this section. The conserved quantity becomes $AA^\intercal - WW^\intercal = \delta I_h$.

Lemma B.3. *Assuming an initialization where $\|\beta(0)\|_F > 0$ and $\Delta = \delta I_h$, then the dynamics of the network can be expressed as $vec(\beta) = -Mvec(X^\intercal P)$, where M is defined as*

$$M = \left(\sqrt{\eta_a \eta_w \beta^\intercal \beta + \frac{\delta^2}{4} I} \otimes I \right) + \left(I \otimes \sqrt{\eta_a \eta_w \beta \beta^\intercal + \frac{\delta^2}{4} I} \right). \tag{97}$$

Proof. With some algebras, one can show the following relations.

$$W^\intercal W = \frac{1}{\eta_a} \left(-\frac{\delta}{2} I + \sqrt{\eta_a \eta_w \beta \beta^\intercal + \frac{\delta^2}{4} I} \right) \tag{98}$$

$$A(t)^\intercal A(t) = \frac{1}{\eta_w} \left(\frac{\delta}{2} I + \sqrt{\eta_a \eta_w \beta^\intercal \beta + \frac{\delta^2}{4} I} \right). \tag{99}$$

We start from

$$\mathrm{vec} \left(\dot{\beta} \right) = - \underbrace{(\eta_w A^\intercal A \oplus \eta_a W^\intercal W)}_{M} \mathrm{vec}(X^\intercal X \beta - X^\intercal Y), \tag{100}$$

where $P = (X\beta - Y)$. Replacing the expressions for $W^{\mathsf{T}}W$ (equation 98) and $A^{\mathsf{T}}A$ (equation 99); the vectorised network function dynamics (equation 80) reads

$$vec(\dot{\beta}) = - \left[\underbrace{\left(\sqrt{\eta_a \eta_w \beta^{\mathsf{T}}\beta + \frac{\delta^2}{4}I} \otimes I \right) + \left(I \otimes \sqrt{\eta_a \eta_w \beta\beta^{\mathsf{T}} + \frac{\delta^2}{4}I} \right)}_{M} \right] vec(X^{\mathsf{T}}P). \quad (101)$$

$\square$

From our expression for $M(\beta)$ we will consider how it simplifies in settings of $\delta \to -\infty, 0, \infty$ allowing us to gain intuition for $M_\delta(\beta)$.

$$M \to \begin{cases} \delta I & \delta \to -\infty. \\ \beta\beta^{\mathsf{T}} \otimes I + I \otimes \beta^{\mathsf{T}}\beta & \delta = 0. \\ \delta I & \delta \to \infty. \end{cases} \quad (102)$$

Lazy. As $\delta \to \pm\infty$ $M \to \delta I_d$, the network transitions into the lazy regime, the dynamics of β converge to the trajectory of linear regression trained by gradient flow and along this trajectory the NTK matrix remains constant, confirming the dynamics are lazy. **Rich.** When $\delta = 0$, $M = \beta\beta^{\mathsf{T}} \otimes I + I \otimes \beta^{\mathsf{T}}\beta$. As a result the NTK changes in both magnitude and direction through training, confirming the dynamics are rich.

Computing $M^{-1} = \nabla^2 q_\delta(\beta)$. Consistent with prior analyses, the natural next step of this derivation would be to compute the inverse of M in order to find the potential q_δ indicative of the implicit bias. However, it is not straight forward to take the inverse of a Kronecker sum. Therefore, we turn to a simplification of our setting where we can precisely identify the influence δ has on the inductive bias .

Assumption B.4. The singular value decomposition of the network function at initialization is $\text{SVD}(\beta(0)) = V\Lambda(0)U^{\mathsf{T}}$ where $\Lambda(0)$ is a diagonal matrix of strictly positive singular values and U, V are the singular vectors of $X^{\mathsf{T}}Y$, ($\text{SVD}(X^{\mathsf{T}}Y) = VSU^{\mathsf{T}}$),

Assumption B.5. The input data $X^{\mathsf{T}}X$ is diagonalizable by V.

Theorem B.6. *Under Assumptions B.4 and B.5, if the gradient flow solution $\beta(\infty)$ satisfies $X^{\mathsf{T}}\beta(\infty) = Y$, then $\beta(\infty) = U\Lambda_* V^{\mathsf{T}}$ where Λ_* solves the constrained optimization,*

$$\Lambda_* = \underset{\substack{\Lambda \in \text{Diag}(\mathbb{R}^d) \\ \Lambda_{ii} > 0 \quad for\ i=1,\ldots,i}}{\arg\min} Q_\delta(\Lambda) \quad s.t. \quad XU\Lambda V^{\mathsf{T}} = Y \quad (103)$$

where $Q_\delta(\Lambda) = \sum_i^h q_\delta(\Lambda_i) - \nabla q_\delta(\Lambda_i(0))\Lambda_i$ and

$q_\delta(x) = \frac{1}{4} \left(2x \sinh^{-1}\left(\frac{\sqrt{\eta_a \eta_w}2x}{\delta} \right) - \sqrt{\eta_a \eta_w 4x^2 + \delta^2} + \delta \right).$

Proof. Under assumptions B.4 and B.5, the network is said to be the task-aligned (Saxe et al. [10]), the eigenvectors of the network function β are constant throughout training and equal to U and V respectivaly. We can therefore describe the dynamics of $\Lambda \in \text{Diag}(\mathbb{R}^d)$ with $\Lambda_{ii} > 0$ for $i = 1, \ldots, d$ where $\Lambda = V^{\mathsf{T}}\beta U$.

Rearranging the expressions for $W^{\mathsf{T}}W$ (equation 98) and $A^{\mathsf{T}}A$ (equation 99) in terms of the singular vectors of the task we get,

$$A(t)^{\mathsf{T}}A(t) = \frac{1}{\eta_w}U \left(\frac{\delta}{2}I + \sqrt{\eta_a \eta_w \Lambda^2 + \left(\frac{\delta}{2}I\right)^2} \right) U^{\mathsf{T}}.$$

$$W^{\mathsf{T}}W = \frac{1}{\eta_a}V \left(-\frac{\delta}{2}I + \sqrt{\eta_a \eta_w \Lambda^2 + \left(\frac{\delta}{2}I\right)^2} \right) V^{\mathsf{T}}.$$

31

Beginning with the network function dynamic β, rewrite the expression in terms of the eigenvectors U and V

$$\dot{\beta} = -(\eta_w(X^\mathsf{T}X\beta - X^\mathsf{T}Y)A^\mathsf{T}A + \eta_a W^\mathsf{T}W(X^\mathsf{T}X\beta - X^\mathsf{T}Y)) \tag{104}$$

$$= -((X^\mathsf{T}XV\Lambda U^\mathsf{T} - VSU^\mathsf{T})U\Lambda_{AA}U^\mathsf{T} + V\Lambda_{WW}V^\mathsf{T}(X^\mathsf{T}XV\Lambda U^\mathsf{T} - VSU^\mathsf{T})) \tag{105}$$

$$= -(V((\Lambda - S)\Lambda_{AA} + \Lambda_{WW}(\Lambda - S))U^\mathsf{T}) \tag{106}$$

Projecting $\dot{\beta}$ onto the singular vectors V^T and U, we now consider the dynamics of $\dot{\Lambda}$

$$\dot{\Lambda} = -((\Lambda - S)\Lambda_{AA} + \Lambda_{WW}(\Lambda - S)) \tag{107}$$

$$= -\left((\Lambda - S)\left(\sqrt{\eta_a\eta_w\Lambda^2 + \left(\frac{\delta}{2}I\right)^2}\right) + \left(\sqrt{\eta_a\eta_w\Lambda^2 + \left(\frac{\delta}{2}I\right)^2}\right)(\Lambda - S)\right) \tag{108}$$

Given that $\sqrt{\eta_a\eta_w\Lambda^2 + \left(\frac{\delta}{2}I\right)^2}$ is a diagonal matrix and $(\Lambda - S)$ is also a diagonal matrix, the expression for $\eta_a\dot{\eta}_w\Lambda$ simplifies accordingly to

$$\dot{\Lambda} = -M_\Lambda(\Lambda - S) \tag{109}$$

with $M_\Lambda = \sqrt{\eta_a\eta_w 4\Lambda^2 + \delta^2 I}$.

As laid out in Appendix A.4.1 the standard analysis for determining the implicit bias of a linear network through mirror flow, if the learning dynamics of the task aligned predictor Λ can be expressed as a *mirror flow* for some strictly convex potential $\Phi_\delta(\beta)$,

$$\dot{\Lambda} = -\left(\nabla^2\Phi_\delta(\Lambda)\right)^{-1}(\Lambda - S) \tag{110}$$

Then, if the gradient flow solution $\beta(\infty)$ satisfies $X^\mathsf{T}\beta(\infty) = Y$, then $\beta(\infty) = U\Lambda_* V^\mathsf{T}$ where Λ_* solves the constrained optimization

$$\Lambda_* = \underset{\substack{\Lambda \in \mathrm{Diag}(\mathbb{R}^d) \\ \Lambda_{ii} > 0 \quad \text{for } i=1,\ldots,h}}{\arg\min} \quad Q_\delta(\Lambda) \quad \text{s.t.} \quad XU\Lambda V^\mathsf{T} = Y \tag{111}$$

where $Q_\delta(\Lambda) = \Phi_\delta(\Lambda) - \nabla\Phi_\delta(\Lambda(0))^\mathsf{T}\Lambda$

The natural next step of this derivation is to compute the inverse of M_Λ in order to find the potential Φ_δ indicative of the implicit bias. $M_\Lambda = \sqrt{\eta_a\eta_w 4\Lambda^2 + \delta^2 I}$ is a diagonal matrix therefore has a known inverse and

$$(M_\Lambda)^{-1} = \nabla^2\Phi_\alpha(\Lambda) = \left(\sqrt{\eta_a\eta_w 4\Lambda^2 + \delta^2 I}\right)^{-1} \tag{112}$$

Solving for the potential $\Phi_\delta(\Lambda)$**.** Consider the hypothesis that a scalar function $q(x)$ exists, allowing us to express $\Phi_\delta(\Lambda) = \sum_{i=1}^h q_\delta(\Lambda_i) + c_\delta$. Here, c_δ is a constant chosen to ensure that $\Phi_\delta(\Lambda_{ii}) > 0$ for all $\Lambda_{ii} > 0$ for $i = 1, \ldots, h$, and that $\Phi_\delta(0) = 0$. The Hessian of this ansatz takes the form,

$$\nabla^2\Phi_\delta(\Lambda) = \left(\sqrt{\eta_a\eta_w 4\Lambda^2 + \delta^2 I}\right)^{-1} \tag{113}$$

We now look for a function $q_\delta(x)$ and its derivatives,

$$q_\delta(\Lambda_i) = \frac{1}{4}\left(\sqrt{\eta_a\eta_w}2\Lambda_i \sinh^{-1}\left(\frac{\sqrt{\eta_a\eta_w}2\Lambda_i}{\delta}\right) - \sqrt{\eta_a\eta_w 4\Lambda_i^2 + \delta^2} + \delta\right) \tag{114}$$

The hessian of $q_\delta(x)$ is

$$\nabla^2\left(\sum_i^h q_\delta(\Lambda_i) + c_\delta\right) = \left(\sqrt{\eta_a\eta_w 4\Lambda^2 + \delta^2}\right)^{-1} \tag{115}$$

We find that the inductive bias is given by a hyperbolic entropy potential evaluated at the singular values of β. In the scenario of aligned networks, the distinction between ℓ^1 and ℓ^2 norms loses significance as the network consistently converges to the same solution. Nevertheless, the learning dynamics of β will differ due to the potential q_δ smooth transition between an ℓ^1 and ℓ^2 penalty, distinguishing the rich and lazy regimes. This potential was initially identified as the inductive bias for diagonal linear networks by Woodworth et al. [14]. $\qquad\square$

B.2 Multi-Layer

We now consider the influence of depth by studying a depth-$(l + 1)$ linear network, $f(x; \theta) = a^\intercal \prod_{i=1}^{l} W_i x$, where $W_1 \in \mathbb{R}^{h \times d}$, $W_i \in \mathbb{R}^{h \times h}$ for $1 < i \leq l$, and $a \in \mathbb{R}^h$. We assume that the dimensions $d = h$ and that all parameters share the same learning rate $\eta = 1$. For this model the predictor coefficients are computed by the product $\beta = \prod_{i=1}^{l} W_i^\intercal a \in \mathbb{R}^d$. Similar to our analysis of a two-layer setting, we assume an isotropic initializations of the parameters.

Definition B.7. There exists a $\delta \in \mathbb{R}$ such that $aa^\intercal - W_l W_l^\intercal = \delta \mathbf{I}_h$ and for all $i \in [l - 1]$ $W_{i+1}^\intercal W_{i+1} = W_i W_i^\intercal$.

This assumption can easily be achieved by setting $a = 0$ and $W_i = \alpha O_i$ for all $i \in [l]$, where $O_i \in \mathbb{R}^{d \times d}$ is an random orthogonal matrix and $\alpha \geq 0$. In this case $\delta = -\alpha^2$. Further, notice this parameterization is naturally achieved in the high-dimensional limit as $d \to \infty$ under a standard Gaussian initialization with a variance inversely proportional with width. As in the two-layer setting, this structure of the initialization will remain conserved throughout gradient flow. We now show how two natural quantities of β, its squared norm $\|\beta\|^2$ and its outer product $\beta\beta^\intercal$, can always be expressed as polynomials of $\|a\|^2$ and $W_1^\intercal W_1$ respectively.

$$\|\beta\|^2 = \|a\|^2 \left(\|a\|^2 - \delta\right)^l. \tag{116}$$

$$\beta\beta^\intercal = (W_1^\intercal W_1)^{l+1} + \delta (W_1^\intercal W_1)^l. \tag{117}$$

Proof. The norm of the regression coefficients is the product $\|\beta\|^2 = a^\intercal \left(\prod_{i=1}^{l} W_i\right) \left(\prod_{i=1}^{l} W_i\right)^\intercal a$. Using the conservation of the initial conditions between consecutive weight matrices, $W_{i+1}^\intercal W_{i+1} = W_i W_i^\intercal$, we can express this telescoped product as $\|\beta\|^2 = a^\intercal (W_l W_l^\intercal)^d a$. When plugging in the conservation between last two layers, this implies $\|\beta\|^2 = a^\intercal (aa^\intercal - \delta I_d)^d a$, which expanded gives the desired result.

The outer product of the regression coefficients is $\beta\beta^\intercal = \left(\prod_{i=1}^{l} W_i\right)^\intercal aa^\intercal \left(\prod_{i=1}^{l} W_i\right)$. Using the conserved initial conditions of the last weights we can factor the outer product as the sum, $\beta\beta^\intercal = \left(\prod_{i=1}^{l} W_i\right)^\intercal W_l W_l^\intercal \left(\prod_{i=1}^{l} W_i\right) + \delta \left(\prod_{i=1}^{l} W_i\right)^\intercal \left(\prod_{i=1}^{l} W_i\right)$. Both these telescoping products factor using the conservation of the initial conditions between consecutive weight matrices giving the desired result. $\qquad\square$

We now demonstrate how the quadratic terms $|a|^2$ and $W_1^\intercal W_1$ significantly influence the dynamics of β, similar to our analysis in the two-layer setting.

Lemma B.8. *The dynamics of β are given by a differential equation $\dot{\beta} = -MX^\intercal \rho$ where M is a positive semi-definite matrix that solely depends on $\|a\|^2$, $W_1^\intercal W_1$, and δ,*

$$M = (W_1^\intercal W_1)^l + \|a\|^2 \left(\sum_{i=0}^{l-1} (\|a\|^2 - \delta)^i (W_1^\intercal W_1)^{l-1-i}\right). \tag{118}$$

Proof. Using a similar telescoping strategy used in the previous proofs we can derive the form of M, which we leave to the reader. $\qquad\square$

Finally, we consider how the expression for M simplifies in the limit as $\delta \to 0$ allowing us to be precise about the inductive bias in this setting.

Theorem B.9. *For a depth-$(l + 1)$ linear network with square width ($d = h$) and isotropic initialization β_0 such that $\|\beta(t)\| > 0$ for all $t \geq 0$, then in the limit as $\delta \to 0$, if the gradient flow solution $\beta(\infty)$ satisfies $X\beta(\infty) = y$, then,*

$$\beta(\infty) = \arg\min_{\beta \in \mathbb{R}^d} \left(\frac{l+1}{l+2}\right) \|\beta\|^{\frac{l+2}{l+1}} - \left(\frac{\beta(0)}{\|\beta(0)\|^{\frac{l}{l+1}}}\right)^\intercal \beta \quad \text{s.t.} \quad X\beta = y. \tag{119}$$

Proof. Whenever $\|\beta\| > 0$ and in the limit as $\delta \to 0$, then we can find a unique expression for $\|a\|^2$ and $W_1^\mathsf{T} W_1$ in terms of $\|\beta\|^2$ and $\beta\beta^\mathsf{T}$,

$$\|a\|^2 = \|\beta\|^{\frac{2}{l+1}}, \qquad W_1^\mathsf{T} W_1 = \|\beta\|^{-\frac{2l}{l+1}} \beta\beta^\mathsf{T}. \tag{120}$$

Plugged into the previous expression for M results in a positive definite rank-one perturbation to the identity,

$$M = \|\beta\|^{\frac{2l}{l+1}} \mathbf{I}_d + l\|\beta\|^{-\frac{2}{l+1}} \beta\beta^\mathsf{T}. \tag{121}$$

Using the Sherman-Morrison formula we find that M^{-1} is

$$M^{-1} = \|\beta\|^{-\frac{2l}{l+1}} \mathbf{I}_d + \left(\frac{l}{l+1}\right) \|\beta\|^{-\frac{4l+2}{l+1}} \beta\beta^\mathsf{T} \tag{122}$$

We can now apply a time-warped mirror flow analysis similar to the analysis presented in Appendix A.4. Consider the time-warping function $g_\delta(\|\beta\|) = \|\beta\|^{-\frac{l}{l+1}}$ and the potential $\Phi(\beta) = \left(\frac{l+1}{l+2}\right) \|\beta\|^{\frac{l+2}{l+1}}$, then its not hard to show $M^{-1} = g_\delta(\|\beta\|)\nabla^2\Phi(\beta)$. This gives the desired result. $\square$

This theorem is a generalization of Proposition 1 derived in [9] for two-layer linear networks in the rich limit to deep linear networks in the rich limit. We find that the inductive bias, $Q(\beta) = (\frac{l+1}{l+2})\|\beta\|^{\frac{l+2}{l+1}} - \|\beta_0\|^{-\frac{l}{l+1}} \beta_0^\mathsf{T}\beta$, strikes a balance between attaining the minimum norm solution and preserving the initialization direction, which with increased depth emphasizes the latter.

C Piecewise Linear Networks

Here, we elaborate on the theoretical results presented in Section 5. Our goal is to extend the tools developed in our analysis of linear networks to piecewise linear networks and understand their limitations. We focus on the dynamics of the input-output map, rather than on the inductive bias of the interpolating solutions. As discussed in Azulay et al. [9], Vardi and Shamir [74], extending a mirror flow style analysis directly to non-trivial piecewise linear networks is very difficult or provably impossible. In this section, we first describe the properties of the input-output map of a piecewise linear function, then describe the dynamics of a two-layer network, and finally discuss the challenges in extending this analysis to deeper networks and potential directions for future work.

C.1 Surface of a piecewise linear network

The input-output map of a piecewise linear network $f(x; \theta)$, with l hidden layers and h hidden neurons per layer, is comprised of potentially $O(h^{dl})$ connected linear regions, each with their own vector of predictor coefficients [63]. The exploration of this complex surface has been the focus of numerous prior works, the vast majority of them focused on counting and bounding the number of linear regions as a function of the width and depth [75, 76, 77, 78, 63, 79, 80, 81]. The central object in all of these studies is the *activation region*,

Definition C.1. For a piecewise linear network $f(x; \theta)$, comprising m hidden neurons with pre-activation $z_i(x; \theta)$ for $i \in [m]$, let the *activation pattern* $\mathcal{A}$ represent an assignment of signs $a_i \in \{-1, 1\}$ to each hidden neuron. The *activation region* $\mathcal{R}(\mathcal{A}; \theta)$ is the subset of input space that generates $\mathcal{A}$,

$$\mathcal{R}(\mathcal{A}; \theta) = \{x \in \mathbb{R}^d \mid \forall i \; a_i z_i(x; \theta) > 0\}. \quad (123)$$

(a) With Biases   (b) Without Biases

Figure 9: **Surface of a ReLU network.** Here we depict the surface of a three-layer ReLU network $f(x; \theta) : \mathbb{R}^2 \to \mathbb{R}$ with twenty hidden units per layer at initialization, comparing configurations with biases (left) and without biases (right). The network with biases partitions input space into convex polytopes that tile input space. The network without biases partitions input space into convex conic sections emanating from the origin. Each region exhibits a distinct activation pattern, allowing the partition to be colored with two colors based on the parity of active neurons. The network operates linearly within each region and maintains continuity across boundaries.

The input-output map is linear within each non-empty activation region and continuous at the boundary between regions. Linearity implies that every non-empty[9] activation region is associated with a *linear predictor* vector $\beta_\mathcal{R} \in \mathbb{R}^d$ such that for all $x \in R(\mathcal{A}; \theta)$, $\beta_\mathcal{R} = \nabla_x f(x; \theta)$. Continuity implies that the boundary between regions is formed by a hyperplane determined by where the pre-activation for a neuron is exactly zero, $\{x : z_i(x; \theta) = 0\}$. When the neighboring regions have different linear predictors[10], then this hyperplane is orthogonal to their difference, which is a vector in the span of the first-layer weights. Taken together, this implies that the union of all activation regions forms a convex partition of input space, as shown in Fig. 9. We now present a surprisingly simple, yet to the best of our knowledge not previously understood property of this partition:

Proposition C.2 (2-colorable). *If $f(x; \theta)$ lacks redundant neurons, implying that every neuron influences an activation region, then the partition of input space can be colored with two distinct colors such that neighboring regions do not share the same color.*

The proof of this proposition is straightforward. There is one color for regions with an even number of active neurons and another for regions with an odd number of active neurons. Because $f(x; \theta)$ lacks redundant neurons, there does not exist a boundary between activation regions where two neurons activations change simultaneously. In this work, we solely utilize this proposition for visualization purposes, as shown in Fig. 9. Nonetheless, we believe it may be of independent interest as it strengthens the connection between the surface of piecewise linear networks and the *mathematics of paper folding*, a connection previously alluded to in the literature [76].

---

[9]While it is trivial to see that for a network $f(x; \theta)$ with m hidden neurons there are 2^m distinct activation patterns, not all activation patterns are attainable. See Raghu et al. [63] for a discussion.

[10]It is possible for neighboring regions to have the same linear predictor. Some works define linear regions as maximally connected component of input space with the same linear predictor [81].

C.2 Dynamics of a two-layer piecewise linear network

We consider the dynamics of a two-layer piecewise linear network without biases, $f(x; \theta) = a^\mathsf{T}\sigma(Wx)$, where $W \in \mathbb{R}^{h \times d}$ and $a \in \mathbb{R}^h$. The activation function is $\sigma(z) = \max\{z, \gamma z\}$ for $\gamma \in [0, 1)$, which includes ReLU $\gamma = 0$ and Leaky ReLU $\gamma \in (0, 1)$. We permit $h > d$, which in the limit as $h \to \infty$, ensures the network possesses the functional expressivity to represent any continuous nonlinear function from $\mathbb{R}^d$ to $\mathbb{R}$ passing through the origin. Following a similar strategy used in Section 4, we consider the contribution to the input-output map from a single hidden neuron $i \in [h]$ with parameters $w_i \in \mathbb{R}^d$ and $a_i \in \mathbb{R}$. As in the linear setting, each hidden neuron is associated with a conserved quantity, $\delta_i = \eta_w a_i^2 - \eta_a \|w_i\|^2$. Unlike in the linear setting, this neuron's contribution to the output $f(x_j; \theta)$ is regulated by whether the input x_j is in the neuron's *active halfspace*, $\{x \in \mathbb{R}^d : w_i^\mathsf{T}x > 0\}$. Let $C \in \mathbb{R}^{h \times n}$ be the matrix with elements $c_{ij} = \sigma'(w_i^\mathsf{T}x_j)$, which determines the activation of the i^{th} neuron for the j^{th} training data point. The subgradient $\sigma'(z) = 1$ if $z > 0$, $\sigma'(z) \in [\gamma, 1]$ if $z = 0$, and $\sigma'(z) = \gamma$ if $z < 0$. These activation functions exhibit positive homogeneity, implying $\sigma(z) = \sigma'(z)z$. Thus, we can express $\sigma(w_i^\mathsf{T}x_j) = c_{ij}w_i^\mathsf{T}x_j$, allowing us to express the gradient flow dynamics for w_i and a_i as

$$\dot{a}_i = -\eta_a w_i^\mathsf{T}\left(\sum_{j=1}^n c_{ij}x_j\rho_j\right), \qquad \dot{w}_i = -\eta_w a_i\left(\sum_{j=1}^n c_{ij}x_j\rho_j\right), \tag{124}$$

where $\rho_j = f(x_j; \theta) - y_j$ is the residual associated with the j^{th} training data point. If we let $\beta_i = a_i w_i$, which determines the contribution of each hidden neuron to the output $f(x_j; \theta)$, then its not hard to see that the gradient flow dynamics of β_i are

$$\dot{\beta}_i = -\underbrace{\left(\eta_w a_i^2 I_d + \eta_a w_i w_i^\mathsf{T}\right)}_{M_i}\underbrace{\left(\sum_{j=1}^n c_{ij}x_j\rho_j\right)}_{\xi_i}. \tag{125}$$

As in the linear setting, the matrix $M_i \in \mathbb{R}^{d \times d}$ appears as a preconditioning matrix on the dynamics Using the exact same derivation presented in Appendix A.3, whenever $a_i^2 > 0$, we can express M_i entirely in terms of β_i and δ_i,

$$M_i = \frac{\sqrt{\delta_i^2 + 4\eta_a\eta_w\|\beta_i\|^2} + \delta_i}{2}I_d + \frac{\sqrt{\delta_i^2 + 4\eta_a\eta_w\|\beta_i\|^2} - \delta_i}{2}\frac{\beta_i\beta_i^\mathsf{T}}{\|\beta_i\|^2}. \tag{126}$$

However, unlike in the linear setting, the vector $\xi_i \in \mathbb{R}^d$ driving the dynamics is not shared for all neurons because of its dependence on c_{ij}. Additionally, the NTK matrix in this setting depends on M_i and C, with elements $K_{jk} = \sum_{i=1}^h c_{ij}x_j^\mathsf{T}\left(\eta_w a_i^2 I_d + \eta_a w_i w_i^\mathsf{T}\right)x_k c_{ik}$. Thus, in order to asses the temporal dynamics of the NTK matrix, we must understand the dynamics of M_i and C. We consider a *signed spherical coordinate* transformation separating the dynamics of β_i into its directional $\hat{\beta}_i = \text{sgn}(a_i)\frac{\beta_i}{\|\beta_i\|}$ and radial $\mu_i = \text{sgn}(a_i)\|\beta_i\|$ components, such that $\beta_i = \mu_i\hat{\beta}_i$. Here, $\hat{\beta}_i$ determines the orientation and direction of the halfspace where the i^{th} neuron is active, while μ_i determines the slope of the linear region in this halfspace. These coordinates evolve according to,

$$\dot{\mu}_i = -\sqrt{\delta_i^2 + 4\eta_a\eta_w\mu_i^2}\hat{\beta}_i^\mathsf{T}\xi_i, \qquad \dot{\hat{\beta}}_i = -\frac{\sqrt{\delta_i^2 + 4\eta_a\eta_w\mu_i^2} + \delta_i}{2\mu_i}\left(I_d - \hat{\beta}_i\hat{\beta}_i^\mathsf{T}\right)\xi_i. \tag{127}$$

These equations can be derived directly from Eq. (124) through chain rule similar to Appendix A.2.1. In fact its worth noting that the this change of coordinates is very similar to the change of coordinates used in the single-neuron analysis. Expressed in terms of the parameters, $\hat{\beta}_i = \frac{w_i}{\|w_i\|}$ and $\mu_i = a_i\|w_i\|$.

D Experimental Details

We used Google Cloud Platform (GCP) nodes to run all experiments. Figure 1 experiments were run on a node with 360 AMD Genoa CPU cores with runtime totaling approximately 90 minutes including averaging over seeds as described below. Neural network training and NTK calculation for Figure 5 was performed on single A100 GPU nodes. Runtime was approximately 20 hours for Figure 5(a), four hours for 5(b), 12 hours for 5(c) (with individual runs ranging from five to 30 minutes depending on the number of datapoints), and 12 hours for 5(d). Figures 2, 3, and 4 are not compute-heavy, and these experiments were run on a personal computer. Overall, we estimate approximately 200 hours of single A100 runtime as well as 100 hours of the 360-core node accounting for failed runs and exploratory experiments.

D.1 Figure 1: Teacher-Student with Two-layer ReLU Networks

For Fig. 1 we consider a student-teacher setup similar to that in [8], with one-hidden layer ReLU networks of the form $f(x; \theta) = \sum_{i=1}^{m} a_i \sigma(w_i^\mathsf{T} x)$, where $f : \mathbb{R}^d \to \mathbb{R}$ and σ is the ReLU activation function. The teacher model, f^{teacher}, has $m = k$ hidden neurons initialized as $w_i^{\text{teacher}} \overset{\text{i.i.d.}}{\sim} \text{Unif}(S^{d-1})$ and $a_i \overset{\text{i.i.d.}}{\sim} \text{Unif}(\{\pm 1\})$ for $i \leq k$.

The student, f^{student}, in turn, has h hidden neurons. We use a symmetrized initialization, as considered in [8], where for $i \leq h/2$, we sample $w_i \overset{\text{i.i.d.}}{\sim} S^{d-1}$ and $a_i \overset{\text{i.i.d.}}{\sim} \text{Unif}(\{\pm 1\})$, and then for $i \geq \frac{h}{2} + 1$ we symmetrize by setting $w_i = w_{i-h/2}$ and $a_i = -a_{i-h/2}$. This ensures that f^{student} predicts 0 on any input at initialization.

Note that the *base* student initialization described thus far is perfectly balanced at each neuron, that is $\delta_i = 0$ for $i \in [m]$; we also define this to be our setting where the scale τ is 1. In order to transform the base initialization into a particular setting of τ and δ, we first solve for the relative layer scaling α in $\delta^2 = \tau^2(\alpha^2 - \alpha^{-2})$ and then scale each w_i by τ/α and each a_i by $\tau\alpha$.

We obtain a training dataset $\{x^{(i)}, y^{(i)}\}_{i=1}^n$ by sampling $x^{(i)} \overset{\text{i.i.d.}}{\sim} S^{d-1}$ and computing noiseless labels as $y^{(i)} = f^{\text{teacher}}(x^{(i)}; \theta^{\text{teacher}})$. The student is then trained with full-batch gradient descent on a mean square loss objective.

Figure 1 (a).

Here the setting is: $d = 2$, $h = 50$, $k = 3$, and $n = 20$. We sample a single teacher and then train four students with the same base initialization but different configurations of τ and δ: $(\tau = 0.1, \delta = 0)$ and $(\tau = 2, \delta = 0)$ for the left subfigure, and $(\tau = 0.1, \delta = 1)$ and $(\tau = 0.1, \delta = -1)$ for the right subfigure. Training is for 1 million steps at a learning rate of 1e-4.

Figure 1 (b).

Here the setting is: $d = 100$, $m = 50$, $k = 3$, and $n = 1000$, as in Fig. 1c of [8]. Training is performed with learning rate of 5e-3/τ^2. Test error is computed as mean square error over a held-out set of 10,000 datapoints.

We sweep over τ over a logarithmic scale in the range $[0.1, 2]$ and δ over a linear scale in the range $[-1, 1]$. We average over 16 random seeds, where the seed controls the sampling of: the teacher weights θ^{teacher}, the base initialization of θ^{student}, and the training data $\{x^{(i)}\}_{i=1}^n$. In this way, each random seed is used for a sweep over all combinations of τ and δ in the sweep; we simply apply the scaling described above to get to each point on the (τ, δ) grid.

The kernel distance computed is as defined in [25], where here we compute it at time t relative to the kernel at initialization, i.e. $S(t) = 1 - \langle K_0, K_t \rangle / (\|K_0\|_F \|K_t\|_F)$. In Fig. 10, we additionally plot Hamming and parameter distances relative to initialization, as well as training loss, while training for ten times longer than in Fig. 1 (b).
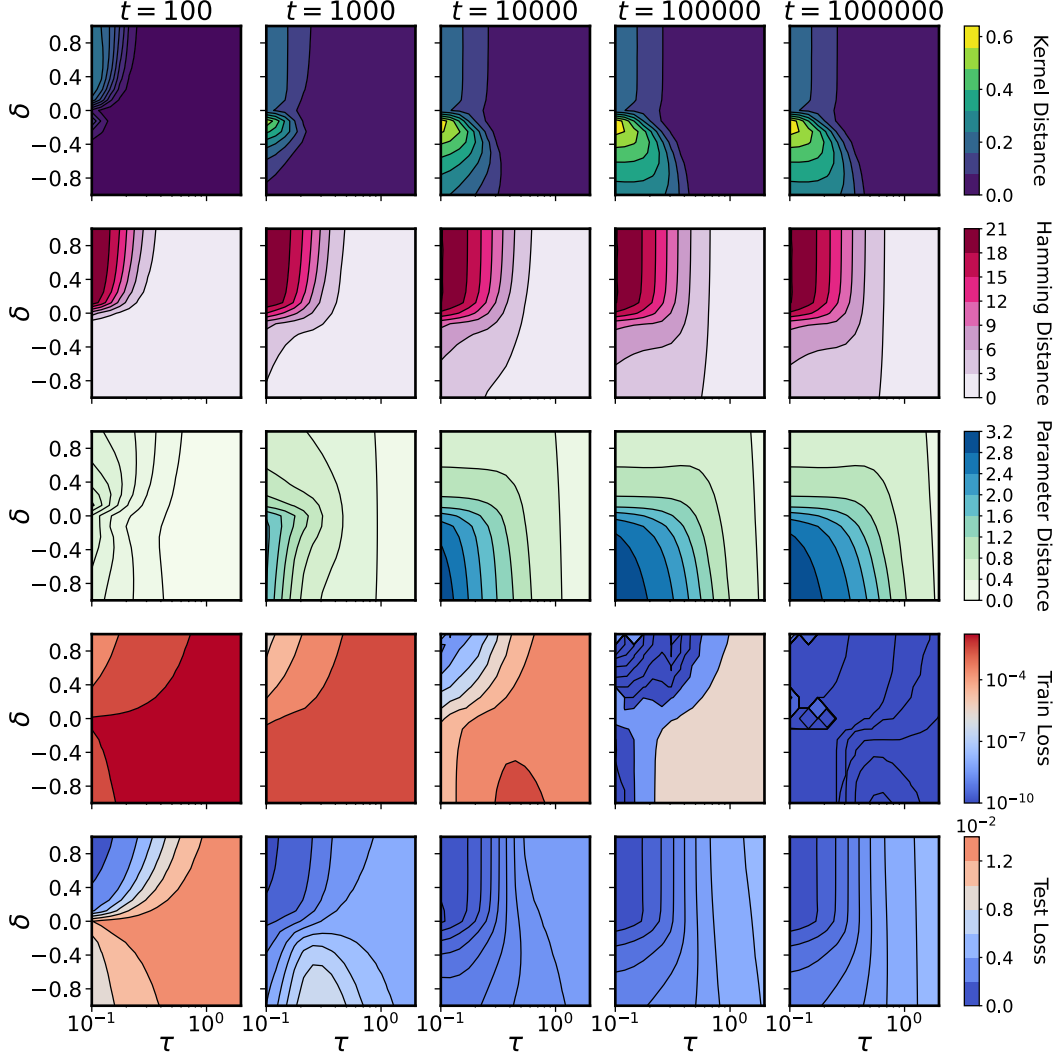
Figure 10: **Supporting figures for Fig. 1 (b).** We plot Hamming distance, parameter distance, and training loss, on top of the test loss and kernel distance considered in Fig. 1 (b), and train for ten times longer than in Fig. 1 (b). We observe that although training loss still drops between 10^5 and 10^6 steps, the test loss and other distances considered remain largely unchanged. Training loss is saturated at 1e-10.

D.2 Figures 2, 3, 4: Single-Neuron Linear Network

Figures 2, 3, and 4 were generated by simulating gradient flow using `scipy.integrate.solve_ivp` function with the RK45 method for solving the ODEs, with a relative tolerance of 1×10^{-6} and time span of $(0, 20)$. In the experiments with full-rank data, we used $X^\intercal X = \mathbf{I}_2$, $\beta_* = \begin{bmatrix} 0 \\ 1 \end{bmatrix}$, and $\beta_0 = \begin{bmatrix} -1 \\ 0 \end{bmatrix}$. For the experiment with low-rank data, we used $X^\intercal X = \begin{bmatrix} 0.25 & 0.5 \\ 0.5 & 1 \end{bmatrix}$, $\beta_* = \begin{bmatrix} 0.44 \\ 0.88 \end{bmatrix}$, and $\beta_0 = \begin{bmatrix} 0.4 \\ 0.05 \end{bmatrix}$. See the discussion in Appendix A.2 for details on how we determined our theoretical predictions. A notebook generating all the figures is provided.

D.3 Figure 5:

Kernel Distance

We trained LeNet-5 [82] (with ReLU nonlinearity and Max Pooling) on MNIST [82]. We use He initialization [83] and divide the first layer weights by α and multiply the last layer weights by α at initialization, which keeps the network functionally the same at initialization. We trained the model

for 500 epochs with a learning rate of 1e-4 and a batch size of 512. The parameter distance is defined as the L_2 distance between all the parameters. To quantify the distance between the activations, we binarize the hidden activation with 1 representing an active neuron. We evaluate Hamming distance over all the binarized hidden activations normalized by the the total number of the activations. We use kernel distance [25], defined as $S(t_1, t_2) = 1 - \langle K_{t_1}, K_{t_2} \rangle / (\|K_{t_1}\|_F \|K_{t_2}\|_F)$, which is a scale invariant measure of similarity between the NTK at two points in time. We subsample 10% of MNIST to evaluate the Hamming distance and kernel distance. All curves in the figure are averaged over 8 runs.

Gabor Filters

We are training a small ResNet based on the CIFAR10 script provided in the DAWN benchmark.[11] The only modification to the provided code base is that we set the weight decay parameter to 0, since this might confound our results. Moreover, we are dividing the convolutional filters weights by a parameter α (after standard initialization) which controls the balancedness of the network. To quantify the smoothness of the filters, we compute the normalized Laplacian of each filter $w_{ij} \in \mathbb{R}^{15 \times 15}$, over input $i = (1, 2, 3)$ and output $j = (1, ..., 64)$ channels

$$\text{smoothness}(w_{ij}) := \left\| \frac{w_{ij}}{\|w_{ij}\|_2} * \Delta \right\|_2^2 \tag{128}$$

where the Laplacian kernel is defined as

$$\Delta := \begin{pmatrix} -0.25 & -0.5 & -0.25 \\ -0.5 & 2 & -0.5 \\ -0.25 & -0.5 & -0.25 \end{pmatrix}. \tag{129}$$

(a) Accuracy   (b) Gabor Smoothness

Figure 11: **Interpreting convolutional filters.** CNN experiments on CIFAR10. We can see in **A)** that all networks achieve comparable training and test accuracy, despite the modification in initialization. However, in **B)** we see that networks with a small initialization ($\alpha < 1$) learn much smoother filters, giving quantiative support to results in Fig. 6. The smoothness is defined as the normalized Laplacian of the filters (see text, eq. 128).

Random Hierarchy Model

We refer to [65], who originally proposed the random hierarchy model (RHM) as a tool for studying how deep networks learn compositional data, for a more in-depth treatment. Here we briefly recap the setup following the notation used in [65].

An RHM essentially lets us build a random classification task with a clear hierarchical structure. The top level of the RHM specifies m equivalent high-level features for *each* class label in $\{1, \ldots, n_c\}$, where each feature has length s and n_c is the number of classes. For example, suppose the vocabulary at the top level is $\mathcal{V}_L = \{a, b, c\}$, $n_c = 2$, $m = 3$, and $s = 2$. Then in a particular instantiation of this RHM, we might have that Class 1 has ab, aa, and ca as equivalent high-level features (this is precisely the example used in Fig.1 of [65]). Class 2 will then have three random high-level features, with the constraint that they are **not** features for Class 1, for example, bb, bc, ac.

---

[11]Code available here.

Each successive level specifies m equivalent lower-level features for each "token" in the vocabulary at the previous level. For example, if $\mathcal{V}_{L-1} = \{d, e, f\}$, we might have that a can be equivalently represented as de, df, or ff; b and c will each have m equivalent representations of their own. We assume that the vocabulary size, v, is the same at all levels. Therefore, sampling an RHM with hyperparameters n_c, m, s, v requires sampling $mn_c + (L-1)mv$ rules.

In order to sample a datapoint from an RHM, we first sample a class label (e.g. Class 1), then uniformly sample one of the highest level features, (e.g. ab), then for each "token" in this feature we sample lower level features (e.g. $a \to de$, $b \to ee$), and so on recursively. The generated sample will therefore have length s^L and a class label. For training a neural network to perform this classification task, each input is converted into a one-hot representation, which will be of shape (s^L, v), and is then flattened.

We use the code released by [65] to train an MLP of width 64 with three hidden layers to learn an RHM with $L = 3$, $n_c = 8$, $m = 4$, $s = 2$, $v = 8$. The main change we make is allowing for scaling the initialization of the first layer by $1/\alpha$ and the initialization the readout layer by α. We then sweep over $\alpha \in \{0.03, 0.1, 0.3, 1, 3, 10\}$ and over the number of datapoints in the training set, which is specified as a fraction of the total number of datapoints the RHM can generate. We average test accuracy, which is by default computed on a held-out set of 20,000 samples, over six random seed configurations, where each configuration seeds the RHM, the neural network, and the data generation.

We train with the default settings used in [65], that is stochastic gradient descent with momentum of 0.9, run for 250 epochs with a learning rate initialized at 6.4 (0.1 times width) and decayed with a cosine schedule down to 80% of epochs. The batch size of 128; we do not use biases or weight decay.

Grokking

We are training a one layer transformer model on the modular arithmetic task in Power et al. [66]. Our experimental code is based on an existing Pytorch implementation.[12] The only modifications to the provided code base is that we use a single transformer layer (instead of the default 2-layer model). Prior analysis in Nanda et al. [69] has shown that this model can learn a minimal (attention-based) circuit that solves the task.

We study the effects on grokking time (defined as ≥ 0.99 accuracy on the validation data) of two manipulations. Firstly, we divide the embedding weights of the positional and token embeddings by the same balancedness parameter α as in the CNN gabor experiments. Secondly, like in Kumar et al. [67], we multiply the output of the model (i.e., the logits) by a factor τ and divide the learning rate by τ^2.


(a) Steps until training (b) Steps until grokked

Figure 12: **Transformer Grokking in Modular Arithmetic Task. A)** Shows the number of training steps required until the training accuracy passes a predefined threshold of 99%; we sample scaling $\tau \in \{0.5, 0.75, 1.0, 1.25, 1.5\}$ [67] and balance $\alpha \in \{0.1, 0.3, 1.0, 3.0, 10\}$ on a regular grid over $n = 5$ random initializations with a maximal computational budget of $m = 30{,}000$ training steps. **B)** Same as **A)**, but reporting the number of training steps required until the test performance passes the predefined threshold of 99%. We clearly see the fastest grokking in an unbalanced rich setting.

---

[12]Code available here.