

# Quantifying Variance in Evaluation Benchmarks

**Lovish Madaan**  <sup>$\alpha, \beta$</sup>     **Aaditya K. Singh**  <sup>$\gamma$</sup>     **Rylan Schaeffer**  <sup>$\delta$</sup>     **Andrew Poulton**  <sup>$\epsilon$</sup>   
**Sanmi Koyejo**  <sup>$\delta$</sup>     **Pontus Stenetorp**  <sup>$\beta$</sup>     **Sharan Narang**  <sup>$\alpha$</sup>     **Dieuwke Hupkes**  <sup>$\alpha$</sup>   
 <sup>$\alpha$</sup>  GenAI, Meta     <sup>$\beta$</sup>  UCL     <sup>$\gamma$</sup>  Gatsby Unit, UCL     <sup>$\delta$</sup>  Stanford University     <sup>$\epsilon$</sup>  Cohere  
{lovish, dieuwkehupkes}@meta.com

## Abstract

Evaluation benchmarks are the cornerstone of measuring capabilities of large language models (LLMs), as well as driving progress in said capabilities. Originally designed to make claims about capabilities (or lack thereof) in fully pretrained models, evaluation benchmarks are now also extensively used to decide between various training choices. Despite this widespread usage, we rarely quantify the variance in our evaluation benchmarks, which dictates whether differences in performance are meaningful. Here, we define and measure a range of metrics geared towards measuring variance in evaluation benchmarks, including seed variance across initialisations, and monotonicity during training. By studying a large number of models – both openly available and pretrained from scratch – we provide empirical estimates for a variety of variance metrics, with considerations and recommendations for practitioners. We also evaluate the utility and tradeoffs of continuous versus discrete performance measures and explore options for better understanding and reducing this variance. We find that simple changes, such as framing choice tasks (like MMLU) as completion tasks, can often reduce variance for smaller scale ( $\sim 7B$ ) models, while more involved methods inspired from human testing literature (such as item analysis and item response theory) struggle to meaningfully reduce variance. Overall, our work provides insights into variance in evaluation benchmarks, suggests LM-specific techniques to reduce variance, and more generally encourages practitioners to carefully factor in variance when comparing models.

## 1 Introduction

Benchmark evaluation datasets are the cornerstone of establishing and defining progress with large language models (LLMs). Virtually any new model release is accompanied by a range of scores on common evaluation benchmarks, illustrating how the model tallies up against previous releases (Mesnard et al., 2024; AI@Meta, 2024; Achiam et al., 2023; Reid et al., 2024). As such, evaluation datasets play an important role in claiming progress and the title of state-of-the-art. Consequently, choices in model development are often based on how they impact performance on benchmarks considered important by the field, giving benchmarks a prominent role in model iteration as well. Yet, despite their importance, benchmark scores are often regarded as a one-dimensional number, and it is rare that they are given a more detailed consideration. While it is well known that benchmarks scores can be heavily influenced by the choice of prompt (Sclar et al., 2023), the distributions of labels in the provided few-shots (Weber et al., 2023) or even the symbols that are used for the different options in a multiple choice setup (Zheng et al., 2023; Alzahrani et al., 2024), papers rarely report more than a single number per benchmark, or specifics on how each number was computed. Furthermore, statistical significance values are scarcely reported on major release papers or leaderboards, or even in papers that study how scores vary across various dimensions. These issues muddy the power of

evaluation datasets, both during development and evaluation: if we cannot ‘trust’ our evaluation results or do not understand what improvements are statistically significant, we cannot make sound comparisons, thus making it more challenging to reliably use benchmarks during model development.

To address this, we present a deep dive into variance in benchmark scores, at much larger scale than any previous work. Across all our experiments, we consider 13 different popular benchmarks and compute their performance over 280 different models, including fully trained public models as well as a set of 7B models and their intermediate checkpoints that we trained from scratch, differing only in their initialisation random seed.

With this, our contributions are three-fold:

- We provide a comprehensive reference guide for what magnitudes of variance are expected for what benchmarks across various circumstances.
- We make suggestions of how variance can be reduced for smaller scale models on choice tasks of important value (MMLU).
- We caution against the use of methods from human standardised testing (item analysis, item response theory) as a means of reducing variance, finding them to be ineffective.

Our work brings to light the often overlooked problem of variance in evaluation benchmarks, quantifies its effects, and provides a set of positive and negative results on how to mitigate it.

## 2 Models and Benchmarks

We run our analysis by comparing benchmark results across a large number of models trained across various setups. In this section, we describe these models and list the benchmarks we investigate.

### 2.1 Models

In our analysis, we use over 280 models for our analysis, including intermediate checkpoints. First, we train ten Llama-2-7B-architecture models from scratch on our own pre-training data mixture inspired by Touvron et al. (2023a) (See Appendix A). These 10 runs are identical, except for the model initialisation seed. The model hyper-parameters, the pre-training data mixture, and the data-loading mechanism is consistent across all these ten runs. We train these models for 210 billion tokens and store 21 checkpoints for each model, leaving us with 10 sets of 21 model snapshots. We refer to these 210 checkpoints as the “seed models.” In addition, we use 41 intermediate and fully-trained models based on the Llama-1 and Llama-2 architecture pre-trained on the same data mixture used for training the seed models.

Finally, we use 32 publicly available models from Huggingface (Wolf et al., 2020): Meta-Llama-3 {8, 70}B (AI@Meta, 2024), Gemma {2, 7}B (Mesnard et al., 2024), DBRX-Base (Databricks, 2024), Mistral 7B (Jiang et al., 2023), Mixtral 8x{7, 22}B (Jiang et al., 2024), Qwen-1.5 {0.5, 1.8, 4, 7, 14, 32, 72, 110}B (Bai et al., 2023), Pythia {1, 1.4, 2.8, 6.9, 12}B (Biderman et al., 2023), Falcon {7, 40}B (Almazrouei et al., 2023), DeepSeek {7, 67}B (Bi et al., 2024), DeepSeek-MoE 16B (Bi et al., 2024), DeepSeek V2 (DeepSeek-AI, 2024), StableLM {1.6, 3, 7}B (StabilityAI, 2024), and MPT {7, 30}B (MosaicML NLP Team, 2023).

The set of models used for the analysis are diverse across architectures, data mixtures, and sizes ranging from 0.5B to 236B total parameters. Details of all models are presented in Table 6.

### 2.2 Benchmarks

We do a comprehensive analysis using 13 large-scale well-established NLP benchmarks: AGIEval (Zhong et al., 2023), AI2 Reasoning Challenge (ARC-C) (Clark et al., 2018), BIG Bench (Hard) (Srivastava et al., 2022; Suzgun et al., 2022), COPA (Roemmele et al., 2011), GSM8k (Cobbe et al., 2021), Hellaswag (Zellers et al., 2019), HumanEval (Chen et al., 2021), MATH (Hendrycks et al., 2021), MMLU (Hendrycks et al., 2020), Natural Questions (Kwiatkowski et al., 2019), PIQA (Bisk et al., 2020), SIQA (Sap et al., 2019), and TriviaQA (Joshi et al., 2017).

These benchmarks are a mix of choice- and generation-based benchmarks, that span various capabilities ranging from general knowledge to coding.

### 3 How much variance do we observe?

We first investigate how much variance there is across different models and datasets. We define a range of metrics for quantifying different kinds of variance.

First, using the 7B models we trained ourselves, we consider variance due to changes in seed, across otherwise identical setups. This *seed variance* gives us a metric useful for performing data ablations – to conclude that pretraining dataset or hyperparameter set B is better than pretraining dataset or hyperparameter set A, we would want the performance increase to be larger than that due to random seed variance across different models trained in setup A. To this end, we also compute a dataset’s *monotonicity*, quantifying how stably it develops during training.

To ground the seed variance numbers, we compare them with bootstrapped 95% confidence intervals on individual models, as well as observed variance across different setups. In all our experiments, we consider both the (discrete) metric preferred for the benchmark and a more continuous representations for the same task.

#### 3.1 Analysis Methodology

For our initial variance analysis, we use both benchmark-level scores (to compute variance and monotonicity) and sample level scores (to estimate 95% confidence intervals). Here, we provide a brief description of the metrics we compute.

**Seed Mean ( $\mu(\mathcal{S}, \mathbb{M})$ )** We compute the performance using metric  $\mathcal{S}$  of the final checkpoint (at 210B tokens) of each of the 10 “fully trained” models in  $\mathbb{M}$  (one for each seed).

**Seed variance ( $\mathbb{E}(\mathcal{S}, \mathbb{M})$ )** Given a benchmark, a preferred metric  $\mathcal{S}$ , and a set of models  $\mathbb{M} = \{M_1, M_2, \dots, M_n\}$ , we define the benchmark seed variance  $\mathbb{E}(\mathcal{S}, \mathbb{M})$  as the standard deviation of the metric  $\mathcal{S}$  scores  $\{\mathcal{S}_{\mathbb{M}} = \mathcal{S}_{M_1}, \mathcal{S}_{M_2}, \dots, \mathcal{S}_{M_n}\}$  for each of the models in  $\mathbb{M}$ .

To estimate the variance expected due only to random seed changes, we take the average of this metric over all checkpoint timesteps  $\mathbb{E}(\mathcal{S}, \mathbb{M}) = \frac{1}{21} \sum_{time=\{10..210B\}} \mathbb{E}(\mathcal{S}, \mathbb{M}^{(time)})$ , where for example  $\mathbb{E}(\mathcal{S}, \mathbb{M}^{(time)})$  corresponds to the standard deviation of performance of the 10 model checkpoints (across seeds) after 200B tokens of training. For each benchmark, we consider both a discrete and a continuous metric.<sup>1</sup> The benchmark details are provided in Table 5 of Appendix A.

**Confidence intervals (95% CI)** We use the bootstrapped library<sup>2</sup> to compute 95% bootstrapped confidence interval (CI) values for each of the benchmarks on all 210 checkpoints from our 10 random seeded pretraining runs. Since bootstrapping is expensive, we also compute analytic interval (for discrete metrics) using the formula:

$$CI_{\text{analytic}}(M) = 1.96 * \sqrt{\frac{\mathcal{S}_M \times (1 - \mathcal{S}_M)}{N}},$$

where  $\mathcal{S}_M$  is the obtained preferred metric score for model M on a given benchmark and  $N$  is the number of test instances present in that benchmark. Empirically, we observe that, for the distributions we consider, bootstrapped and Analytic CIs converge when the number of bootstrap samples is large.

**Monotonicity values ( $\text{mon}_{\text{disc}}$  /  $\text{mon}_{\text{cont}}$ )** We compute the extent to which the scores for a benchmark develop monotonically during training. We define monotonicity for seed  $i$  as the Kendall Rank correlation between the list of scores  $[\mathcal{S}_{M_i^{10B}}, \mathcal{S}_{M_i^{20B}}, \dots, \mathcal{S}_{M_i^{210B}}]$  and a monotonically increasing or decreasing array of the same length, for discrete and continuous metrics, respectively.

#### 3.2 Results

In this section, we present our comprehensive analysis for two scenarios.

<sup>1</sup>With the exception of the datasets Big Bench (Hard), MATH, Natural Questions, and TriviaQA.

<sup>2</sup><https://github.com/facebookarchive/bootstrapped>

Table 1: **Variance values on 7B seed models.** Benchmarks are listed in alphabetical order. We report means -  $\mu(\mathcal{S}, \mathbb{M})$ , standard deviations -  $\mathbb{E}(\mathcal{S}, \mathbb{M})$ , confidence intervals - 95% CI, and monotonicities -  $\text{mon}_{\text{disc}}$ ,  $\text{mon}_{\text{cont}}$ . We also report size and chance level performance for reference—note all generative tasks have a chance level performance of 0.  $\mathbb{E}(\mathcal{S}, \mathbb{M})$  is generally lower than 95% CI. We also observe that  $\text{mon}_{\text{cont}} > \text{mon}_{\text{disc}}$  for all benchmarks.

Benchmark	Size	Chance	$\mu(\mathcal{S}, \mathbb{M})$	$\mathbb{E}(\mathcal{S}, \mathbb{M})$	95% CI	$\text{mon}_{\text{disc}}$	$\text{mon}_{\text{cont}}$
AGIEval	2546	20	23.44	0.77	1.63	0.37	0.29
ARC-C <sup>3</sup>	1165	25	39.71	0.80	2.74	0.88	0.91
Big Bench (Hard)	6511	0	29.10	0.87	1.07	0.77	-
COPA	100	50	78.80	2.15	8.30	0.56	0.90
GSM8k	1319	0	4.10	0.41	0.87	0.74	0.30
Hellaswag	10042	25	70.08	0.21	0.93	0.99	0.99
HumanEval	164	0	11.89	1.11	3.98	0.79	0.98
MATH	5000	0	1.52	0.23	0.28	0.52	-
MMLU	14042	25	25.86	0.57	0.72	0.09	0.15
MMLU-Cloze	14042	25	37.47	0.22	0.79	0.95	0.96
Natural Questions	3610	0	16.43	0.60	1.04	0.91	-
PIQA	1838	50	76.93	0.41	1.99	0.87	0.93
SIQA	1954	33	46.69	0.55	2.21	0.66	0.81
TriviaQA	11313	0	42.69	0.45	0.83	0.99	-

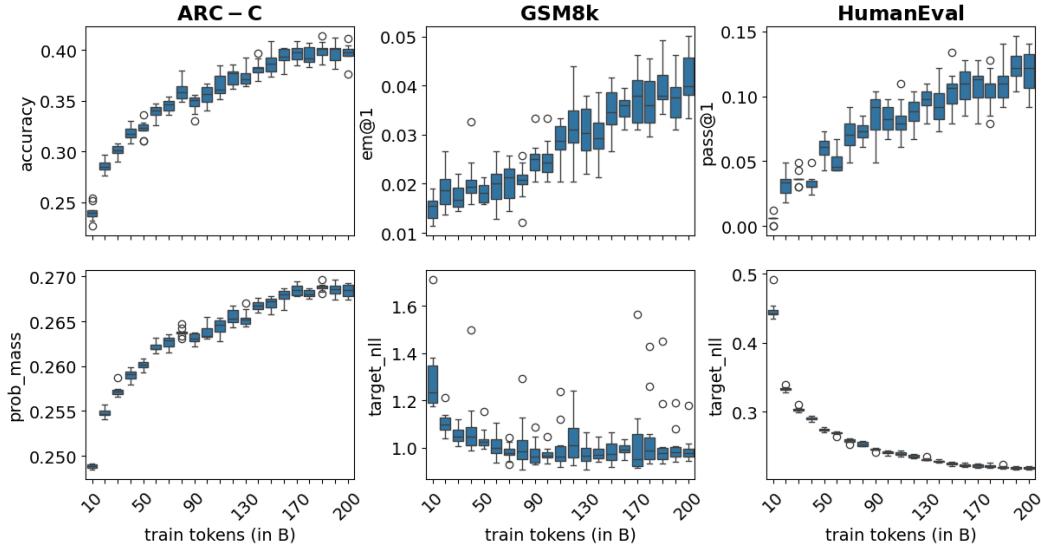
Table 2: **7B seed models.** Comparison between discrete (Disc) and continuous (Cont) metrics along with the signal to noise ratio (SNR). The means -  $\mu(\mathcal{S} = \text{Disc}, \mathbb{M})$ ,  $\mu(\mathcal{S} = \text{Cont}, \mathbb{M})$  and standard deviations (Disc Std, Cont Std) reported here (and used to calculate SNR) are computed across the final checkpoints across the 10 seeds.

Benchmark	$\mu(\mathcal{S} = \text{Disc}, \mathbb{M})$	Disc Std	Disc SNR	$\mu(\mathcal{S} = \text{Cont}, \mathbb{M})$	Cont Std	Cont SNR
AGIEval	23.44	0.93	25.20	0.2267	0.0009	254.93
ARC-C	39.71	0.87	45.89	0.2684	0.0007	381.64
COPA	78.80	2.04	38.63	0.5376	0.0008	662.41
GSM8k	4.10	0.52	7.88	0.9948	0.0653	15.24
Hellaswag	70.08	0.12	608.23	0.2833	0.0001	1921.15
HumanEval	11.89	1.75	6.79	0.2186	0.0018	124.08
MMLU	25.86	0.49	52.45	0.2511	0.0007	347.57
MMLU-Cloze	37.47	0.12	302.73	0.2698	0.0004	678.42
PIQA	76.93	0.39	198.98	0.5168	0.0003	1641.14
SIQA	46.69	0.51	91.87	0.3656	0.0009	387.11

**Seed variance** In Table 1, we report the observed variance across our 7B seed models in which the training setup is same across all init seeds, including a deterministic data ordering. We contextualise these numbers with the per-model 95% confidence interval, reported in the form of an average of 210 (one for each model) confidence interval sizes. The latter is easily computable from a single training run, whereas the former requires multiple (expensive) training runs with different seeds.

For some benchmarks (e.g. AGIEval, MMLU), scores are around chance accuracy ( $\sim 25\%$ ) even after training for 210B tokens. Benchmarks with few test examples (like COPA and HumanEval) exhibit high variance (both seed variance and 95% CIs). Generally, the 7B seed variance is well below the 95% CI for the same benchmark, though the ratio of the two is quite variable. Having access to the former value, which is smaller but closer to what would be needed to, for instance, compare two data mixes, may allow practitioners to make more fine-grained decisions during model development.

Motivated by prior work which suggests a move to continuous metrics (Srivastava et al., 2022; Schaeffer et al., 2023; Du et al., 2024; Schaeffer et al., 2024), we show a comparison of discrete and continuous metrics along with their signal to noise ratios (SNR) in Table 2. To maintain consistency, we used probability mass of the predicted answer for all choice-based benchmarks and NLL of the correct answer for generation-based benchmarks; more details are provided in Appendix A. We observe that the SNR is considerably higher for continuous metrics for all benchmarks, suggesting that they may be better when comparing models in the sense that they are less confounded by noise. These results may thus help in building better scaling laws for downstream evaluation tasks (Achiam et al., 2023), along with accurate comparisons between two models that have performances lying within the confidence interval for the discrete metric.



**Figure 1: Development of model performance over time.** Boxplots for both discrete and continuous metrics depicting the model improvement over time for ARC-C, GSM8k, and HumanEval. Top row depicts discrete metrics for each of the benchmarks, and the bottom row is composed of the continuous metrics.

**Monotonicity** In Table 1, we list the monotonicity values for each of the continuous and discrete metrics listed in Table 5. Higher monotonicity values are indicative of evaluations that more stably represent model improvement. In almost all cases, the monotonicity is better for the continuous metrics than for the discrete metrics, mirroring our findings with SNR above. However, for some benchmarks, such as HellaSwag and TriviaQA, the difference is minimal, likely since these benchmarks saturate earlier in training. Likewise, for benchmarks where performance remains at chance level we observe very low monotonicities.

In Figure 1, we visualise the development of discrete and continuous metrics and their seed variance during training, for ARC-C, GSM8k, and HumanEval. Generally (with the exception of GSM8k), continuous metrics have better predictive scaling compared to the discrete metrics because they have higher monotonicity and SNR. Interestingly, we see that the variance remains relatively constant as performance increases, suggesting that the estimates may extrapolate well to models trained for longer. Overall, these results suggests that monitoring continuous metrics could be more fruitful during model development than tracking discrete metrics.

### 3.3 The curious case of MMLU

Motivated by prior work considering the inconsistency of multiple choice benchmarks (Wang et al., 2024; Alzahrani et al., 2024), we examined two formulations of MMLU: (Standard) MMLU and MMLU-Cloze.

Standard MMLU refers to the prompting format where the choices along with the choice texts are present for the few-shot examples as well as the question in the prompt text. To evaluate the sample, we append the choice letters (“A”, “B”, “C”, or “D”) at the end of the prompt text, and pick the choice that has the lowest negative log-likelihood (NLL). For MMLU-Cloze, just the correct choice’s text is present for the few-shots, and we pick the choice that gives the lowest NLL after appending the choice texts at the end of the prompt. The complete prompts used for the two cases are detailed in Appendix B.

In Figure 2, we plot performance over training and see that standard MMLU is at chance performance even after training on 210B tokens. The cloze formulation performs better, and importantly has lower seed variance and much higher monotonicity (0.95 instead of 0.09, see Table 1). This result seems surprising, given that the cloze format is not standard. Further investigation yields that fully-trained large models tend to have better performance on standard MMLU compared to MMLU-cloze (e.g.

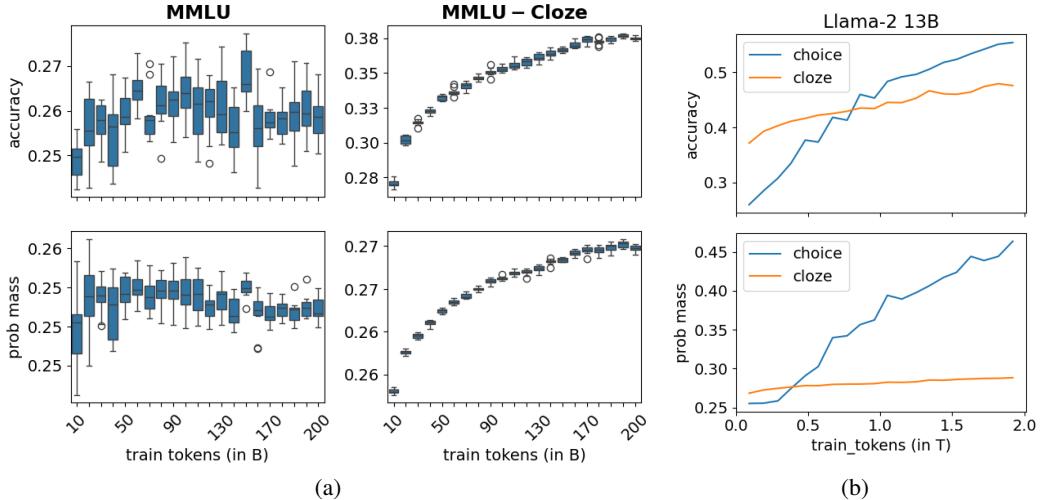


Figure 2: **Development of model performance over time.** In (a), we show the boxplots for the two MMLU variants. The top row is for the discrete metric (accuracy) and bottom row for the continuous metric (probability mass of the correct answer). In (b), we show the comparison of the standard (choice) and cloze variants on a Llama-2 13B model trained from scratch.

78.7% on standard MMLU vs. 60.6% for MMLU-Cloze for LLaMa 3 70B). Despite this difference in absolute performance, we find the performance on standard and cloze formats is highly correlated for fully trained large models (Pearson correlation of 0.92 on the 70 models listed in § 2.1).

To understand this dichotomy better, we train a Llama-2-13B-like model from scratch on our pre-training mix. We observe a sudden jump in performance at around 800B tokens (for both discrete and continuous metrics), after which standard MMLU performs better than MMLU-cloze (see Figure 2).

Given these results, we encourage researchers to use cloze formulations when doing pretraining ablations, as they are less confounded by noise during early stages of training, but still seem predictive of final performance on the standard MMLU format.

## 4 Understanding variance through the lens of item analysis

In the previous section, we computed the empirically occurring variances for commonly used evaluation benchmarks, considering benchmark-level scores, and we showed how looking at continuous metrics or cloze formulations of tasks can boost SNR.

As another avenue of possibly reducing variance, and to better understand it, we take inspiration from *item analysis*, a common method used to assess the usefulness of individual test questions on standardised tests administered to humans (Livingston, 2011; University of Washington, 2024). Item analysis focuses on metrics of individual samples (e.g. difficulty) to understand the types of questions on tests in terms of how individuals (in our case, models) perform on them.

### 4.1 Method

In applying item analysis to benchmarks, we consider two metrics. *Item difficulty* refers to the average score on an item across models; *Item discrimination* refers to the correlation between models’ performances on a single data point and models’ overall performances. Intuitively, items with either high or low difficulty will have low discrimination (as all models will be wrong or right, respectively).

As we wish to make recommendations about evaluation datasets that extend to future models, we split our 70 models into train and test sets. We consider two splits: “random” and “difficulty”. As the name suggests, in the random split, we split models randomly; In the difficulty split, we hold out the best performing 14 models. The full lists of models in each split can be found in Appendix D.1. We then calculate item analysis metrics on individual data points for the train and test sets. As is often done with human testing, we also consider the use of removing data points with low item discrimination,

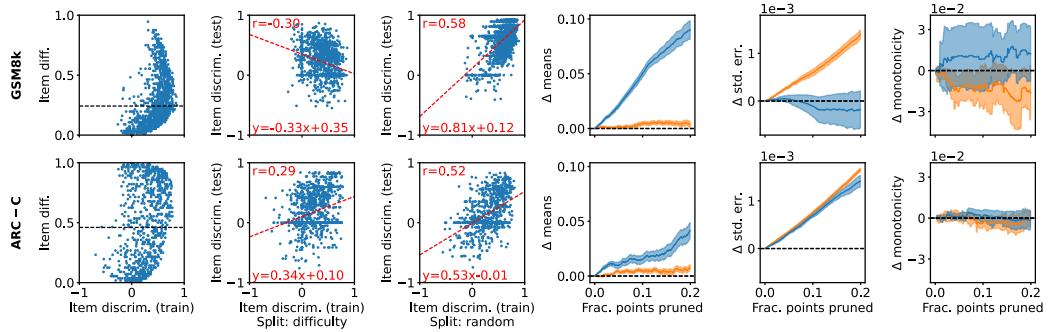


Figure 3: Item analysis results on GSM8k and ARC-C. Results on additional benchmarks provided in Appendix D.2. **First column** shows a scatter plot of item difficulty (x-axis) vs item discrimination (y-axis). **Second column** shows a scatter plot of item discrimination calculated over models from the train or test set of the difficulty split. **Third column** is the same as the second, except on the random split. As expected (since train and test splits come from the same distribution), discrimination on train models for this split is positively correlated to discrimination on test models. **Fourth, fifth, and sixth columns** show the effects of iteratively removing up to 20% of items (based on discrimination) on the mean (fourth column), standard error (fifth column) of model performance on the test set from the difficulty split by looking at the delta. Error bars indicate 95% confidence intervals in the delta. Monotonicity (sixth column) is calculated over the 10 runs from § 2.1. Orange curves show effects from randomly removing points, as a baseline.

and observe the effects this has on evaluation metrics such as mean, standard error of the mean (std. err.),<sup>4</sup> and monotonicity.

## 4.2 Results

In Figure 3, we show results for two illustrative benchmarks: ARC-C and GSM8k. Full results across other benchmarks can be found in Appendix D.2. Overall, we find that item discrimination scores may not provide much useful signal for the field of language model evaluations (unlike their widespread usage in human standardised testing). This is especially true given that state-of-the-art models perform better and better, and we would like tests to stay informative when models improve. To illustrate this, we show how high discrimination on train (weaker) models often does not correspond to high discrimination on test (stronger) models (Figure 3, second column). Stripping around  $x = 0$  corresponds to items that train set models always get *wrong* (yielding 0 discrimination) but are informative on test set models. Similarly, stripping around  $y = 0$  corresponds to items that test models always get *right* (yielding 0 discrimination) but are informative on the train set. If we instead consider item discriminations on a *random* split of models (Figure 3, third column), we see stronger correlations, indicating that the low correlation is in fact due to the difference in item discrimination on weaker and stronger models.

In Appendix D.3, we qualitatively inspect examples with negative item discrimination (which are thus anti-correlated with overall model performance), but are not able to discern any clear patterns for most benchmarks (a notable exception being Hellaswag, see Figure 8). While these negative results suggest item discriminations may not be the most informative means of understanding (or reducing) variance on stronger models, we consider further application to explore the causal effect.

Specifically, we consider pruning data points with low item discrimination, with the hopes that this will reduce variance or improve monotonicity. More precisely, we prune data points with low item discrimination on the train set of models from the difficulty split and we visualise metrics calculated using the pruned subset on the test set of models from the difficulty split. Results are presented in the three rightmost columns of Figure 3. Overall, while we find modest improvements in both standard error (a decrease) and monotonicity (an increase), the drift in the estimated accuracy is mildly concerning. It may be acceptable for the purpose of comparing models, but may also provide an overestimate of capabilities if considering the absolute score. One hypothesis for this discrepancy with human testing could be that item discrimination for human tests typically does not consider

<sup>4</sup>Note that the confidence intervals of § 3.1 are 1.96 times the standard error.

Table 3: **Variance values for Tiny Benchmark (across seeds).** Full represents the full benchmark, and IRT/IRT++ use the 100 examples proposed in Polo et al. (2024).  $\mathbb{E}(\mathcal{S}, \mathbb{M})$  is the seed variance defined in § 3.1, which is represented as  $\mathbb{E}$  in this table.

Benchmark	Full $\mu$	IRT $\mu$	IRT++ $\mu$	Full $\mathbb{E}$	IRT $\mathbb{E}$	IRT++ $\mathbb{E}$
ARC-C	39.71	46.21	42.32	0.80	1.80	1.86
GSM8k	4.10	3.21	4.62	0.41	1.16	1.49
Hellaswag	70.08	71.80	68.81	0.21	2.06	2.42

Table 4: **Monotonicity values for Tiny Benchmark.** We list the monotonicity values for both discrete ( $\text{mon}_{\text{disc}}$ ) and continuous ( $\text{mon}_{\text{cont}}$ ) metrics for the 7B seed models from § 3.2. Full represents the full benchmark, and IRT/IRT++ use the 100 examples proposed in Polo et al. (2024).

Benchmark	$\text{mon}_{\text{disc}}$ (Full/IRT/IRT++)	$\text{mon}_{\text{cont}}$ (Full/IRT/IRT++)
ARC-C	0.88 / 0.64 / 0.63	0.91 / 0.78 / 0.82
GSM8k	0.74 / 0.32 / 0.30	0.30 / 0.24 / 0.24
Hellaswag	0.99 / 0.84 / 0.80	0.99 / 0.93 / 0.94

out-of-distribution splits – it takes into account the entire spectrum of scores. However, even beyond the difficulty split, we similarly find little-to-no benefits on the random split (see Figure 7). As a result, we overall would not suggest the use of item analysis-based methods for understanding variance in language model evaluations, though the underlying cause for this mismatch remains an open question for future work.

## 5 The false promise of item response theory for LLMs

In a similar category to item analysis, *item response theory* (Cai et al., 2016; van der Linden, 2018; Brzezińska, 2020; Lord and Novick, 1968) describes a set of statistical models used to analyse human abilities on standardised test data. In the recent past, the method has become popular as a means of understanding model performance on a set of evaluation samples (Lalor et al., 2016; Vania et al., 2021; Rodriguez et al., 2021). Most recently, Polo et al. (2024) used IRT to cluster evaluation points with the aims of reducing eval benchmark size (and thus, cost of running).

Following our mixed findings applying item analysis, we apply the IRT method from Polo et al. (2024) to our models and the overlapping set of evaluation benchmarks. For a brief summary of the IRT method, we refer to Appendix E.1. Specifically, we go beyond the comparisons drawn in prior work and consider how our defined variance metrics (§ 3) change under this model. We believe the application to evaluating intermediate checkpoints during pretraining is especially relevant, as that’s the application where smaller evaluation datasets could have the most efficiency gains (as opposed to one-time evaluations of larger models).

In Tables 3 and 4, we report various metrics on the discrete performance measure for GSM8k, Hellaswag, and ARC-C. We find that simply using the performance on the 100 datapoints selected by Polo et al. (2024) for each benchmark can lead to quite large deviations in the mean (an overestimation by 7% for ARC-C). The full IRT++ method obtains less deviation, replicating prior findings (Polo et al., 2024). However, both methods suffer from greatly increased seed variance (final two columns, Table 3), indicating that the tiny-benchmarks method may have limited use during pretraining ablations as it makes model comparisons more likely to be confounded by randomness from the initialisation and data ordering seed. This increased variance is also reflected in the monotonicity metrics – we see a decrease in monotonicity in Table 4, indicating that performance oscillates more during training (see Figure 9).

Beyond the smaller scale models, we also considered the use of tiny-benchmarks for evaluating larger models, like the ones used for item analysis in Section 4. In Figure 4, we find that IRT-based methods generalise relatively well when it comes to the average performance metric (with the IRT++ estimator performing better), but have much larger standard error of the mean. This increased error cautions against the use of IRT-based subsets for model evaluations that will be used to compare different models. To quantify how this increased standard error of the mean may affect model rankings, we also compute the Kendall rank correlation on our 70 models using the performance estimate obtained

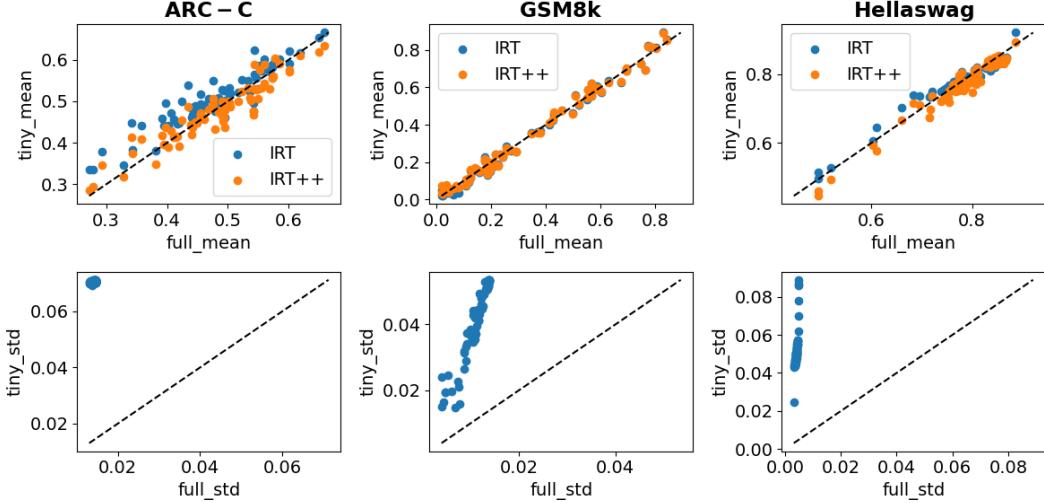


Figure 4: **Tiny Benchmarks Means and Standard Errors of the mean (proportional to 95% CI).**

from using the full dataset, as well as the IRT and IRT++ methods. In Table 7, we find that the correlation can drop as low as 0.76, corresponding to 12% of model pairwise comparisons giving the opposite result when using the IRT or IRT++ method (versus the full dataset mean estimate). Furthermore, we find that the number of flips is relatively higher on models that perform better, suggesting that IRT-based methods may not scale well (similar to item analysis). These findings reinforce the promise of IRT-based methods for a point estimate of the mean (relatively low error, Figure 4), but caution against the use of IRT-based methods when *comparing* models due to the increased variance of the estimate.

## 6 Related work

While a significant body of work exists proposing natural language processing (NLP) benchmarks to evaluate the capabilities of models, there is comparatively less work studying the benchmarks themselves. Before the era of chat large language models, Marie et al. (2021) conducted a large scale meta-analysis of 769 research papers published from 2010 to 2020 and identified troubling trends, including one that partially motivates our work: models are frequently declared superior to competitors based on small differences in performance scores, without proper hypothesis testing that takes into account natural fluctuations in benchmark scores. Spiritually similar claims were made by Dehghani et al. (2021) in their provocatively titled paper “The Benchmark Lottery”. Kocmi et al. (2021) further leveraged large-scale human experiments to evaluate benchmarks with automated metrics and concluded that commonly used metrics such as BLEU score had led to poor deployment decisions. Their conclusion was echoed by a meta-analysis of 3500 NLP benchmark scores published on Papers with Code (Blagec et al., 2022).

More recently, with accelerating progress in NLP, researchers have begun to study benchmarks in earnest to understand their properties and limitations (Gehrman et al., 2023). Von Werra et al. (2022) proposed a framework to evaluate benchmarks themselves and provided a mechanism for researchers to share their benchmarking analyses. Certain papers have studied specific aspects of benchmarks, focusing on the sensitivity of language models to various factors. Sclar et al. (2023) tested how sensitive language models are to differently formatted prompts, while Wang et al. (2024) and Alzahrani et al. (2024) find that models are inconsistent across changes in the format of Multiple-Choice Question Answering (MCQA) benchmarks. Our work builds on these works by focusing on the inherent variance in benchmarks (e.g. due to model seed) that practitioners should consider when making decisions, and suggesting minor modifications (e.g. in how a task is scored or formulated) that can reduce this variance.

With the aims of improving efficiency in model development cycles, recent work proposes reducing the size of evaluation benchmarks by picking representative samples (Vivek et al., 2023; Polo et al., 2024). Polo et al. (2024) show that methods from human standardised testing (specifically, item

response theory; Lord and Novick, 1968) can be combined with clustering to subselect evaluation benchmarks without incurring too much deviation from the mean. However, they do not consider the increased *variance* from their method nor how small deviations in means can compound when comparing multiple models. We go beyond their work by considering the use of additional methods from human standardised testing literature (item analysis; Livingston, 2011), as well as showing that such methods generally do not meaningfully reduce variance.

Perhaps most similar to ours is the work of Xiang et al. (2022), who study different sources of variance in NLP benchmarks and offers cautionary advice about when one should (not) be confident in benchmark scores. Their approach is limited to the machine translation setting; here we quantify and study variance in 13 different NLP benchmarks (covering general knowledge, reasoning, coding, and math) across 280 models, including many frontier LLMs.

## 7 Conclusion

As language models become more and more prevalent, it has become increasingly important to get a sense of their capabilities. One of the primary ways to assess these capabilities is through the use of evaluation benchmarks, where a model is scored on a series of examples. These scores are often directly compared, without consideration of the *variance*. This obscures the interpretation of evaluation results, in assessing final models as well as making decisions during model development. In this work, we aimed to quantify evaluation benchmark variance across a range of settings (from pretraining intermediate checkpoints, to the largest frontier LLMs) using a diverse set of metrics (seed variance, confidence intervals, and monotonicity). Beyond quantifying variance, we also experimented with various techniques used in human standardised testing (item analysis; University of Washington (2024), item response theory; Cai et al. (2016)), but generally found these methods to be ineffective on the models and benchmarks we considered, in terms of reducing variance. Future work could explore such avenues further, and it is possible that as models reach closer and closer to human-level performance these methods will provide more useful insights. On the other hand, in line with recent work advocating for a *teleological* approach to measuring capabilities (McCoy et al., 2023), we demonstrated LLM-specific techniques (e.g. the use of continuous metrics or cloze-formatted tasks) can improve the signal-to-noise ratio in our evals. Such techniques are not available when assessing humans, but provide a unique opportunity for LLM evaluations, especially when performing pretraining ablations. We hope our work spurs future work in this direction of reducing variance, in addition to serving as an empirical guide for model practitioners to use when comparing models and assessing performance.

## References

- Josh Achiam, Steven Adler, Sandhini Agarwal, Lama Ahmad, Ilge Akkaya, Florencia Leoni Aleman, Diogo Almeida, Janko Altenschmidt, Sam Altman, Shyamal Anadkat, et al. Gpt-4 technical report. *arXiv preprint arXiv:2303.08774*, 2023.
- AI@Meta. Llama 3 model card. 2024. URL [https://github.com/meta-llama/llama3/blob/main/MODEL\\_CARD.md](https://github.com/meta-llama/llama3/blob/main/MODEL_CARD.md).
- Ebtessam Almazrouei, Hamza Alobaidli, Abdulaziz Alshamsi, Alessandro Cappelli, Ruxandra Cojocaru, Merouane Debbah, Etienne Goffinet, Daniel Heslow, Julien Launay, Quentin Malartic, Badreddine Noune, Baptiste Pannier, and Guilherme Penedo. Falcon-40B: an open large language model with state-of-the-art performance. 2023.
- Norah Alzahrani, Hisham Abdullah Alyahya, Yazeed Alnumay, Sultan Alrashed, Shaykhah Alsubaie, Yusef Almushaykeh, Faisal Mirza, Nouf Alotaibi, Nora Altwaresh, Areeb Alowisheq, et al. When benchmarks are targets: Revealing the sensitivity of large language model leaderboards. *arXiv preprint arXiv:2402.01781*, 2024.
- Jinze Bai, Shuai Bai, Yunfei Chu, Zeyu Cui, Kai Dang, Xiaodong Deng, Yang Fan, Wenbin Ge, Yu Han, Fei Huang, Binyuan Hui, Luo Ji, Mei Li, Junyang Lin, Runji Lin, Dayiheng Liu, Gao Liu, Chengqiang Lu, Keming Lu, Jianxin Ma, Rui Men, Xingzhang Ren, Xuancheng Ren, Chuanqi Tan, Sinan Tan, Jianhong Tu, Peng Wang, Shijie Wang, Wei Wang, Shengguang Wu, Benfeng Xu, Jin Xu, An Yang, Hao Yang, Jian Yang, Shusheng Yang, Yang Yao, Bowen Yu, Hongyi Yuan, Zheng

Yuan, Jianwei Zhang, Xingxuan Zhang, Yichang Zhang, Zhenru Zhang, Chang Zhou, Jingren Zhou, Xiaohuan Zhou, and Tianhang Zhu. Qwen technical report. *arXiv preprint arXiv:2309.16609*, 2023.

Xiao Bi, Deli Chen, Guanting Chen, Shanhua Chen, Damai Dai, Chengqi Deng, Honghui Ding, Kai Dong, Qiushi Du, Zhe Fu, et al. Deepseek llm: Scaling open-source language models with longtermism. *arXiv preprint arXiv:2401.02954*, 2024.

Stella Biderman, Hailey Schoelkopf, Quentin Anthony, Herbie Bradley, Kyle O'Brien, Eric Hallahan, Mohammad Aftab Khan, Shivanshu Purohit, USVSN Sai Prashanth, Edward Raff, Aviya Skowron, Lintang Sutawika, and Oskar van der Wal. Pythia: A suite for analyzing large language models across training and scaling, 2023.

Yonatan Bisk, Rowan Zellers, Ronan Le Bras, Jianfeng Gao, and Yejin Choi. Piqa: Reasoning about physical commonsense in natural language. 2020.

Kathrin Blagec, Georg Dorffner, Milad Moradi, Simon Ott, and Matthias Samwald. A global analysis of metrics used for measuring performance in natural language processing, 2022.

Justyna Brzezińska. Item response theory models in the measurement theory. *Communications in Statistics - Simulation and Computation*, 49(12):3299–3313, 2020. doi: 10.1080/03610918.2018.1546399. URL <https://doi.org/10.1080/03610918.2018.1546399>.

Li Cai, Kilchan Choi, Mark Hansen, and Lauren Harrell. Item response theory. *Annual Review of Statistics and Its Application*, 3(Volume 3, 2016):297–321, 2016. ISSN 2326-831X. doi: <https://doi.org/10.1146/annurev-statistics-041715-033702>. URL <https://www.annualreviews.org/content/journals/10.1146/annurev-statistics-041715-033702>.

Mark Chen, Jerry Tworek, Heewoo Jun, Qiming Yuan, Henrique Ponde de Oliveira Pinto, Jared Kaplan, Harri Edwards, Yuri Burda, Nicholas Joseph, Greg Brockman, Alex Ray, Raul Puri, Gretchen Krueger, Michael Petrov, Heidy Khlaaf, Girish Sastry, Pamela Mishkin, Brooke Chan, Scott Gray, Nick Ryder, Mikhail Pavlov, Alethea Power, Lukasz Kaiser, Mohammad Bavarian, Clemens Winter, Philippe Tillet, Felipe Petroski Such, Dave Cummings, Matthias Plappert, Fotios Chantzis, Elizabeth Barnes, Ariel Herbert-Voss, William Hebgen Guss, Alex Nichol, Alex Paino, Nikolas Tezak, Jie Tang, Igor Babuschkin, Suchir Balaji, Shantanu Jain, William Saunders, Christopher Hesse, Andrew N. Carr, Jan Leike, Josh Achiam, Vedant Misra, Evan Morikawa, Alec Radford, Matthew Knight, Miles Brundage, Mira Murati, Katie Mayer, Peter Welinder, Bob McGrew, Dario Amodei, Sam McCandlish, Ilya Sutskever, and Wojciech Zaremba. Evaluating large language models trained on code. 2021.

Peter Clark, Isaac Cowhey, Oren Etzioni, Tushar Khot, Ashish Sabharwal, Carissa Schoenick, and Oyvind Tafjord. Think you have solved question answering? try arc, the ai2 reasoning challenge. *arXiv preprint arXiv:1803.05457*, 2018.

Karl Cobbe, Vineet Kosaraju, Mohammad Bavarian, Mark Chen, Heewoo Jun, Lukasz Kaiser, Matthias Plappert, Jerry Tworek, Jacob Hilton, Reiichiro Nakano, et al. Training verifiers to solve math word problems. *arXiv preprint arXiv:2110.14168*, 2021.

Databricks. Dbrx technical blog. 2024. URL <https://www.databricks.com/blog/introducing-dbrx-new-state-art-open-lm>.

DeepSeek-AI. Deepseek-v2: A strong, economical, and efficient mixture-of-experts language model, 2024.

Mostafa Dehghani, Yi Tay, Alexey A. Gritsenko, Zhe Zhao, Neil Houlsby, Fernando Diaz, Donald Metzler, and Oriol Vinyals. The benchmark lottery, 2021.

Zhengxiao Du, Aohan Zeng, Yuxiao Dong, and Jie Tang. Understanding emergent abilities of language models from the loss perspective, 2024.

Sebastian Gehrmann, Elizabeth Clark, and Thibault Sellam. Repairing the cracked foundation: A survey of obstacles in evaluation practices for generated text. *Journal of Artificial Intelligence Research*, 77:103–166, 2023.

Dan Hendrycks, Collin Burns, Steven Basart, Andy Zou, Mantas Mazeika, Dawn Song, and Jacob Steinhardt. Measuring massive multitask language understanding. *arXiv preprint arXiv:2009.03300*, 2020.

Dan Hendrycks, Collin Burns, Saurav Kadavath, Akul Arora, Steven Basart, Eric Tang, Dawn Song, and Jacob Steinhardt. Measuring mathematical problem solving with the math dataset. 2021.

Albert Q Jiang, Alexandre Sablayrolles, Arthur Mensch, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Florian Bressand, Gianna Lengyel, Guillaume Lample, Lucile Saulnier, et al. Mistral 7b. *arXiv preprint arXiv:2310.06825*, 2023.

Albert Q. Jiang, Alexandre Sablayrolles, Antoine Roux, Arthur Mensch, Blanche Savary, Chris Bamford, Devendra Singh Chaplot, Diego de las Casas, Emma Bou Hanna, Florian Bressand, Gianna Lengyel, Guillaume Bour, Guillaume Lample, Lélio Renard Lavaud, Lucile Saulnier, Marie-Anne Lachaux, Pierre Stock, Sandeep Subramanian, Sophia Yang, Szymon Antoniak, Teven Le Scao, Théophile Gervet, Thibaut Lavril, Thomas Wang, Timothée Lacroix, and William El Sayed. Mixtral of experts, 2024.

Mandar Joshi, Eunsol Choi, Daniel Weld, and Luke Zettlemoyer. triviaqa: A Large Scale Distantly Supervised Challenge Dataset for Reading Comprehension. *arXiv e-prints*, art. arXiv:1705.03551, 2017.

Tom Kocmi, Christian Federmann, Roman Grundkiewicz, Marcin Junczys-Dowmunt, Hitokazu Matsushita, and Arul Menezes. To ship or not to ship: An extensive evaluation of automatic metrics for machine translation, 2021.

Tom Kwiatkowski, Jennimaria Palomaki, Olivia Redfield, Michael Collins, Ankur Parikh, Chris Alberti, Danielle Epstein, Illia Polosukhin, Matthew Kelcey, Jacob Devlin, Kenton Lee, Kristina N. Toutanova, Llion Jones, Ming-Wei Chang, Andrew Dai, Jakob Uszkoreit, Quoc Le, and Slav Petrov. Natural questions: a benchmark for question answering research. *Transactions of the Association of Computational Linguistics*, 2019.

John P Lalor, Hao Wu, and Hong Yu. Building an evaluation scale using item response theory. In *Proceedings of the Conference on Empirical Methods in Natural Language Processing. Conference on Empirical Methods in Natural Language Processing*, volume 2016, page 648. NIH Public Access, 2016.

Samuel A Livingston. Item analysis. In *Handbook of test development*, pages 435–456. Routledge, 2011.

F. M. Lord and M. R. Novick. *Statistical theories of mental test scores*. Addison-Wesley, 1968.

Benjamin Marie, Atsushi Fujita, and Raphael Rubino. Scientific credibility of machine translation research: A meta-evaluation of 769 papers. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*. Association for Computational Linguistics, 2021.

R. Thomas McCoy, Shunyu Yao, Dan Friedman, Matthew Hardy, and Thomas L. Griffiths. Embers of autoregression: Understanding large language models through the problem they are trained to solve, 2023.

Thomas Mesnard, Cassidy Hardin, Robert Dadashi, Surya Bhupatiraju, Shreya Pathak, Laurent Sifre, Morgane Rivière, Mihir Sanjay Kale, Juliette Love, Pouya Tafti, Léonard Hussenot, Aakanksha Chowdhery, Adam Roberts, Aditya Barua, Alex Botev, Alex Castro-Ros, Ambrose Slone, Amélie Héliou, Andrea Tacchetti, Anna Bulanova, Antonia Paterson, Beth Tsai, Bobak Shahriari, Charline Le Lan, Christopher A. Choquette-Choo, Clément Crepy, Daniel Cer, Daphne Ippolito, David Reid, Elena Buchatskaya, Eric Ni, Eric Noland, Geng Yan, George Tucker, George-Cristian Murraru, Grigory Rozhdestvenskiy, Henryk Michalewski, Ian Tenney, Ivan Grishchenko, Jacob Austin, James Keeling, Jane Labanowski, Jean-Baptiste Lespiau, Jeff Stanway, Jenny Brennan, Jeremy Chen, Johan Ferret, Justin Chiu, and et al. Gemma: Open models based on gemini research and technology. *arXiv preprint arXiv:2403.08295*, 2024.

MosaicML NLP Team. Introducing mpt-7b: A new standard for open-source, commercially usable llms, 2023. URL [www.mosaicml.com/blog/mpt-7b](http://www.mosaicml.com/blog/mpt-7b). Accessed: 2023-05-05.

Felipe Maia Polo, Lucas Weber, Leshem Choshen, Yuekai Sun, Gongjun Xu, and Mikhail Yurochkin. tinybenchmarks: evaluating llms with fewer examples, 2024.

Machel Reid, Nikolay Savinov, Denis Teplyashin, Dmitry Lepikhin, Timothy Lillicrap, Jean-baptiste Alayrac, Radu Soricut, Angeliki Lazaridou, Orhan Firat, Julian Schrittweiser, et al. Gemini 1.5: Unlocking multimodal understanding across millions of tokens of context. *arXiv preprint arXiv:2403.05530*, 2024.

Pedro Rodriguez, Joe Barrow, Alexander Miserlis Hoyle, John P Lalor, Robin Jia, and Jordan Boyd-Graber. Evaluation examples are not equally informative: How should that change nlp leaderboards? In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics and the 11th International Joint Conference on Natural Language Processing (Volume 1: Long Papers)*, pages 4486–4503, 2021.

Melissa Roemmele, Cosmin Adrian Bejan, and Andrew S Gordon. Choice of plausible alternatives: An evaluation of commonsense causal reasoning. In *2011 AAAI Spring Symposium Series*, 2011.

Maarten Sap, Hannah Rashkin, Derek Chen, Ronan Le Bras, and Yejin Choi. Social iq: Commonsense reasoning about social interactions. In *Proceedings of the 2019 Conference on Empirical Methods in Natural Language Processing and the 9th International Joint Conference on Natural Language Processing (EMNLP-IJCNLP)*, 2019.

Rylan Schaeffer, Brando Miranda, and Sanmi Koyejo. Are emergent abilities of large language models a mirage?, 2023.

Rylan Schaeffer, Hailey Schoelkopf, Brando Miranda, Gabriel Mukobi, Varun Madan, Adam Ibrahim, Herbie Bradley, Stella Biderman, and Sanmi Koyejo. Why has predicting downstream capabilities of frontier ai models with scale remained elusive?, 2024.

Melanie Sclar, Yejin Choi, Yulia Tsvetkov, and Alane Suhr. Quantifying language models’ sensitivity to spurious features in prompt design or: How i learned to start worrying about prompt formatting. *arXiv preprint arXiv:2310.11324*, 2023.

Aarohi Srivastava, Abhinav Rastogi, Abhishek Rao, Abu Awal Md Shoeb, Abubakar Abid, Adam Fisch, Adam R Brown, Adam Santoro, Aditya Gupta, Adrià Garriga-Alonso, et al. Beyond the imitation game: Quantifying and extrapolating the capabilities of language models. *arXiv preprint arXiv:2206.04615*, 2022.

StabilityAI. Stablelm technical report. 2024. URL <https://stability.wandb.io/stability-llm/stable-lm/reports/StableLM-3B-4E1T--Vm1ldzoyMjU4?accessToken=u3zujipenkx5g7rtcj9qojjgxpconyjktjkli2po09nffrffdhchq045vp0wyfo>.

Mirac Suzgun, Nathan Scales, Nathanael Schärlí, Sebastian Gehrmann, Yi Tay, Hyung Won Chung, Aakanksha Chowdhery, Quoc V Le, Ed H Chi, Denny Zhou, et al. Challenging big-bench tasks and whether chain-of-thought can solve them. *arXiv preprint arXiv:2210.09261*, 2022.

Hugo Touvron, Thibaut Lavril, Gautier Izacard, Xavier Martinet, Marie-Anne Lachaux, Timothée Lacroix, Baptiste Rozière, Naman Goyal, Eric Hambro, Faisal Azhar, Aurelien Rodriguez, Armand Joulin, Edouard Grave, and Guillaume Lample. Llama: Open and efficient foundation language models, 2023a.

Hugo Touvron, Louis Martin, Kevin Stone, Peter Albert, Amjad Almahairi, Yasmine Babaei, Nikolay Bashlykov, Soumya Batra, Prajjwal Bhargava, Shruti Bhosale, et al. Llama 2: Open foundation and fine-tuned chat models. *arXiv preprint arXiv:2307.09288*, 2023b.

University of Washington. Understanding item analyses, 2024. URL <https://www.washington.edu/assessment/scanning-scoring/scoring/reports/item-analysis/>.

Wim J. van der Linden, editor. *Handbook of Item Response Theory: Three Volume Set*. CRC Press, 2018.

Clara Vania, Phu Mon Htut, William Huang, Dhara Mungra, Richard Yuanzhe Pang, Jason Phang, Haokun Liu, Kyunghyun Cho, and Samuel R Bowman. Comparing test sets with item response theory. *arXiv preprint arXiv:2106.00840*, 2021.

Rajan Vivek, Kawin Ethayarajh, Diyi Yang, and Douwe Kiela. Anchor points: Benchmarking models with much fewer examples, 2023.

Leandro Von Werra, Lewis Tunstall, Abhishek Thakur, Sasha Luccioni, Tristan Thrush, Aleksandra Piktus, Felix Marty, Nazneen Rajani, Victor Mustar, and Helen Ngo. Evaluate & evaluation on the hub: Better best practices for data and model measurements. In Wanxiang Che and Ekaterina Shutova, editors, *Proceedings of the 2022 Conference on Empirical Methods in Natural Language Processing: System Demonstrations*, pages 128–136, Abu Dhabi, UAE, December 2022. Association for Computational Linguistics. doi: 10.18653/v1/2022.emnlp-demos.13. URL <https://aclanthology.org/2022.emnlp-demos.13>.

Haochun Wang, Sendong Zhao, Zewen Qiang, Bing Qin, and Ting Liu. Beyond the answers: Reviewing the rationality of multiple choice question answering for the evaluation of large language models. *arXiv preprint arXiv:2402.01349*, 2024.

Lucas Weber, Elia Bruni, and Dieuwke Hupkes. Mind the instructions: a holistic evaluation of consistency and interactions in prompt-based learning. *arXiv preprint arXiv:2310.13486*, 2023.

Thomas Wolf, Lysandre Debut, Victor Sanh, Julien Chaumond, Clement Delangue, Anthony Moi, Pierrick Cistac, Tim Rault, Rémi Louf, Morgan Funtowicz, et al. Transformers: State-of-the-art natural language processing. In *Proceedings of the 2020 conference on empirical methods in natural language processing: system demonstrations*, pages 38–45, 2020.

Jiannan Xiang, Huayang Li, Yahui Liu, Lemao Liu, Guoping Huang, Defu Lian, and Shuming Shi. Investigating data variance in evaluations of automatic machine translation metrics. *arXiv preprint arXiv:2203.15858*, 2022.

Rowan Zellers, Ari Holtzman, Yonatan Bisk, Ali Farhadi, and Yejin Choi. Hellaswag: Can a machine really finish your sentence? *arXiv preprint arXiv:1905.07830*, 2019.

Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, and Minlie Huang. Large language models are not robust multiple choice selectors. In *The Twelfth International Conference on Learning Representations*, 2023.

Wanjun Zhong, Ruixiang Cui, Yiduo Guo, Yaobo Liang, Shuai Lu, Yanlin Wang, Amin Saied, Weizhu Chen, and Nan Duan. Agieval: A human-centric benchmark for evaluating foundation models, 2023.

## A Models and Benchmarks Details

For pre-training the 7B Llama-2 like checkpoints, we use a pre-training mix of publicly available data. We apply filtering to remove documents containing a high amount of personal information. We use a learning rate of  $3.0 \times 10^4$ , sequence length of 4096, and a batch size of  $4.1M$  tokens to train the 7B models for 50000 steps. We use 256 80GiB A100 GPUs for a single pre-training run for 50k steps on our internal cluster. We do 10 such runs with different seeds. Each step takes 4.3 seconds.

For running the evaluations, we use 8 GPUs for each evaluation job comprising multiple evaluation datasets in a single job. A single evaluation job takes on average takes 3.5 hours for 13 benchmarks.

In Table 5, we provide the discrete metric (preferred), the continuous metric, and the number of samples for each of the benchmarks we consider. We can choose any continuous metric like character NLL, raw NLL, probability mass, log of probabilities, etc. for the benchmarks, but to maintain consistency, we choose probability mass of the predicted answer for choice-based tasks and negative log likelihood (NLL) of the target answer for generation-based benchmarks. Choice-based benchmarks are evaluated by appending the possible option choice letters or choice texts and then choosing the option with the lowest NLL. Generation-based benchmarks involve free-form generation, where the answer is extracted from the model’s response using various post-processing techniques.

Table 5: **Benchmark Details** Details of all benchmarks used in the paper alphabetically. Exact Match (EM) is computed for 1 generation (maj@1). Prob Mass is the probability mass of the predicted answer and Target NLL represents the NLL of the target answer. CoT represents chain of thought prompting.

Benchmark	License	# samples	# few-shot	Disc Metric	Cont Metric
AGIEval (Zhong et al., 2023)	MIT	2546	3-5	Acc	Prob Mass
ARC-C (Clark et al., 2018)	Apache 2.0	1165	0	Acc	Prob Mass
Big Bench Hard (Srivastava et al., 2022)	Apache 2.0	6511	3 (CoT)	EM	-
COPA (Roemmel et al., 2011)	BSD 2-Clause	100	0	Acc	Prob Mass
GSM8k (Cobbe et al., 2021)	MIT	1319	8 (CoT)	EM	Target NLL
Hellaswag Zellers et al. (2019)	MIT	10042	0	Acc	Prob Mass
HumanEval (Chen et al., 2021)	MIT	164	0	Pass@1	Target NLL
MATH (Hendrycks et al., 2021)	MIT	5000	4 (CoT)	EM	-
MMLU (Hendrycks et al., 2020)	MIT	14042	5	Acc	Prob Mass
Natural Questions (Kwiatkowski et al., 2019)	MIT	3610	5	EM	-
PIQA (Bisk et al., 2020)	Academic Free	1838	0	Acc	Prob Mass
SIQA Sap et al. (2019)	-	1954	0	Acc	Prob Mass
TriviaQA Joshi et al. (2017)	Apache 2.0	11313	5	EM	-

Table 6: **Model Details** Details of all models in the paper categorized by model family along with the number of parameters.

Model Family	Models	Model Sizes (# params)
Meta-Llama (AI@Meta, 2024; Touvron et al., 2023b,a)	Llama-1, Llama-2, Llama-3	7-70B
Google (Mesnard et al., 2024)	Gemma	2-7B
Databricks (Databricks, 2024)	DBRX-Base	132B
Mistral (Jiang et al., 2023, 2024)	Mistral, Mixtral	7-141B
Qwen (Bai et al., 2023)	Qwen-1.5	0.5-110B
EleutherAI (Biderman et al., 2023)	Pythia	1-12B
TII-UAE (Almazrouei et al., 2023)	Falcon	7-40B
DeepSeek (Bi et al., 2024; DeepSeek-AI, 2024)	DeepSeek, DeepSeek-MoE, DeepSeek-V2	7-236B
StabilityAI (StabilityAI, 2024)	StableLM	1.6-7B
MosaicML (MosaicML NLP Team, 2023)	MPT	7-30B

## B MMLU prompt formats

We use the following prompt variations for the standard and cloze versions of MMLU. We list down the preamble and the shot formatting for both cases. The final question is formatted like the few shot examples without the gold choice letter or text.

### B.1 MMLU

**Preamble:**

The following are multiple choice questions (with answers) about *<subject>*.

**Shot formatting:**

*<question>*

- A. *<choice A text>*
- B. *<choice B text>*
- C. *<choice C text>*
- D. *<choice D text>*

Answer: *<gold choice letter>*

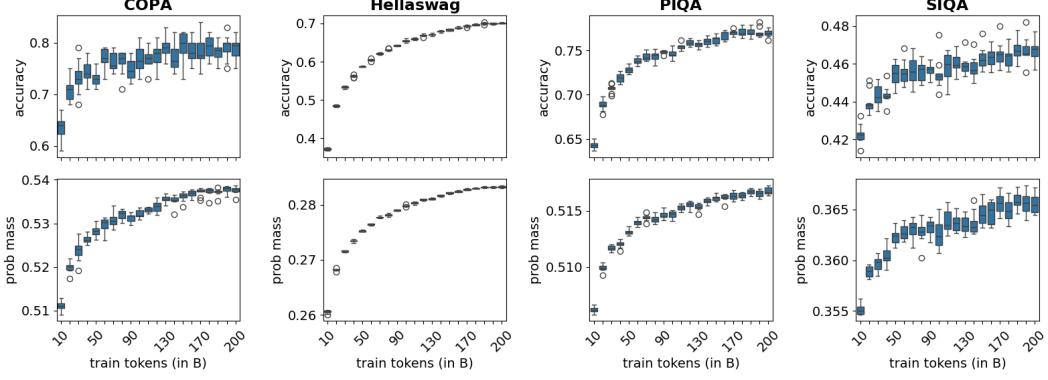
### B.2 MMLU-cloze

**Preamble:**

**Shot formatting:**

*<question>*

Answer: *<gold choice text>*



**Figure 5: Development of model performance over time.** Boxplots for both discrete and continuous metrics depicting the model improvement over time for COPA, Hellaswag, PIQA, and SIQA. Top row depicts discrete metrics for each of the benchmarks, and the bottom row is composed of the continuous metrics.

## C Variance Analysis Additional Results

In this section, we present additional results on model performance development for the remaining benchmarks - COPA, Hellaswag, PIQA, and SIQA (see Figure 5). This supplements the results presented in Figure 1 and Figure 2. We observe similar trends except for SIQA. The error bars for both discrete and continuous metrics are similar, however, the continuous metric plot has less number of outliers.

## D Item analysis additional results

### D.1 Splits

We used 70 base models for the item analysis results. We provide the splits used below.

**Difficulty split (train):** LLaMa 3 8B, Mistral 7B, Qwen {0.5, 1.8, 4}B, LLaMa 2 7B, LLaMa 2 13B, LLaMa 2 70B, DeepSeek 7B, DeepSeek MoE 16B, Falcon 7B, Falcon 40B, Gemma 2B, Gemma 7B, LLaMa 1 {7, 13, 33, 65} B, MPT 30B, Pythia {1, 1.4, 2.8, 6.9, 12}B, StableLM {3, 7}B. In addition to these open source models, we use 30 internal checkpoints from LLaMa-architecture models we pre-trained on our internal data mix.

**Difficulty split (test):** LLaMa 3 70B, Mixtral 8x{7,22}B, Qwen 1.5 {7, 13, 32, 72, 110}B, DBRX, DeepSeek 67B, and 4 internal held out models.

**Random split (train):** LLaMa 3 {8, 70}B, Mistral 7B, Mixtral 8x{7,22}B, Qwen 1.5 {0.5, 1.8, 4, 7, 13, 32, 72}B, LLaMa 2 7B, LLaMa 2 13B, LLaMa 2 70B, DBRX, DeepSeek MoE 16B, DeepSeek 67B, Falcon 40B, Gemma 2B, Gemma 7B, LLaMa 1 {7, 33, 65} B, MPT 30B, Pythia {1, 1.4, 2.8, 6.9, 12}B, StableLM 3B. In addition to these open source models, we use 25 internal checkpoints from LLaMa-architecture models we pre-trained on our internal data mix.

**Random split (test):** DeepSeek 7B, Falcon 7B, Qwen 1.5 110B, LLaMa 1 13B, StableLM 7B, and 9 internal checkpoints.

### D.2 Additional results

We present results on additional benchmarks, in a similar format to Figure 3, in Figure 6. Furthermore, we provide extended results on the random split of models in Figure 7.

### D.3 Inspection of samples with low item discrimination

We provide the 3 items from GSM8k, ARC-C and Hellaswag with the lowest item discrimination.

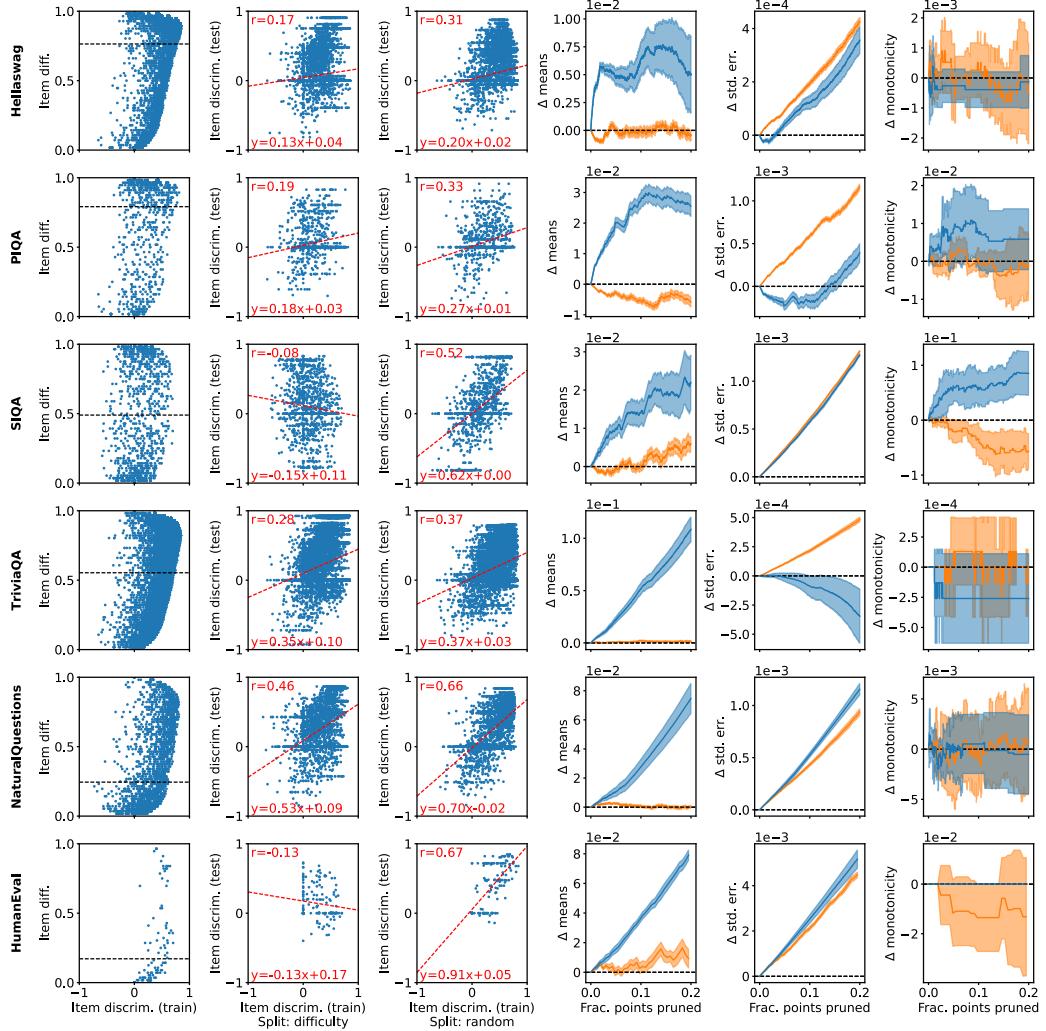


Figure 6: Item analysis results on six additional benchmarks, in the same format as Figure 3.

**For GSM8k:**

**Question:**

Aaron and Vanessa were relay race partners on a running team. Aaron was able to run each mile twice as fast as Vanessa, but Vanessa was able to run twice as far as Aaron did. If Vanessa ran 4 miles and Aaron completed his part of the race in 16 minutes, how long in minutes did Vanessa take to complete her part?

**Answer:** 64

**Item Discrimination:** -0.264

**Item Difficulty:** 0.1